

Control Flow Optimizations

Compilers for High Performance
Architectures



Class Outline

- Control flow analysis
- Program profiling
- Control-flow optimizations
 - Branch to unconditional branch
 - Unconditional branch to branch
 - Branch to next basic block
 - Basic block merging
 - Branch to same target
 - Branch target expansion
 - Unreachable code elimination
 - Code hoisting
 - Partial predication
 - Full predication

Control Flow

- Control flow is only known at run-time
 - Branch speculation
 - Wrong path squash
- Compiler only knows the static control flow
 - The actual values are unknown
 - The actual control flow is unknown
 - Control flow analysis
 - Determine properties of the program control flow

Basic Block (BB)


- Group operations into units with equivalent execution conditions
- Basic block – a sequence of consecutive operations in which flow of control enters at the beginning and leaves at the end without halt or possibility of branching except at the end
 - Straight-line sequence of instructions
 - If one operation is executed in a BB, they all are
- Finding BB's
 - The first operation starts a BB
 - Any operation that is the target of a branch starts a BB
 - Any operation that immediately follows a branch starts a BB

Identifying BBs - Example

```
L1: r7 = load(r8)
L2: r1 = r2 + r3
L3: beq r1, 0, L10
L4: r4 = r5 * r6
L5: r1 = r1 + 1
L6: beq r1 100 L2
L7: beq r2 100 L10
L8: r5 = r9 + 1
L9: r7 = r7 & 3
L10: r9 = load (r3)
L11: store(r9, r1)
```

Control Flow Graph (CFG)



- Directed graph, $G=(V,E)$ where each vertex V is a basic block, and there is an edge $E=(v1,v2)$ if BB $v1$ can branch to BB $v2$
 - Standard representation used by many compilers
 - Often have 2 pseudo V 's
 - entry node
 - exit node



Dominator


- Dominator – Given a CFG, a node x dominates a node y , if every path from the Entry block to y contains x
- 3 properties of dominators
 - Each BB dominates itself
 - If x dominates y , and y dominates z , then x dominates z
 - If x dominates z and y dominates z , then either x dominates y or y dominates x
- Intuition
 - Given some BB, which blocks are guaranteed to have executed prior to executing the BB

Dominator Examples




Dominator Analysis

- Compute $\text{dom}(\text{BB}_i)$ = set of BBs that dominate BB_i
- Initialization
 - $\text{Dom}(\text{entry}) = \text{entry}$
 - $\text{Dom}(\text{everything else}) = \text{all nodes}$
- Iterative computation
 - while change, do
 - change = false
 - for each BB (except the entry BB)
 - $\text{tmp}(\text{BB}) = \text{BB} + \{\text{intersect of Dom of all predecessor BB's}\}$
 - if ($\text{tmp}(\text{BB}) \neq \text{dom}(\text{BB})$)
 - $\text{dom}(\text{BB}) = \text{tmp}(\text{BB})$
 - change = true



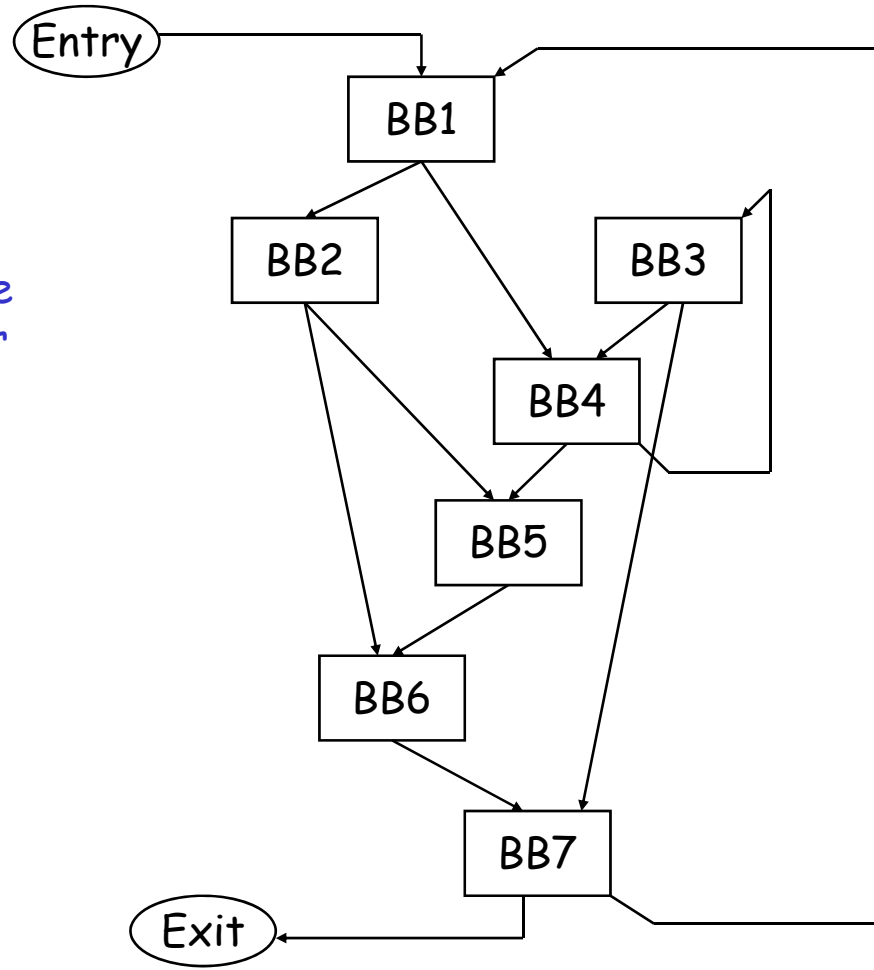
Immediate Dominator

- Each node n has a unique immediate dominator m that is the last dominator of n on any path from the initial node to n
- Closest node that dominates



Class Problem 1



Calculate the
DOM set for
each BB



Post Dominator


- Reverse of dominator
- Post Dominator – Given a CFG($V, E, \text{Entry}, \text{Exit}$), a node x post dominates a node y , if every path from y to the Exit contains x
- Intuition
 - Given some BB, which blocks are guaranteed to have executed after executing the BB

Post Dominator Examples




Post Dominator Analysis

- Compute $\text{pdom}(\text{BB}_i)$ = set of BBs that post dominate BB_i
- Initialization
 - $\text{Pdom}(\text{exit}) = \text{exit}$
 - $\text{Pdom}(\text{everything else}) = \text{all nodes}$
- Iterative computation
 - while change, do
 - change = false
 - for each BB (except the exit BB)
 - $\text{tmp}(\text{BB}) = \text{BB} + \{\text{intersect of pdom of all successor BB's}\}$
 - if ($\text{tmp}(\text{BB}) \neq \text{pdom}(\text{BB})$)
 - $\text{pdom}(\text{BB}) = \text{tmp}(\text{BB})$
 - change = true




Immediate Post Dominator

- Each node n has a unique immediate post dominator m that is the first post dominator of n on any path from n to the Exit
- Closest node that post dominates




Class Problem 2

Calculate the
PDOM set for
each BB




Why Do We Care About Dominators?

- Loop detection – next subject
- Dominator
 - Guaranteed to execute before
 - Redundant computation – an op is redundant if it is computed in a dominating BB
 - Most global optimizations use dominance info
- Post dominator
 - Guaranteed to execute after
 - Make a guess (ie 2 pointers do not point to the same locn)
 - Check they really do not point to one another in the post dominating BB



Weighted CFG

- Add profiling information to the CFG
 - Basic block execution counts
 - Edge execution counts




Program Profiling


- Many optimizations require run-time knowledge
 - Help classic optimizations
 - Apply optimizations to the common case
 - Make the right choice
 - Code predication
 - Data speculation
- Goals
 - High accuracy of data
 - Profile using a representative input
 - Low overhead
 - Do not change the actual runtime behavior: Race conditions, Spin-locks

Profiling usage example


Original:



Transformed:



After CSE:



Basic block profiling


- Code instrumentation
 - Add code to the beginning of each basic block
 - Counts the number of times the BB executed
 - Add code to the exit point of the program to dump data
- Assign an identifier to each basic block
- Allocate a counters in memory
 - Consecutive array
 - Assign blocks a unique index
 - Group counters to optimize data cache misses
 - Assign blocks a unique address

Edge profiling


- Code instrumentation
 - Add code to the beginning of each basic block
 - Build the edge identifier using previous and current BB identifier
 - Set current BB as "previous BB"
 - Add code to the exit point of the program to dump data
- Assign an identifier to each basic block
- Allocate a counters in memory
 - Many more edges than blocks
 - Higher cache miss rate
 - Higher performance impact

BB and Edge profiling example

Basic bloc profiling



Edge profiling



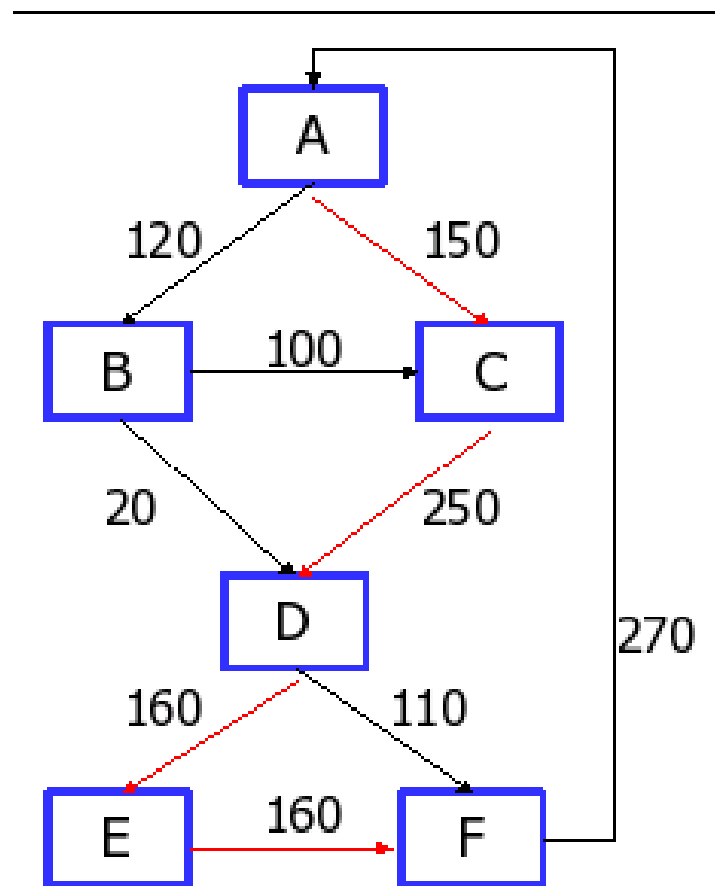
Path profiling

- Edge profiling not always correct

- Common path should be ACDEF
- But ...

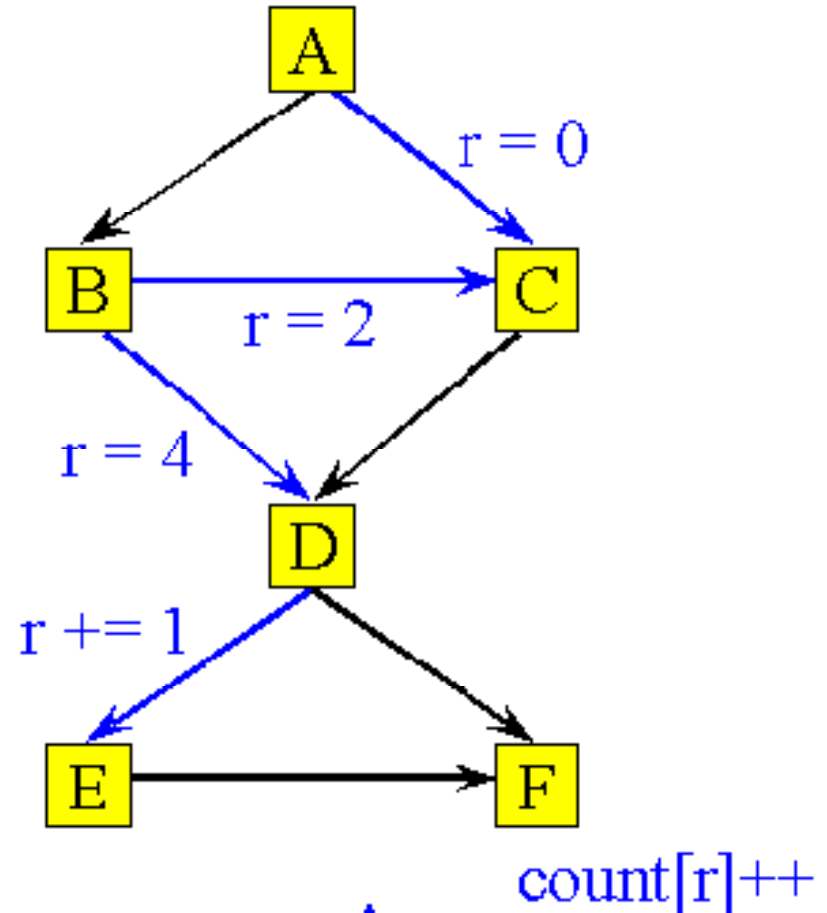
Path	prof1	prof2
ACDF	90	110
ACDEF	60	40
ABCDF	0	0
ABCDEF	100	100
ABDF	20	0
ABDEF	0	20

- ... also satisfy the edge profile



Efficient path profiling

- Trace the path through a region
 - Instrument all basic blocks
 - Maintain a list structure
 - Very expensive!
- Build a unique path ID as we progress
 - Instrument only selected edges
 - No need for a list structure
 - Still near 40% overhead vs. 15% for Edge profiling
- Statistic profiling
 - Turn profiling on and off to obtain data at random intervals
 - Aggregate data is still representative
 - Profiling overhead is reduced



Control Flow Optimizations for Acyclic Code

- Generally quite simplistic
- Goals
 - Reduce the number of dynamic branches
 - Make larger basic blocks
 - Reduce code size
- Classic control flow optimizations
 - Branch to unconditional branch
 - Unconditional branch to branch
 - Branch to next basic block
 - Basic block merging
 - Branch to same target
 - Branch target expansion
 - Unreachable code elimination


Control Flow Optimizations (1)

1. Branch to unconditional branch

L1: if (a < b) goto L2

...

L2: goto L3



L1: if (a < b) goto L3

...

L2: goto L3 → may be deleted


2. Unconditional branch to branch

L1: goto L2

...

L2: if (a < b) goto L3

L4:



L1: if (a < b) goto L3

goto L4:


...

L2: if (a < b) goto L3 → may be deleted


L4:

Control Flow Optimizations (2)


3. Branch to next basic block




Branch is unnecessary



4. Basic block merging




Merge BBs when single edge between



Control Flow Optimizations (3)



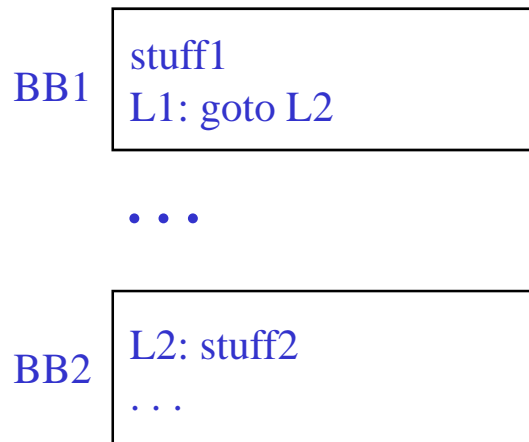
5. Branch to same target

...
L1: if (a < b) goto L2
goto L2



...
L1: goto L2

6. Branch target expansion




What about expanding a conditional branch?

Unreachable Code Elimination

Algorithm

```
Mark procedure entry BB visited
to_visit = procedure entry BB
while (to_visit not empty) {
    current = to_visit.pop()
    for (each successor block of current) {
        Mark successor as visited;
        to_visit += successor
    }
}
Eliminate all unvisited blocks
```



Which BB(s) can be deleted?

Class Problem 3

Maximally optimize the control flow of this code


L1: if (a < b) goto L11
L2: goto L7
L3: goto L4
L4: stuff4
L5: if (c < d) goto L15
L6: goto L2
L7: if (c < d) goto L13
L8: goto L12
L9: stuff 9
L10: if (a < c) goto L3
L11: goto L9
L12: goto L2
L13: stuff 13
L14: if (e < f) goto L11
L15: stuff 15
L16: rts

Code hoisting

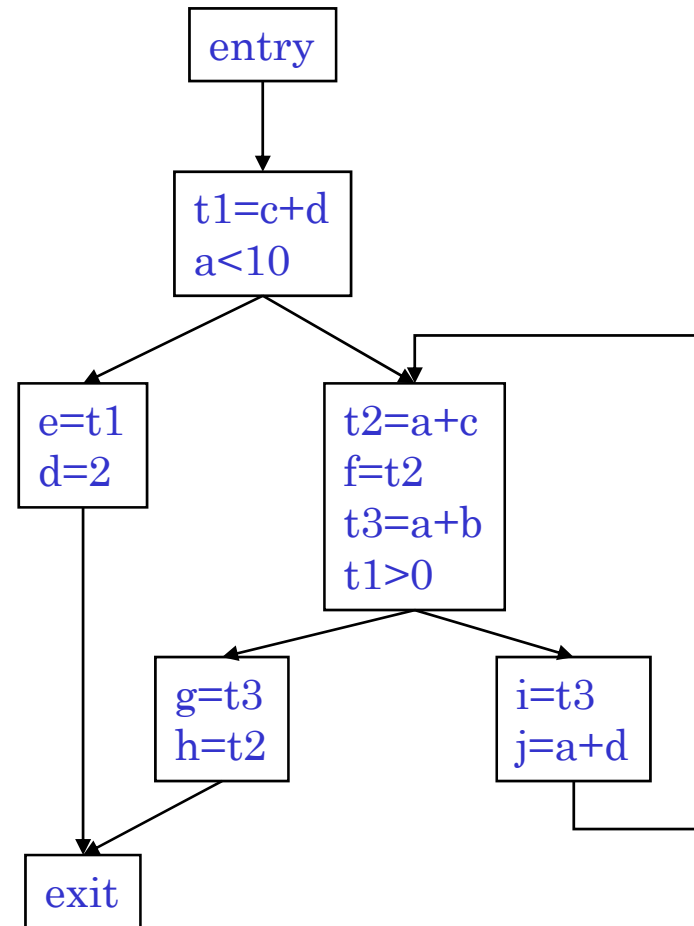
- Detect expressions which are always evaluated after a point in program execution
 - Regardless of control flow
- Move computation to the earliest point in the code where it can be executed
 - Add move / rename operations as needed
- In addition we can perform speculative code hoisting
 - Speculative loads
 - Speculative computation
 - speculative stores (dangerous)

Code hoisting example

- Before



- After



(partial) Predicated execution

- Conditional moves
 - Replace conditional code for code with no branches
 - Not all conditional codes can be converted
 - Limited to arithmetic operations
 - Can not store results to memory

Branching code:

```
        if (a > b) goto L1
        max = b
        goto L2
L1:     max = a
L2:     ...
```

Partially predicated code:

```
t1 = a > b
max = b
max = a   if t1
```

Predicated Execution


- Hardware mechanism that allows operations to be conditionally executed
- Add an additional boolean source operand (predicate)
 - ADD r1, r2, r3 if p1
 - if (p1 is True), $r1 = r2 + r3$
 - else if (p1 is False), do nothing (Add treated like a NOP)
 - p1 referred to as the guarding predicate
 - Predicated on True means always executed
 - Omitted predicated also means always executed
- Provides compiler with an alternative to using branches to selectively execute operations
 - If statements in the source
 - Realize with branches in the assembly code
 - Could also realize with conditional instructions
 - Or use a combination of both

Predicated Execution Example

Traditional branching code


```
a = b + c
if (a > 0)
    e = f + g
else
    e = f / g
h = i - j
```

BB1		add a, b, c
BB1		bgt a, 0, L1
BB3		div e, f, g
BB3		jump L2
BB2	L1:	add e, f, g
BB4	L2:	sub h, i, j



Predicated code

	BB1	add a, b, c if T
p2 → BB2	BB1	p2 = a > 0 if T
p3 → BB3	BB1	p3 = a <= 0 if T
	BB3	div e, f, g if p3
	BB2	add e, f, g if p2
	BB4	sub h, i, j if T




What about nested if-then-else's?

Nested if-then-else's

```
a = b + c
if (a > 0)
  if (a > 25)
    e = f + g
  else
    e = f * g
else
  e = f / g
h = i - j
```

BB1		add a, b, c
BB1		bgt a, 0, L1
BB3		div e, f, g
BB3		jump L2
BB2	L1:	bgt a, 25, L3
BB6		mpy e, f, g
BB6		jump L2
BB5	L3:	add e, f, g
BB4	L2:	sub h, i, j



Traditional branching code

Nested if-then-else's (2)

```
a = b + c
if (a > 0)
  if (a > 25)
    e = f + g
  else
    e = f * g
else
  e = f / g
h = i - j
```

```
BB1  add a, b, c if T
BB1  p2 = a > 0 if T
BB1  p3 = a <= 0 if T
BB3  div e, f, g if p3
BB2  p5 = a > 25 if p2
BB2  p6 = a <= 25 if p2
BB6  mpy e, f, g if p6
BB5  add e, f, g if p5
BB4  sub h, i, j if T
```

```
BB1
BB2
BB3
BB4
BB5
BB6
```

Predicated code

What do we assume to make this work ??


if p2 is False, both p5 and p6 are False

So, predicate setting instruction should set result to False if guarding predicate is false!!!

Benefits/Costs of Predicated Execution

- Benefits
 - Remove branches (both conditional and unconditional)
 - Remove branch mispredictions
 - Overlap execution of if-then-else statements
 - Branches tend to sequentialize operations
 - Predicates can be computed/used in parallel
- Costs
 - Useless instructions executed
 - Code size (extra operand, can't fit into 32-bits)
 - Possibly longer schedule lengths
- The real story
 - Must be applied selectively or you get worse performance than not using it at all

Benefits/Costs of Predicated Execution (2)



BB1
BB2
BB3
BB4
BB5
BB6
BB7

Benefits:

- No branches, no mispredicts
- Can freely reorder independent operations in the predicated block
- Overlap BB2 with BB5 and BB6

Costs (execute all paths)

- worst case schedule length
- worst case resources required

Class Problem 4

```
if (a > 0) {  
    r = t + s  
    if (b > 0 || c > 0)  
        u = v + 1  
    else if (d > 0)  
        x = y + 1  
    else  
        z = z + 1  
}
```

- a. Draw the CFG
- b. Predicate the code removing all branches