

Loop Scheduling

Compilers for High Performance
Architectures



Loop Scheduling

- Dependence Graph

```
do i=0,N
```

```
  y(i)=y(i)+a*x(i)
```

```
endo
```

```
do i=0,N
```

```
  A: t1 = x(i)    ; load
```

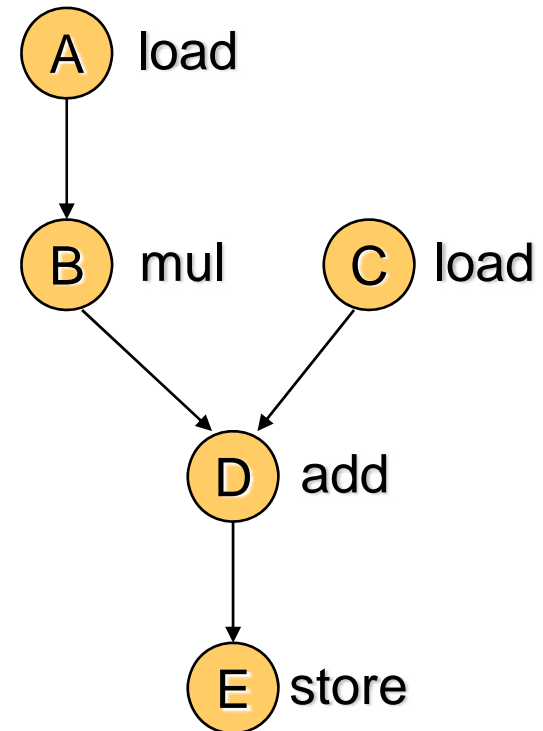
```
  B: t2 = a*t1    ; mul
```

```
  C: t3 = y(i)    ; load
```

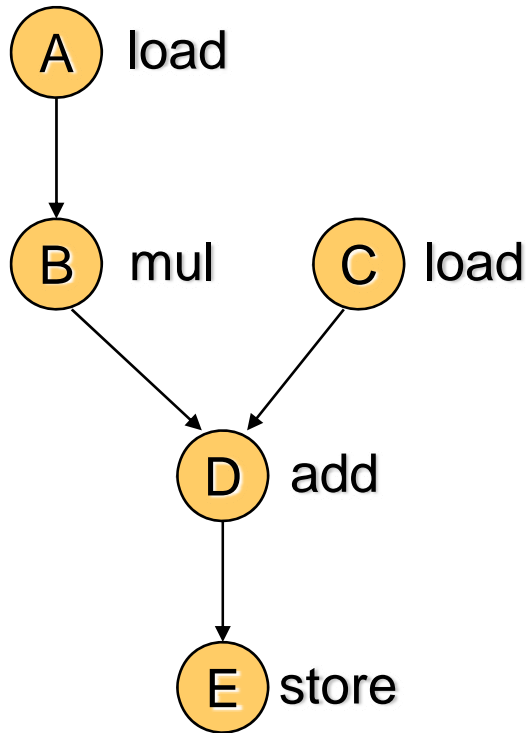
```
  D: t4 = t2 + t3 ; add
```

```
  E: y(i) = t4    ; store
```

```
enddo
```



Loop Scheduling



5 pipelined units
2 load units 2 cycles
1 add unit 2 cycles
1 mul unit 2 cycles
1 store unit 1 cycle

Cycle	Load 1	Load 2	add	mul	store
0	A	C			
1					
2			B		
3					
4				D	
5					
6					E

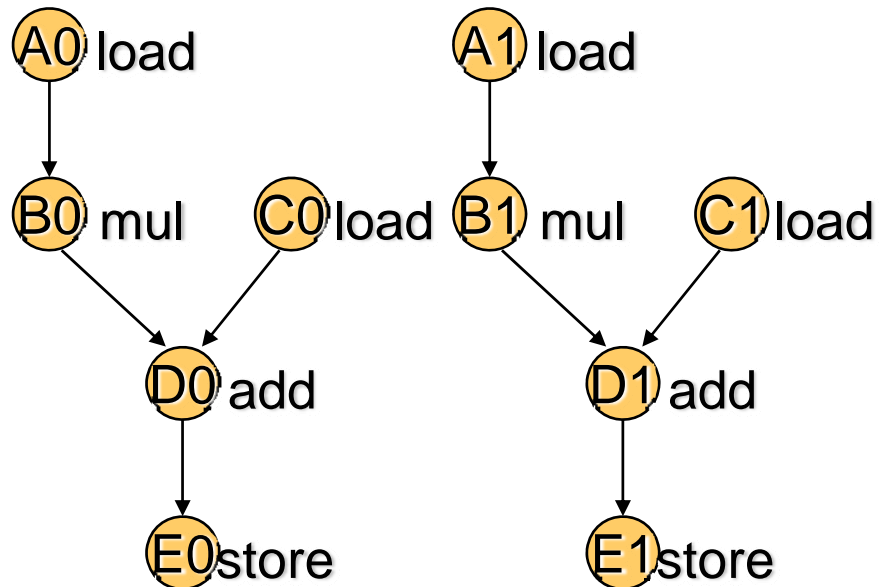
7 cycles per iteration

Maximum ILP = 5 Inst/cycle

Achieved ILP = 0.71 Inst/cycle

Schedule the unrolled loop

```
do i=0,N
  y(i+0)=y(i+0)+a*x(i+0)
  y(i+1)=y(i+1)+a*x(i+1)
enddo
```



Cycle	Load 1	Load 2	add	mul	store
0	A0	C0			
1	A1	C1			
2			B0		
3			B1		
4				D0	
5				D1	
6					E0
7					E1

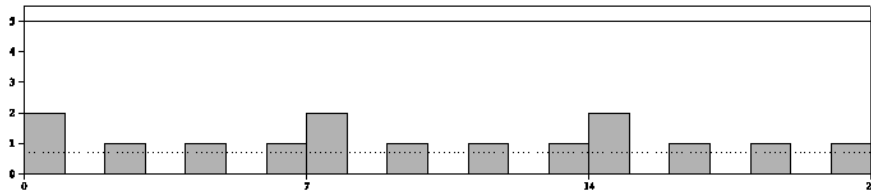
8 cycles per 2 iterations

Maximum ILP = 5 Inst/cycle

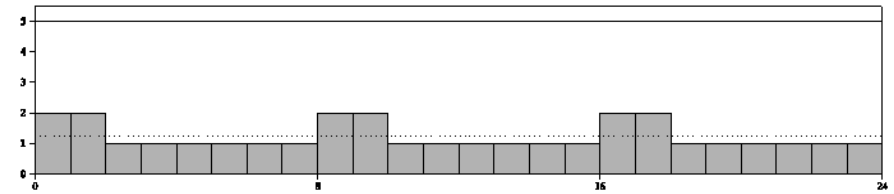
Achieved ILP = 1.25 Inst/cycle

ILP versus unroll degree

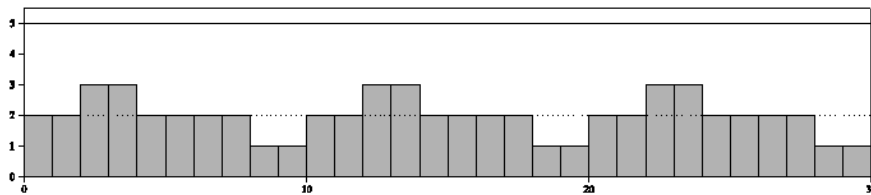
No unroll



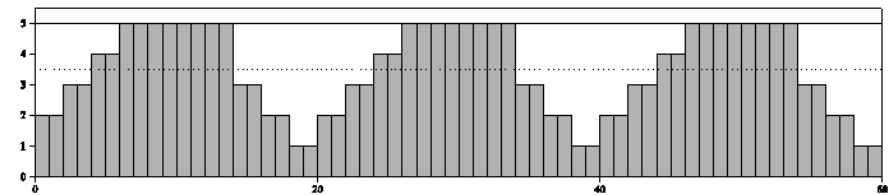
Unroll = 2



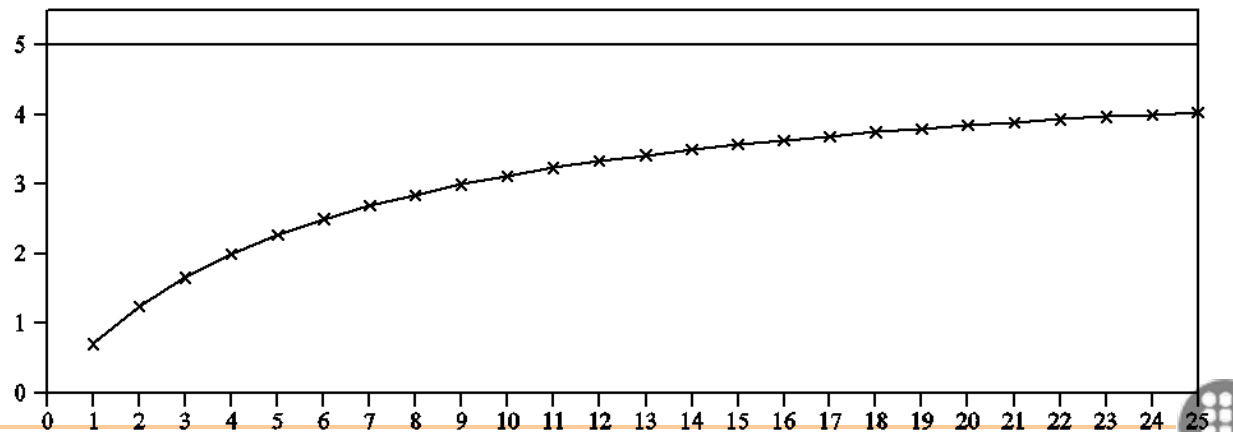
Unroll = 4



Unroll = 14



ILP versus
Unroll degree



ILP vs unroll degree

- To achieve maximum parallelism
- Fully unroll the loop
- We have found a repeating pattern !!!

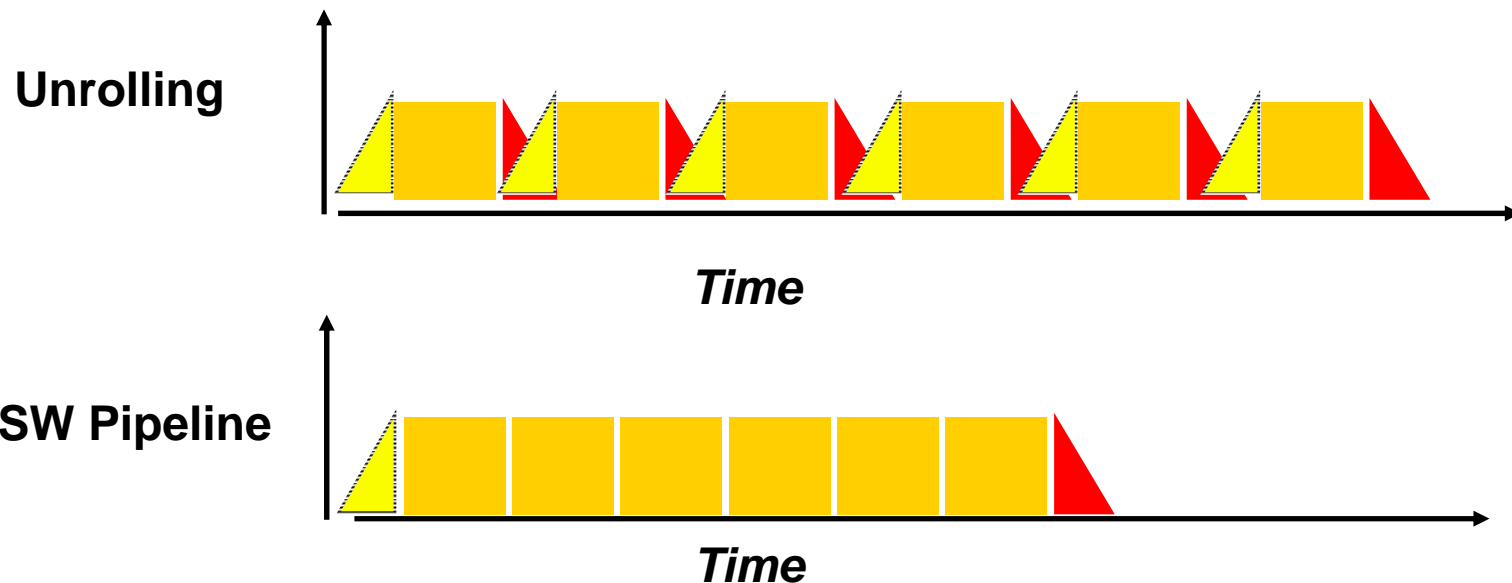
Cycle	Load 1	Load 2	add	mul	store
0	A0	C0			
1	A1	C1			
2	A2	C2	B0		
3	A3	C3	B1		
4	A4	C4	B2	D0	
5	A5	C5	B3	D1	
6	A6	C6	B4	D2	E0
7	A7	C7	B5	D3	E1

i-1	Ai-1	Ci-1	Bi-3	Di-5	Ei-7
i	Ai	Ci	Bi-2	Di-5	Ei-6

n	An	Cn	Bn-2	Dn-5	En-7
n+1			Bn-1	Dn-4	En-6
n+2			Bn	Dn-3	En-5
n+3				Dn-2	En-4
N+4				Dn-1	En-3
N+5				Dn	En-2
N+6					En-1
n+7					En

Software Pipelining

- What is Software Pipelining?
 - A technique for scheduling instructions to exploit ILP in inner loops
 - Many programs spend most of their time in inner loops (e.g. 80% in Perfect Club)
 - We can exploit parallelism between loop iterations
 - Achieves same effect as fully unrolling the loop
- More efficient than unrolling (but also more limited)



Software pipelining

- A repeating pattern can be detected or constructed
 - kernel
- By iterating over the kernel we can achieve the same effect as fully unrolling the loop
- We need code before/after the loop to fill/drain the "pipe":
 - prologue
 - epilogue

Cycle	Load 1	Load 2	add	mul	store
0	A0	C0			
1	A1	C1			
2	A2	C2	B0		
3	A3	C3	B1		
4	A4	C4	B2	D0	
5	A5	C5	B3	D1	

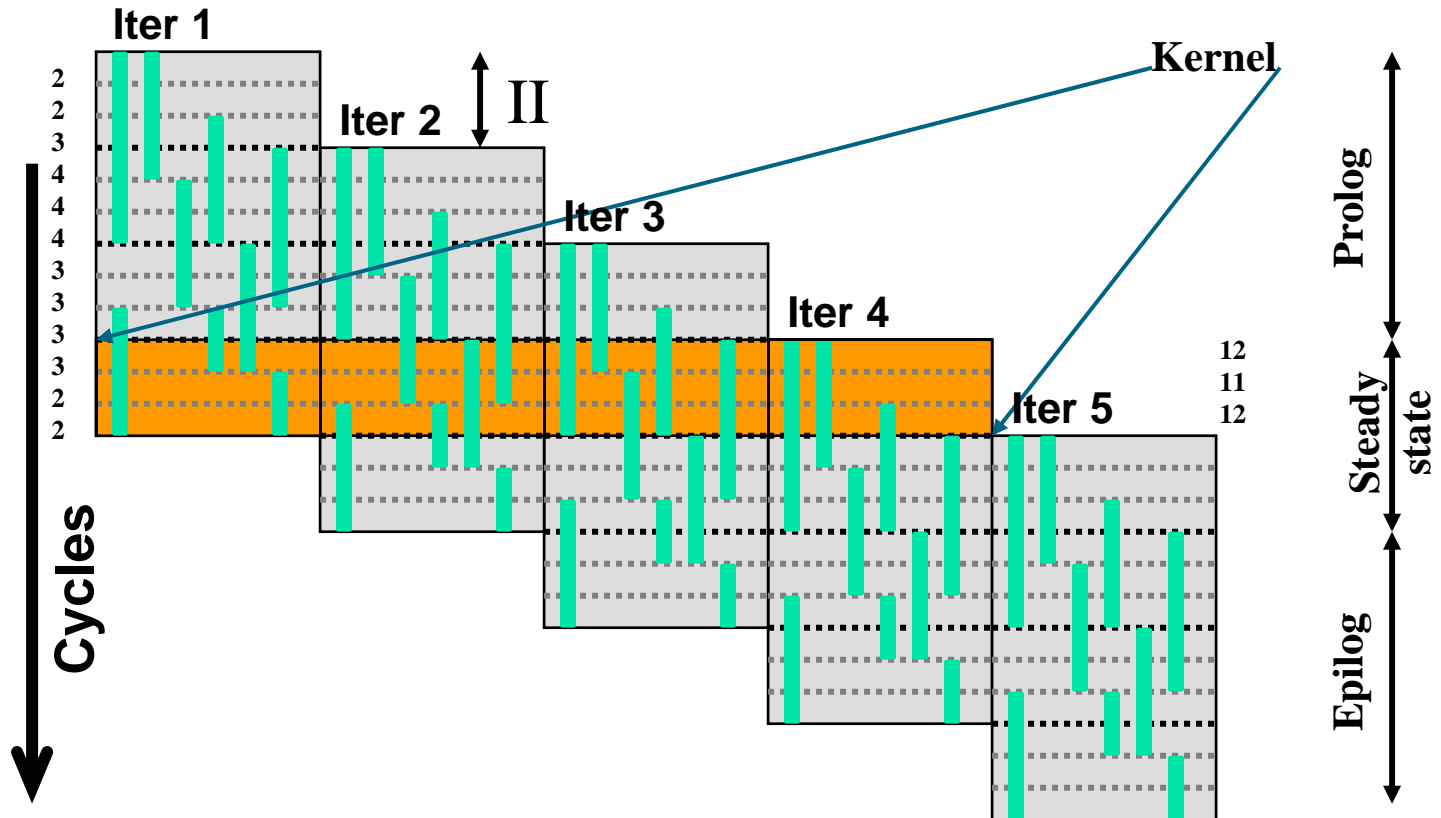
do i=6, N

i	Ai	Ci	Bi-2	Di-5	Ei-6
---	----	----	------	------	------

endo

n+1			Bn-1	Dn-4	En-6
n+2			Bn	Dn-3	En-5
n+3				Dn-2	En-4
N+4				Dn-1	En-3
N+5				Dn	En-2
N+6					En-1
n+7					En

Software pipelining



Software pipelining

- The schedule of an iteration can be divided in stages
- Every II (Initiation Interval) cycles a new iteration is started
- Every II cycles active iterations transition to the next stage
- Several iterations overlap their execution in different stages
- Very similar to the behavior of hardware pipelines hence the term “Software Pipelining”

Software Pipelining vs. Unrolling

- Sw pipelining optimizes the loop throughput
 - Assumption: latency of an iteration is not important
 - Unrolling optimizes the latency of n iterations
- Sw pipelining limits code expansion to prologue / epilogue
 - Even Prologue/Epilogue could be removed with additional hardware support (rotating registers and predicates)
- Sw pipelining is practically limited to loops without control flow
 - Control flow in the loop must be predicated out
 - This may be very inefficient for unbalanced paths
- Efficient sw pipelining requires hardware support
 - Rotating registers, predicates and special branches

Software pipelining

- Main Software pipelining objective:
 - Find an instruction schedule for the kernel with the smallest initiation interval so that throughput is maximized
- Intuitive approaches: Keep unrolling and scheduling the loop until a repeating pattern is found
 - Detecting patterns is difficult
 - Results may be suboptimal
 - Convergence to a valid schedule is not guaranteed
- Direct approach: Build directly the schedule
 - NP complexity
 - Optimal approaches
 - Heuristic approaches
 - Modulo scheduling: best known and practical approach

Modulo Scheduling

Compilers for High Performance
Architectures



Class outline

- Modulo scheduling
- Added difficulties for modulo scheduling
- Iterative modulo scheduling
- Code generation schemas
 - Code size
- Hardware support for modulo scheduling
 - Rotating register files
 - Rotating predicates
 - Code Schemas with hardware support
- Register requirements of software pipelined loops

Modulo Scheduling

- A software pipelining framework
 - Produce a schedule for one iteration so that no data dependence and no resource conflict arises when the same schedule is repeated at regular intervals
 - The constant interval between the start of successive iterations is called initiation interval (II)
- Modulo scheduling steps
 - Choose a candidate initial II (MII or minimum initiation interval)
 - MII is a lower bound of the II
 - Starting from MII, try to find a legal schedule
 - Increases the II until we can find a legal schedule

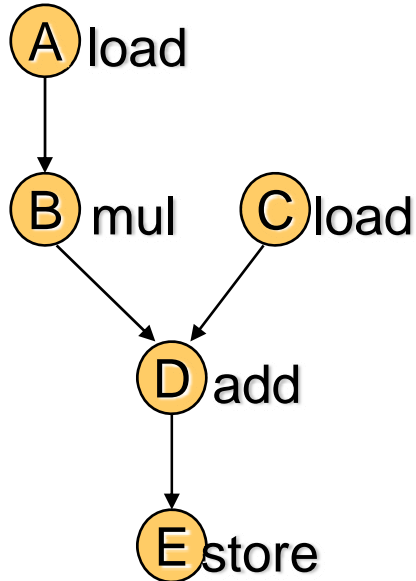
Modulo Scheduling Notes

- Scheduling a loop body within a MS framework is very different than scheduling acyclic regions
 - MS is not concerned with the schedule length
 - The II is what determines the loop performance: as long as a schedule satisfies II requirements, length is irrelevant
- To enforce constraints, MS uses a “modulo reservation table”
 - It checks for resource conflicts in the current cycle and in all schedule cycles that differ by II from the current cycle
- Modulo scheduling is near-optimal in 96-99% of simple loops

Initiation Interval

- Choosing a sensible MII improves the speed of the modulo scheduler
 - We can compute a good MII as the maximum of
 - A resource-constrained II (ResMII, derived from a mapping from the loop body to the target resources)
 - A recurrence-constrained II (RecMII, derived from the inter-iteration data dependence graph of the loop)

Modulo Scheduling example

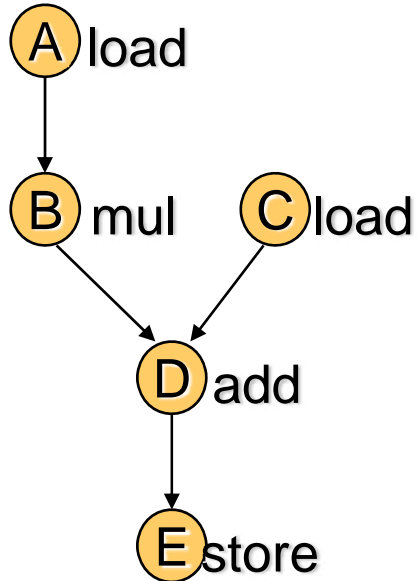


- 2 pipelined units
 - 1 MEM unit
 - load 2 cycles
 - store 1 cycle
 - 1 ALU unit
 - mul 3 cycles
 - add 3 cycles

- Computing the MII
 - Resources limit the II in this example
 - MEM = 3 uses / 1 resource
 - ALU = 2 uses / 1 resource
 - MII = 3

$$resMII = \max_i \left(\left\lceil \frac{resUsage_i}{resCount_i} \right\rceil \right) = \max \left(\left\lceil \frac{2}{1}, \frac{3}{1} \right\rceil \right) = 3$$

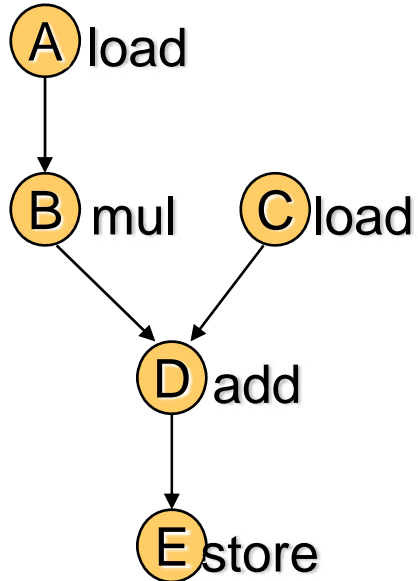
Modulo Scheduling example



Stage	Cycle	MEM	ALU
0	0	A	
	1		
	2		
1	3		
	4		
	5		
2	6		
	7		
	8		
3	9		
	10		
	11		

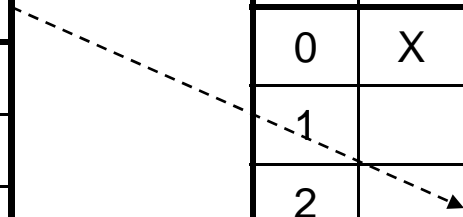
MRT		
Cycle	MEM	ALU
0	X	
1		
2		

Modulo Scheduling example

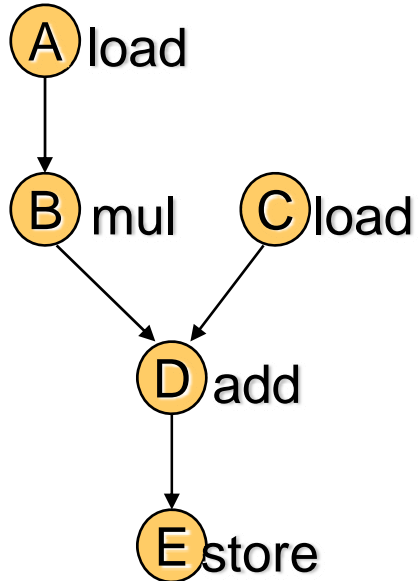


Stage	Cycle	MEM	ALU
0	0	A	
	1		
	2		B
1	3		
	4		
	5		
2	6		
	7		
	8		
3	9		
	10		
	11		

MRT		
Cycle	MEM	ALU
0	X	
1		
2		X



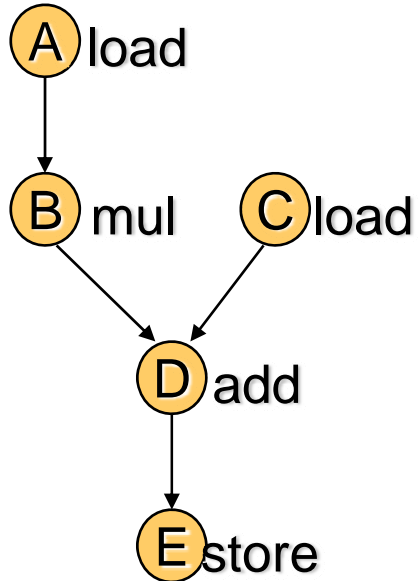
Modulo Scheduling example



Stage	Cycle	MEM	ALU
0	0	A	
	1	C	
	2		B
1	3		
	4		
	5		
2	6		
	7		
	8		
3	9		
	10		
	11		

MRT		
Cycle	MEM	ALU
0	X	
1	X	
2		X

Modulo Scheduling example

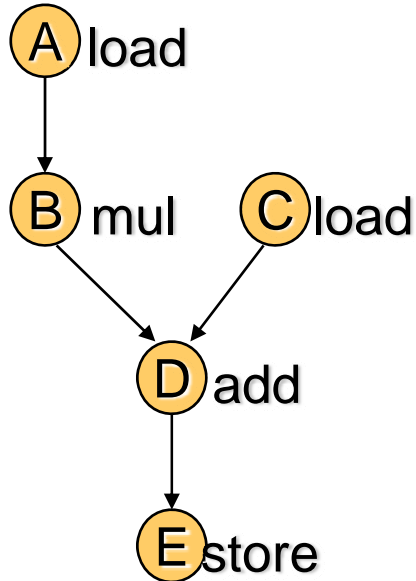


Stage	Cycle	MEM	ALU
0	0	A	
	1	C	
	2		B
1	3		
	4		
	5		D
2	6		
	7		
	8		
3	9		
	10		
	11		

MRT		
Cycle	MEM	ALU
0	X	
1	X	
2		X

Conflict!

Modulo Scheduling example

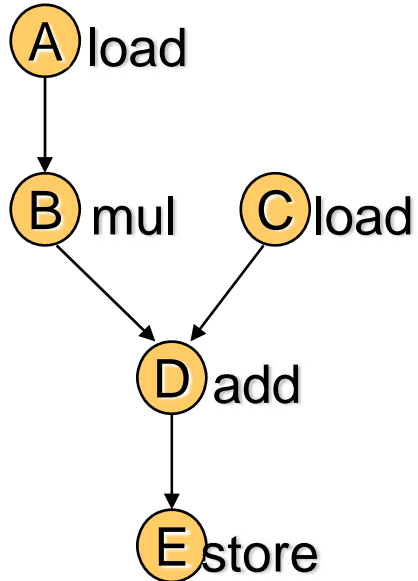


Stage	Cycle	MEM	ALU
0	0	A	
	1	C	
	2		B
1	3		
	4		
	5		
2	6		D
	7		
	8		
3	9		
	10		
	11		

MRT		
Cycle	MEM	ALU
0	X	X
1	X	
2		X



Modulo Scheduling example

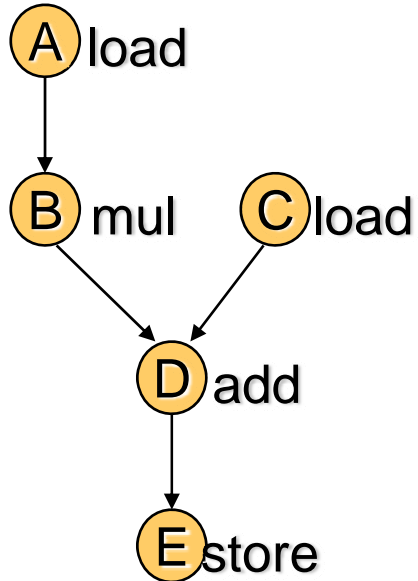


Stage	Cycle	MEM	ALU
0	0	A	
	1	C	
	2		B
1	3		
	4		
	5		
2	6		D
	7		
	8		
3	9	E	
	10		
	11		

MRT		
Cycle	MEM	ALU
0	X	X
1	X	
2		X

Conflict!!!

Modulo Scheduling example

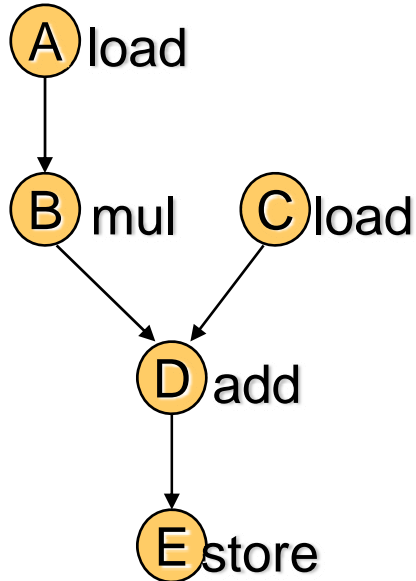


Stage	Cycle	MEM	ALU
0	0	A	
	1	C	
	2		B
1	3		
	4		
	5		
2	6		D
	7		
	8		
3	9		
	10	E	
	11		

MRT		
Cycle	MEM	ALU
0	X	X
1	X	
2		X

Conflict!!!

Modulo Scheduling example



An individual iteration takes more time:

- 12 cycles
- vs 9 cycles

But N iterations take:

- $9 + 3 \times N$
- vs $9 \times N$

Stage	Cycle	MEM	ALU
0	0	A	
	1	C	
	2		B
1	3		
	4		
	5		
2	6		D
	7		
	8		
3	9		
	10		
	11	E	

MRT		
Cycle	MEM	ALU
0	X	X
1	X	
2	X	X

Recurrence constraints and MII

The maximum computation rate of a loop is bounded by the following ratio

$$r_{\text{opt}} = \min_C \{ D_C / W_C \}$$

where C is a dependence cycle, D_C is the total dependence distance along C and W_C is total execution time of C . i.e.

$$D_C = \sum_{e \in C} d_e$$

(d_e is the dependence distance along the edge e in C)

$$W_C = \sum_{e \in C} w_e$$

(w_e is the edge weight (latency) along the edge e in C)

Computing RecMII

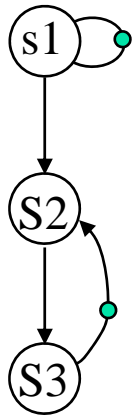
$$RecMII = \max_c \left(\frac{\sum_{e \in C} Lat_e}{\sum_{e \in C} Dist_e} \right)$$

- Computing MIIrec and identifying recurrences:
 - HRMS/SMS
 - require identifying each individual recurrence circuit and sort them by criticality. An **exponential** number of recurrence circuits can appear
 - IMS & slack: MinDist Matrix
 $O(V^3) \sim O(V^4)$
 - LXMS: Marks back-edges during graph construction $O(V^2)$
 - Computes MIIrec and identifies recurrence subgraphs
 $O(V * \log V) \sim O(V^2)$

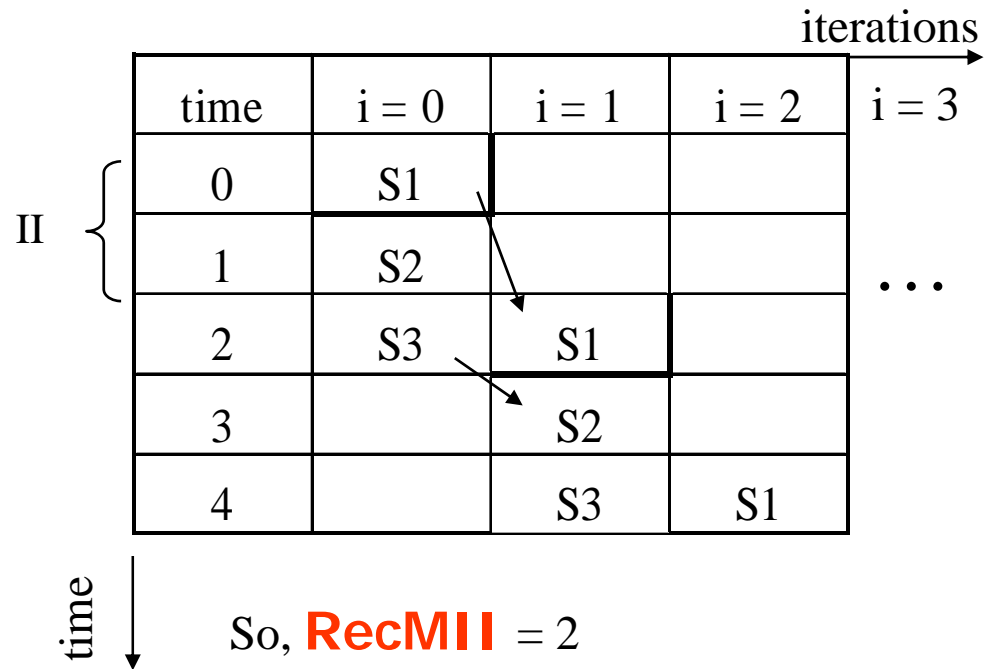
Example of recurrence constraints

```

for i = 0 to N - 1 do
S1:    a[i] = a[i - 1] + R[i];
S2 :   b[i] = a[i] + c[i - 1];
S3 :   c[i] = b[i] + 1;
end;
    
```

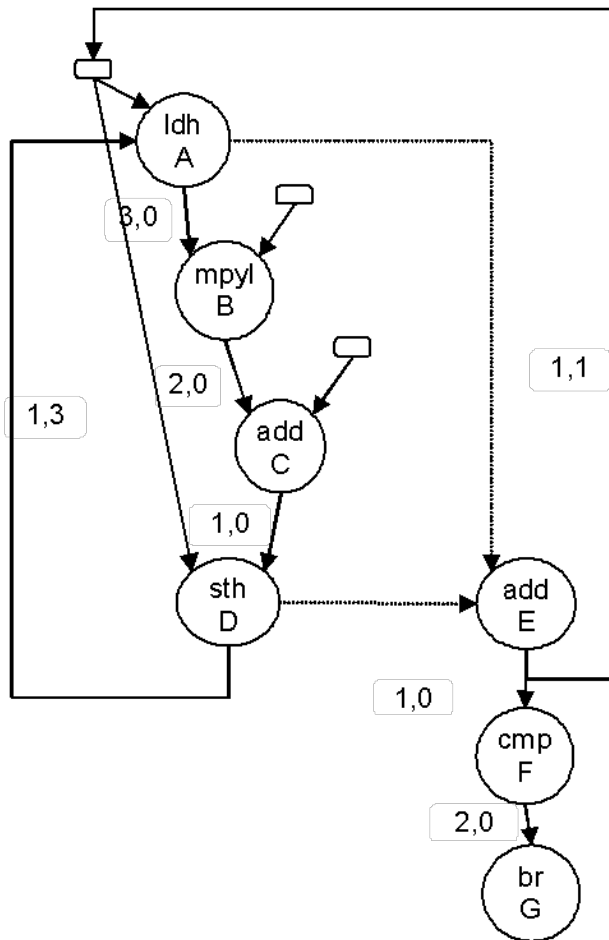


Note: We use a token here to represent a flow dependence of distance = 1



Assume each operation takes 1 cycle!

Class Problem: compute MII



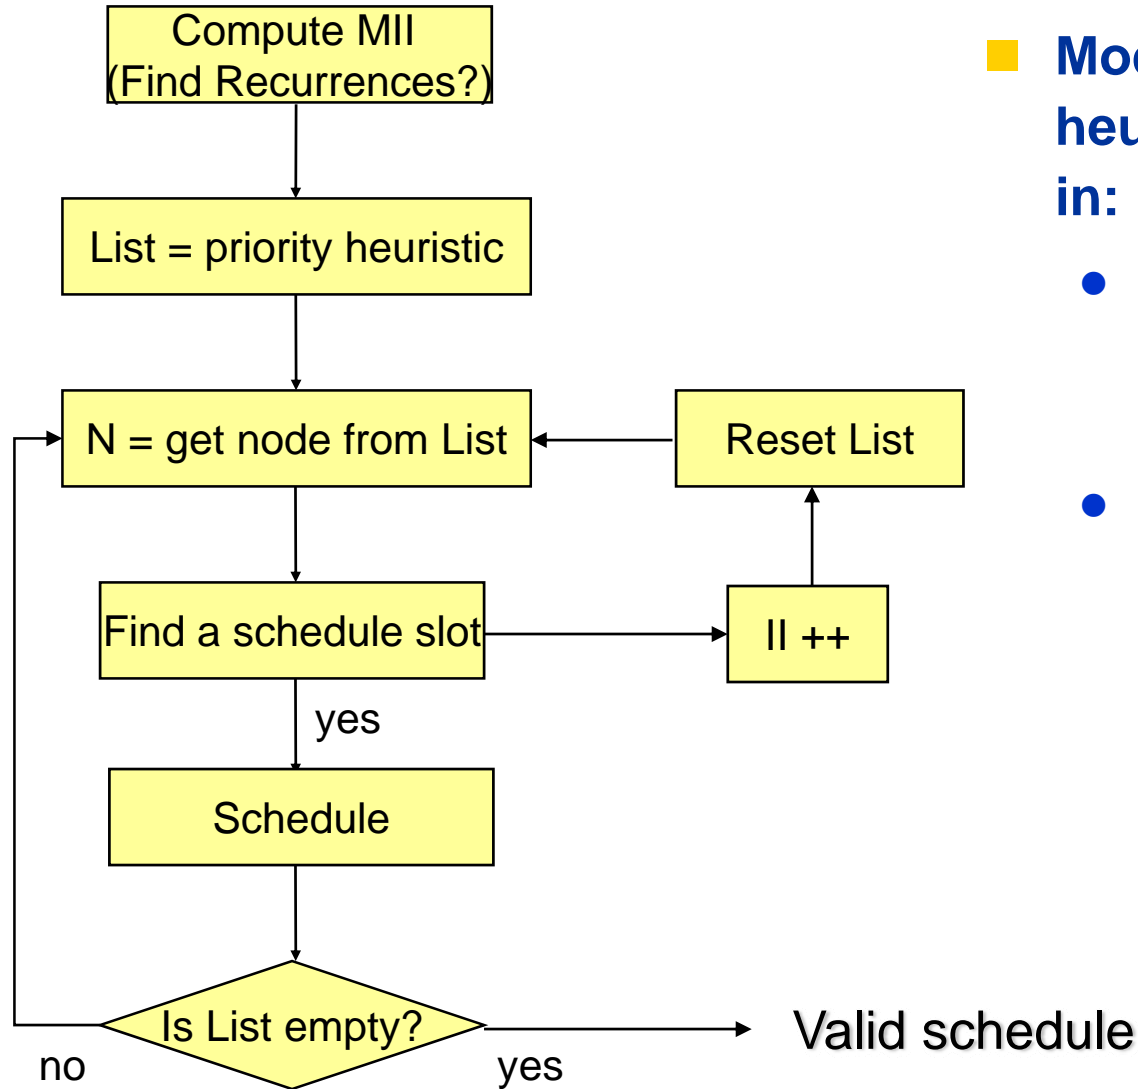
■ Arcs represent

- latency, distance

■ Functional Units

- 1 MEM unit
 - ✓ ldh
 - ✓ sth
- 4 ALU units
 - ✓ add
 - ✓ cmp
 - ✓ br
- 2 MULT units
 - ✓ mpy

Basic modulo scheduling algorithm



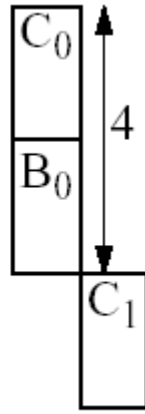
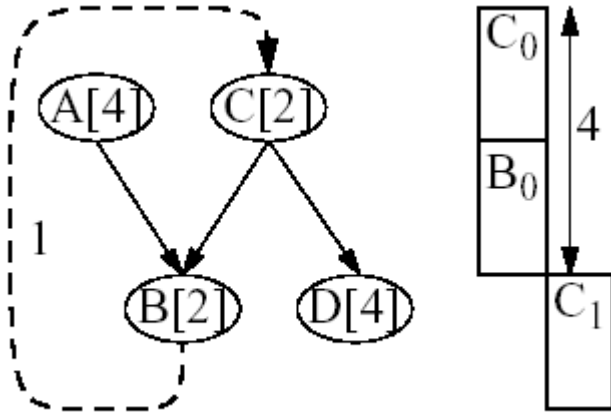
■ Modulo scheduling heuristics mostly differ in:

- Priority heuristic
 - ✓ E.g. adapted list scheduling
- Finding a schedule slot

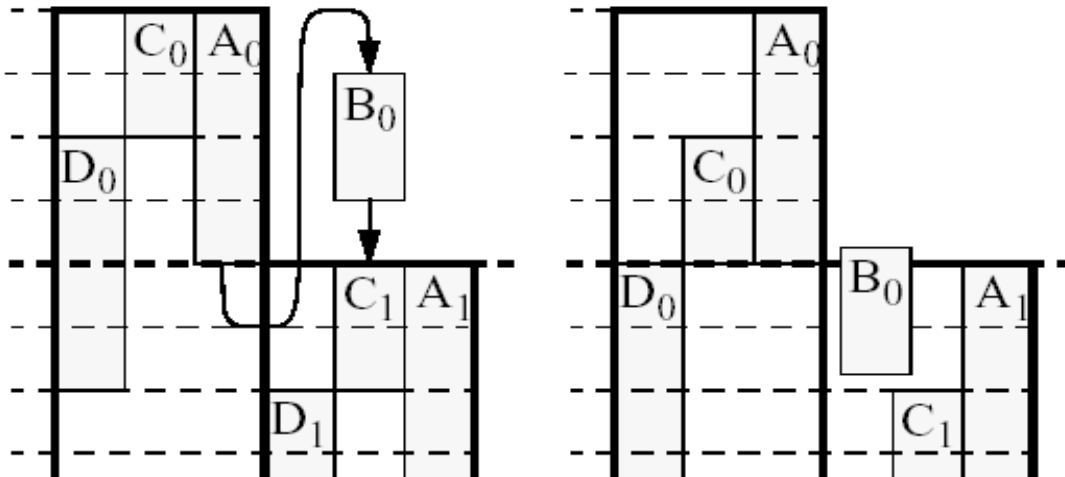
Additional difficulties for MS

- Recurrence constraints difficult the schedule
- Operations with complex reservation tables
 - Non-pipelined functional units (block reservation tables)
 - Operations that use multiple resources or units
- Those constraints combined may make it impossible to achieve the MII
 - Thus MII is only a lower bound

Recurrence constraints

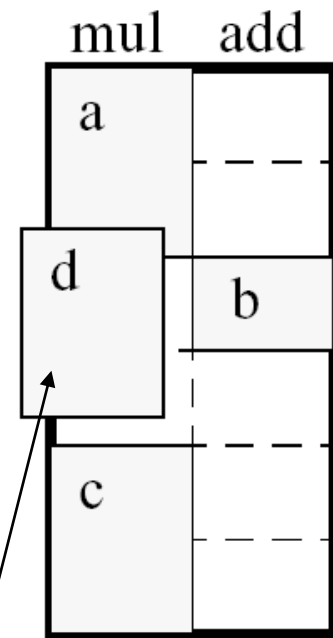
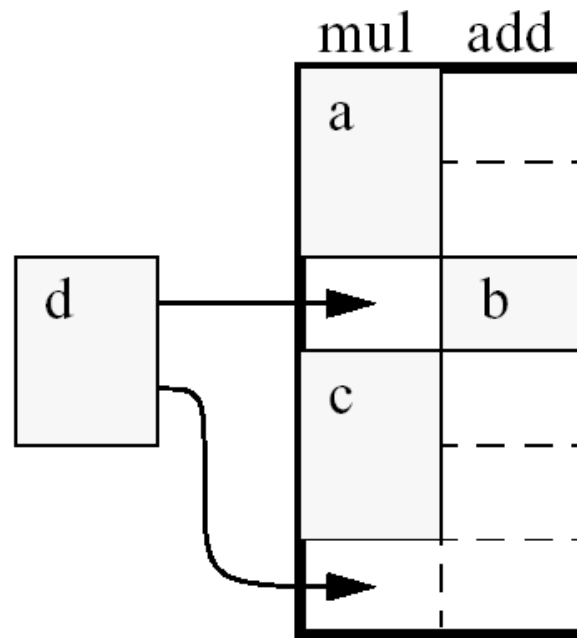
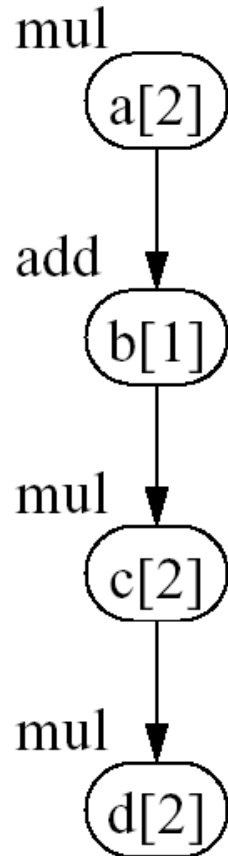


- Assume operations fully pipelined and 3 functional units
 - MII = 4 (recurrences)
- Recurrences difficult the scheduling process
 - A greedy list scheduling might not work
 - Scheduling order {A,C,D,B}
 - No space for B
 - Delaying C by 2 cycles makes space for B



Block reservation tables

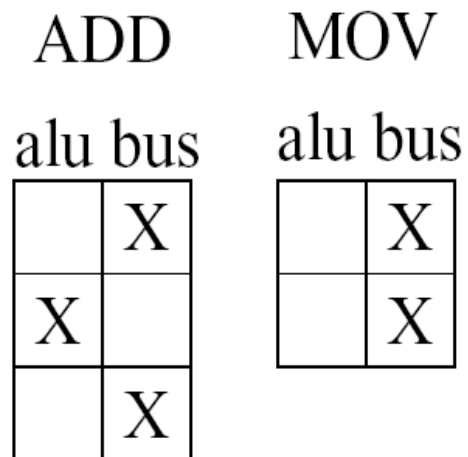
MII = 6



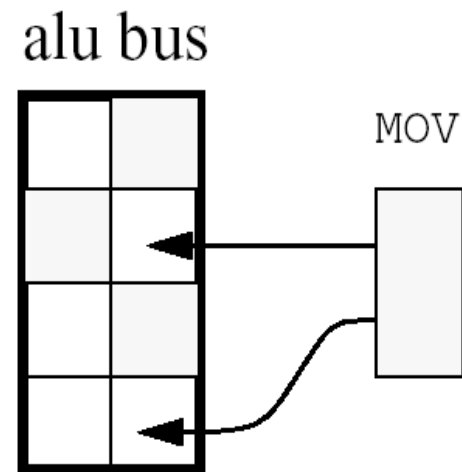
Use 2 slots in next stage

Complex reservation tables

It is impossible to achieve $II = MII = 4$
Optimum $II = 5$

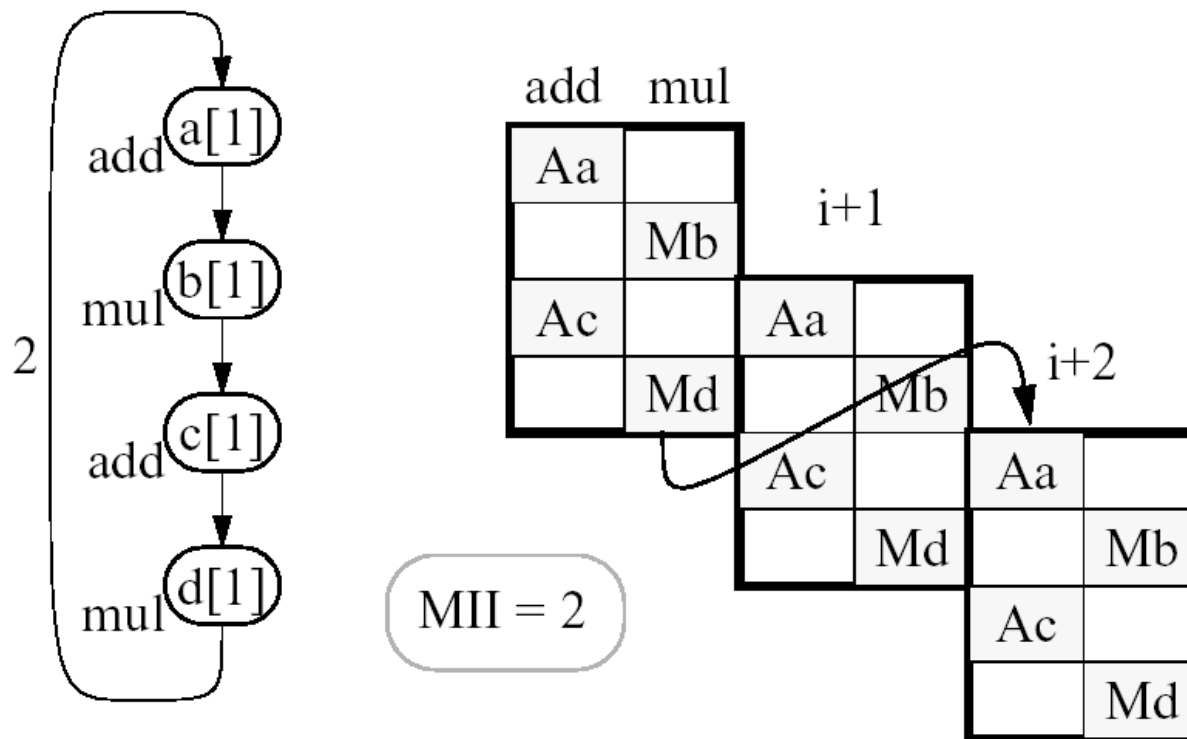


$MI = 4$



Resource + recurrence constraints

- Scheduling all OPs with minimum delay, resource constraints are violated
- Delaying some OP violates the recurrence



Iterative Modulo Scheduling

- Tries to solve all constraints by using a limited form of backtracking
 - If an operation does not fit in the schedule:
 - Choose a schedule slot
 - Un-schedule all conflicting operations -> go back to pending List
 - And Schedule
- A popular misconception is that IMS is slow
 - Operations are scheduled & unscheduled more times
 - However, in practice scheduling time is reduced:
 - Many loops do not require un-scheduling
 - For most loops that do, a few un-scheduled operations may save from rescheduling the whole loop with increased II
 - Faster and better performance
 - Of course some loops will consume its un-scheduling budget and end up increasing the II anyway
 - the amount of un-scheduling must be limited to avoid excessive compilation time

Priority Function

Height-based priority works well for acyclic scheduling, makes sense that it will work for loops as well

Acyclic:

$$\text{Height}(X) = \begin{cases} 0, & \text{if } X \text{ has no successors} \\ \text{MAX}_{\text{for all } Y = \text{succ}(X)} ((\text{Height}(Y) + \text{Delay}(X, Y)), & \text{otherwise} \end{cases}$$

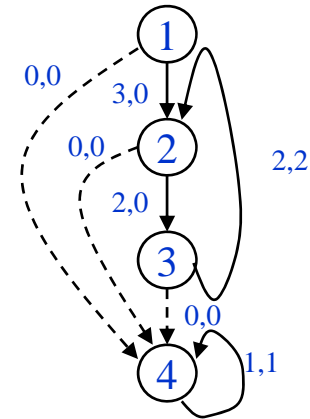
Cyclic:

$$\text{HeightR}(X) = \begin{cases} 0, & \text{if } X \text{ has no successors} \\ \text{MAX}_{\text{for all } Y = \text{succ}(X)} ((\text{HeightR}(Y) + \text{EffDelay}(X, Y)), & \text{otherwise} \end{cases}$$

$$\text{EffDelay}(X, Y) = \text{Delay}(X, Y) - \Pi * \text{Distance}(X, Y)$$

Calculating Height

1. Insert pseudo edges from all nodes to branch with latency = 0, distance = 0 (dotted edges)
2. Compute Π , For this example assume $\Pi = 2$
3. $\text{HeightR}(4) =$
4. $\text{HeightR}(3) =$
5. $\text{HeightR}(2) =$
6. $\text{HeightR}(1) =$



The Scheduling Window

With cyclic scheduling, not all the predecessors may be scheduled, so a more flexible earliest schedule time is:

$$E(Y) = \underset{\text{for all } X = \text{pred}(Y)}{\text{MAX}} \begin{cases} 0, & \text{if } X \text{ is not scheduled} \\ \text{MAX}(0, \text{SchedTime}(X) + \text{EffDelay}(X, Y)), & \text{otherwise} \end{cases}$$

where $\text{EffDelay}(X, Y) = \text{Delay}(X, Y) - \Pi * \text{Distance}(X, Y)$

Every Π cycles a new loop iteration will be initialized, thus every Π cycles the pattern will repeat. Thus, you only have to look in a window of size Π , if the operation cannot be scheduled there, then it cannot be scheduled.

$$\text{Latest schedule time}(Y) = L(Y) = E(Y) + \Pi - 1$$

Modulo Scheduling - Driver

```
compute MII
II = MII
budget = BUDGET_RATIO * number of ops
while (schedule is not found) do
    iterative_schedule(II, budget)
    II++
```

- Budget_ratio is a measure of the amount of backtracking that can be performed before giving up and trying a higher II
 - Empirical results show that ideal values are around 1.5-2
 - Ideal varies with architecture and benchmarks

Modulo Scheduling – Iterative Scheduler

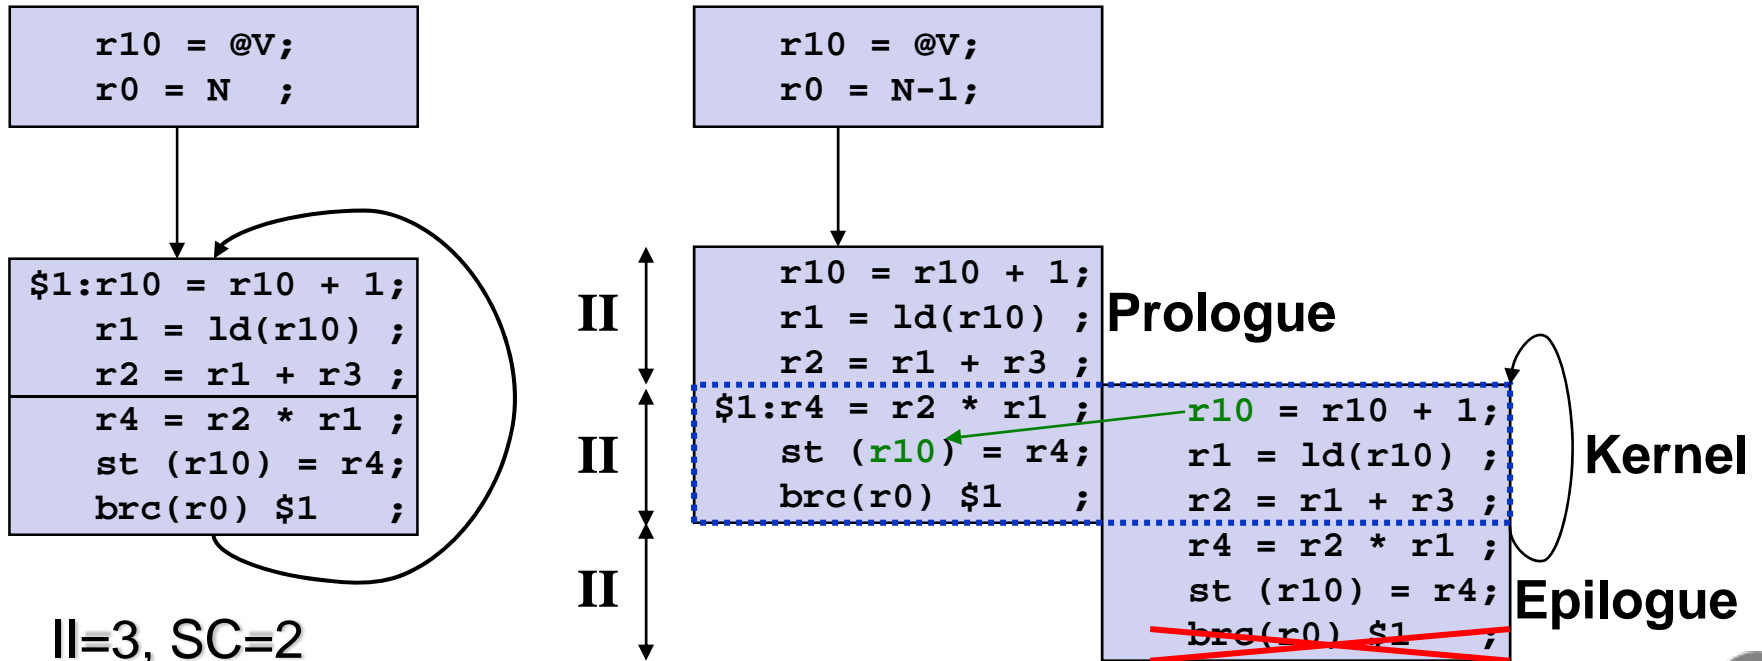
```
iterative_schedule(II, budget)
    compute op priorities
    while (there are unscheduled ops and budget > 0) do
        op = unscheduled op with the highest priority
        min = early time for op ( $E(Y)$ )
        max = min + II - 1
        t = find_slot(op, min, max)
        schedule op at time t
        /* Backtracking phase - undo previous
           scheduling decisions */
        Unschedule all previously scheduled ops that
            conflict with op
    budget--
```

Modulo Scheduling – Find_slot

```
find_slot(op, min, max)
/* Successively try each time in the range */
for (t = min to max) do
    if (op has no resource conflicts in MRT at t)
        return t
/* Op cannot be scheduled in its specified range */
/* So schedule this op and displace all conflicting
ops */
if (op has never been scheduled or min > previous
    scheduled time    of op)
    return min
else
    return MIN(1 + prev scheduled time of op, max)
```

Code generation Schemas

- Once a valid schedule is found code must be generated
 - Prologue code
 - Kernel code
 - Epilogue code
- Intuitively easy to generate
 - Just overlap the schedule for SC (Stage Count) iterations
- Example is incorrect as is
 - R10 is rewritten by iteration (i+1) before using value generated in iteration (i)
 - Renaming does not work since the kernel is repeatedly executed

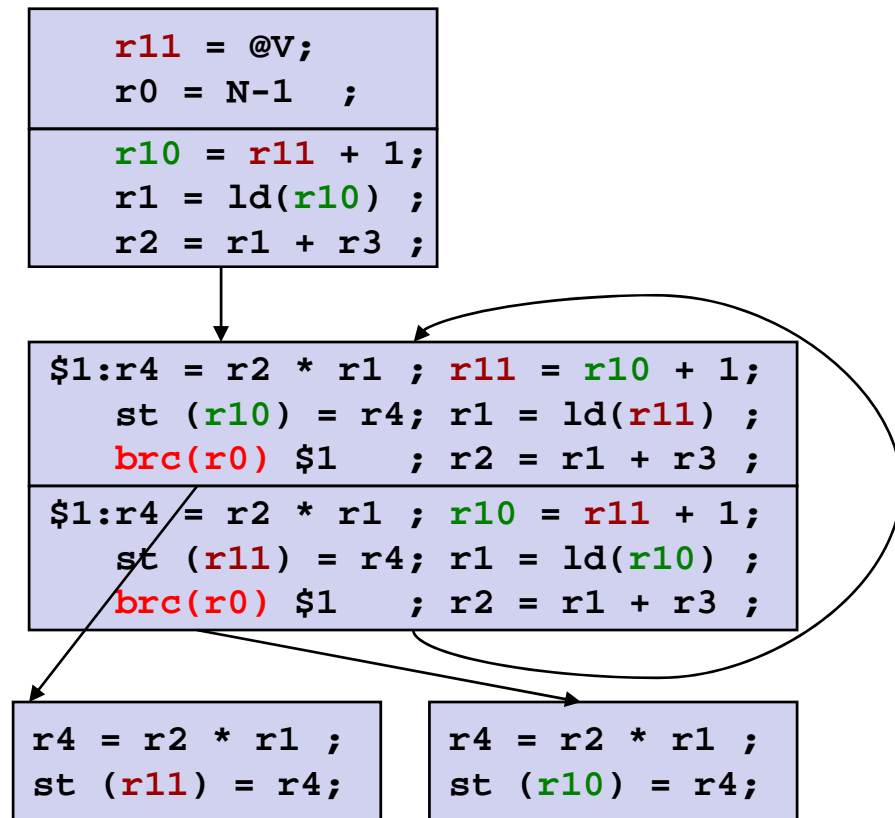


Modulo Variable Expansion

- Unroll the kernel K_{min} times

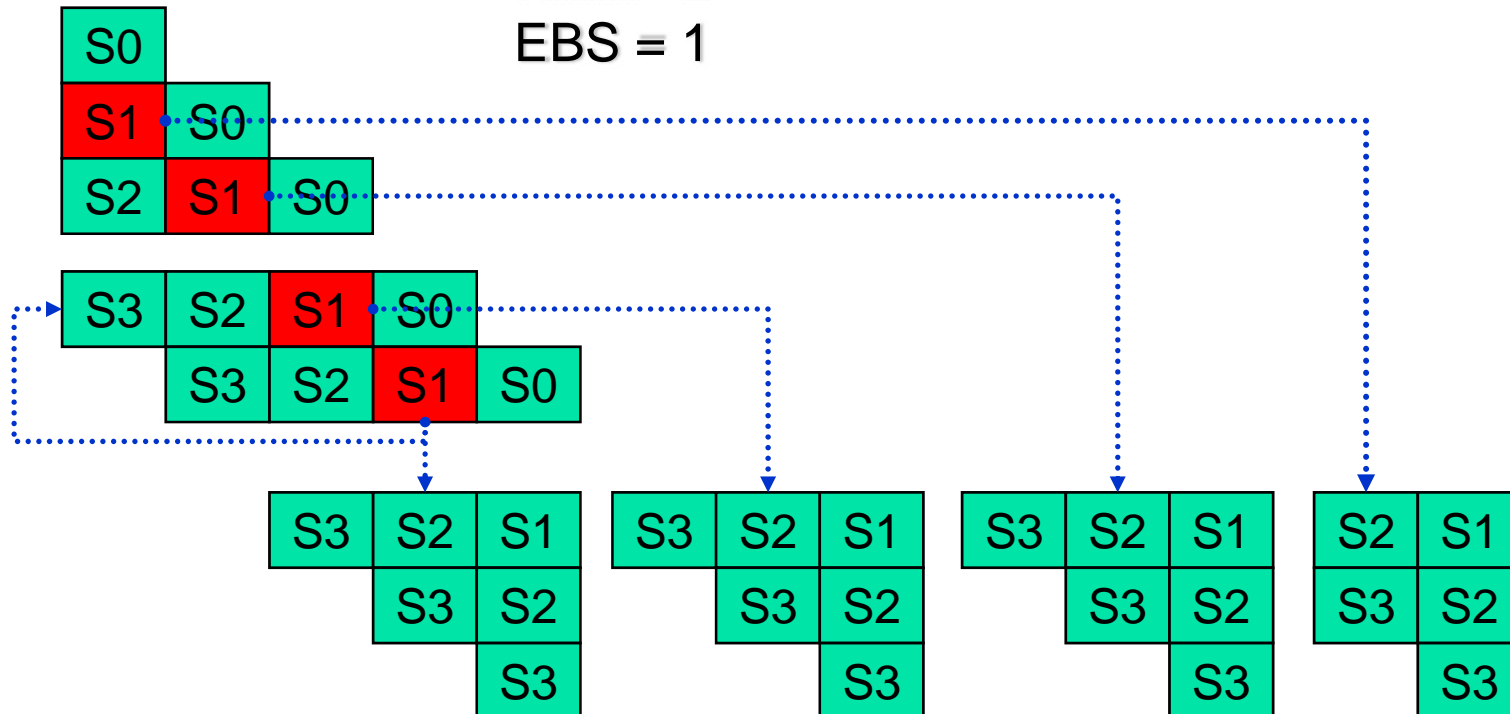
$$K_{min} = \max_i \left(\left\lceil \frac{Lifetime_i}{II} \right\rceil \right)$$

- Rename the overlapping registers
 - So that Def-Use chains correspond to the original iterations of the code
- Multiple versions of the epilogue are required
 - Each one uses different registers
- Noticeable code expansion
 - SC = Stage Count
 - K_{min} = minimum MVE unroll
 - EBS = Exit Branch Stage



Basic code generation schema

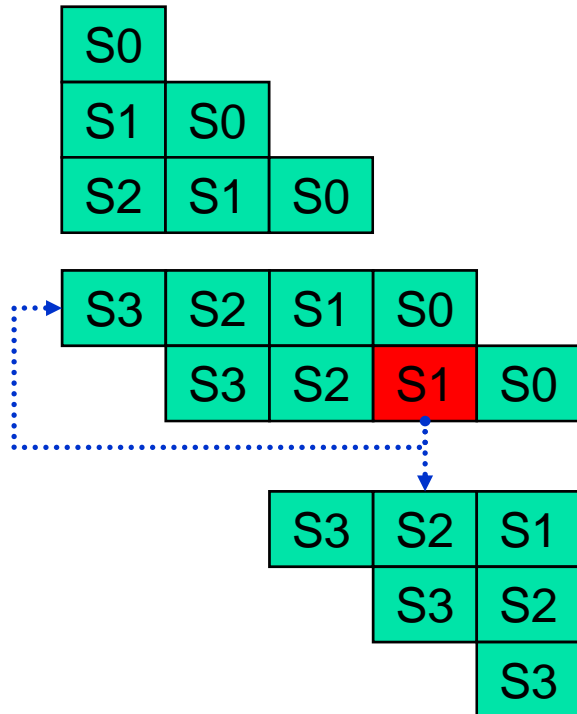
SC = 4
Kmin= 2
EBS = 1



Non Speculative code schema only for counted loops

Exit branch not necessarily scheduled in last stage

Preconditioned code schema



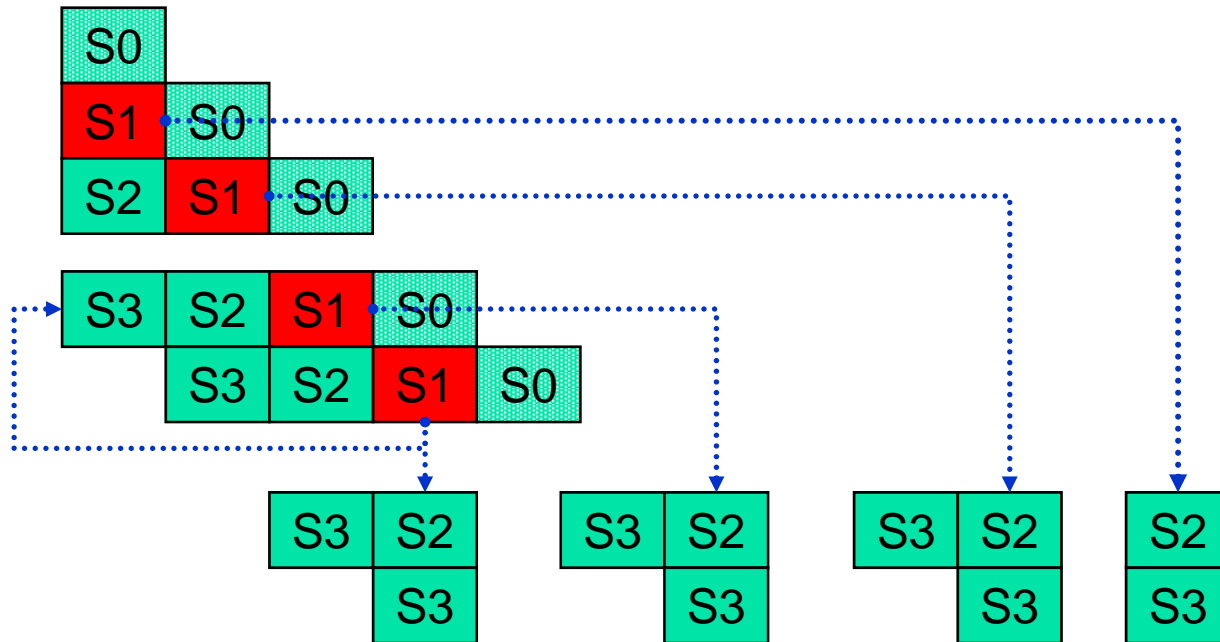
- Requires a non-pipelined preconditioning loop:
 - Execute short trip loops
 - Execute remainder iterations
- Iteration count
 - L = iterations of original loop
 - N = iterations executed by the preconditioning loop
 - M = Iterations executed by modulo scheduled loop

$$N = \begin{cases} L & \text{if } L < SC - 1 \\ \lceil L - (SC - 1) \rceil \bmod Kmin & \text{otherwise} \end{cases}$$

$$M = L - N$$

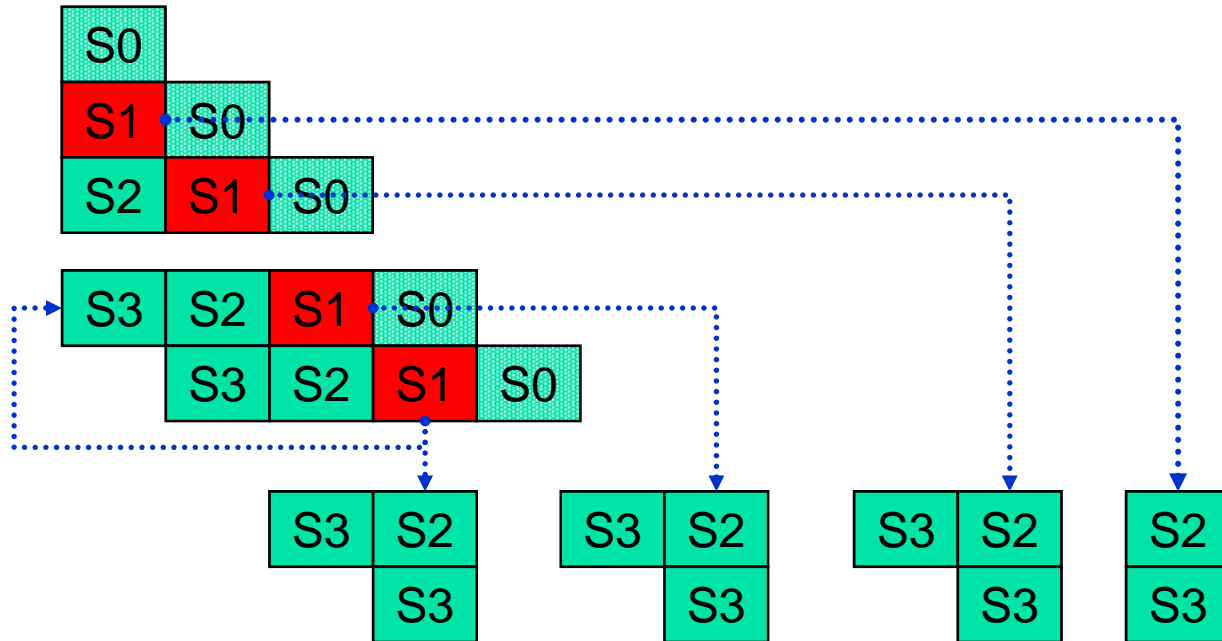
- Disadvantages
 - Adds a whole loop body to the code size.
 - Only for counted loops

Speculative code schema



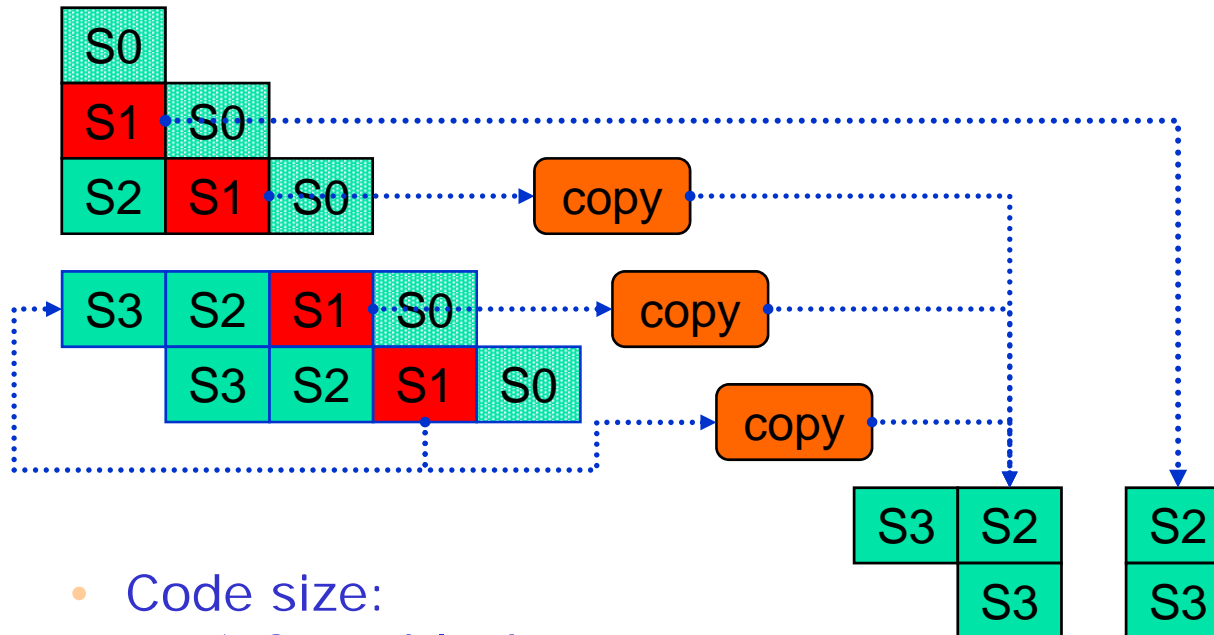
- Counted loops and loops with conditional exit.
 - [Rau et al., Micro25] [Lavery & Hwo, Micro 29]

Collapse kernel epilogs



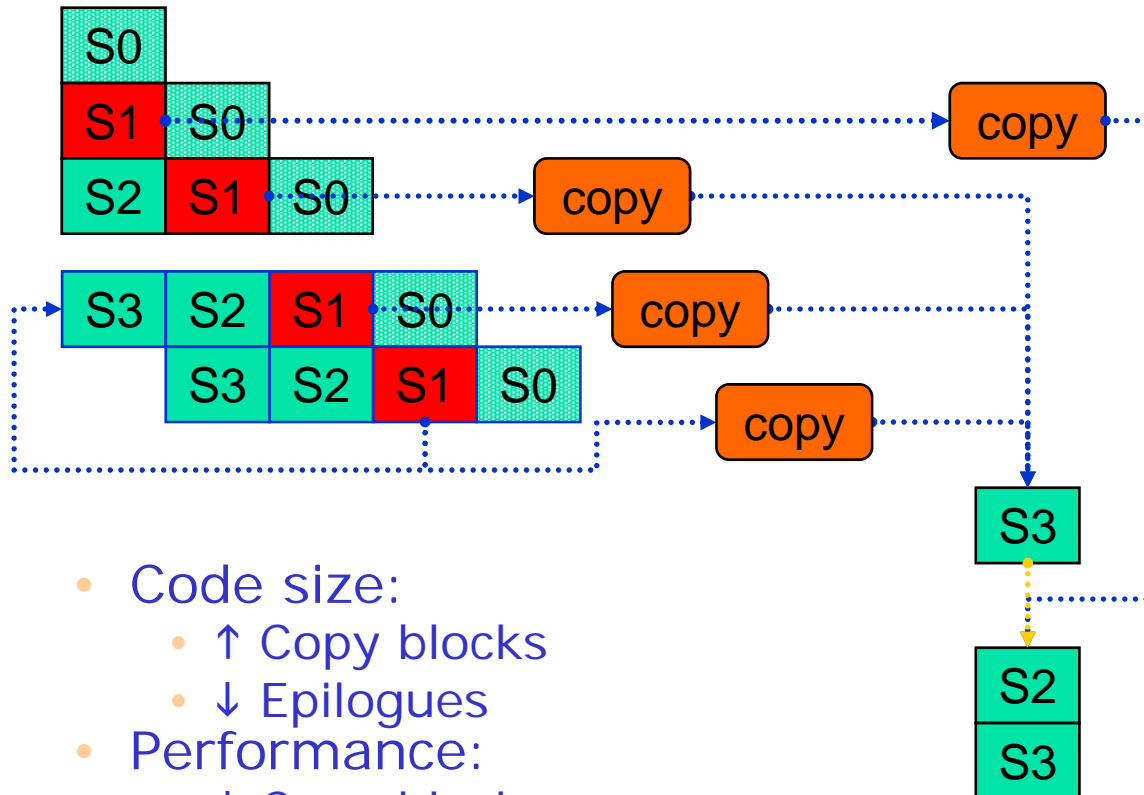
- The three kernel epilogs are identical except the registers they use
- The latter the exit branch -> more speculation -> smaller epilogs
 - Exit branch in last stage -> epilogs are empty

Collapse kernel epilogues



- Code size:
 - ↑ Copy blocks
 - ↓ Epilogues
- Performance:
 - ↓ Copy blocks
 - ↑ Epilog overlaps with next BB

Collapse prologue epilogues



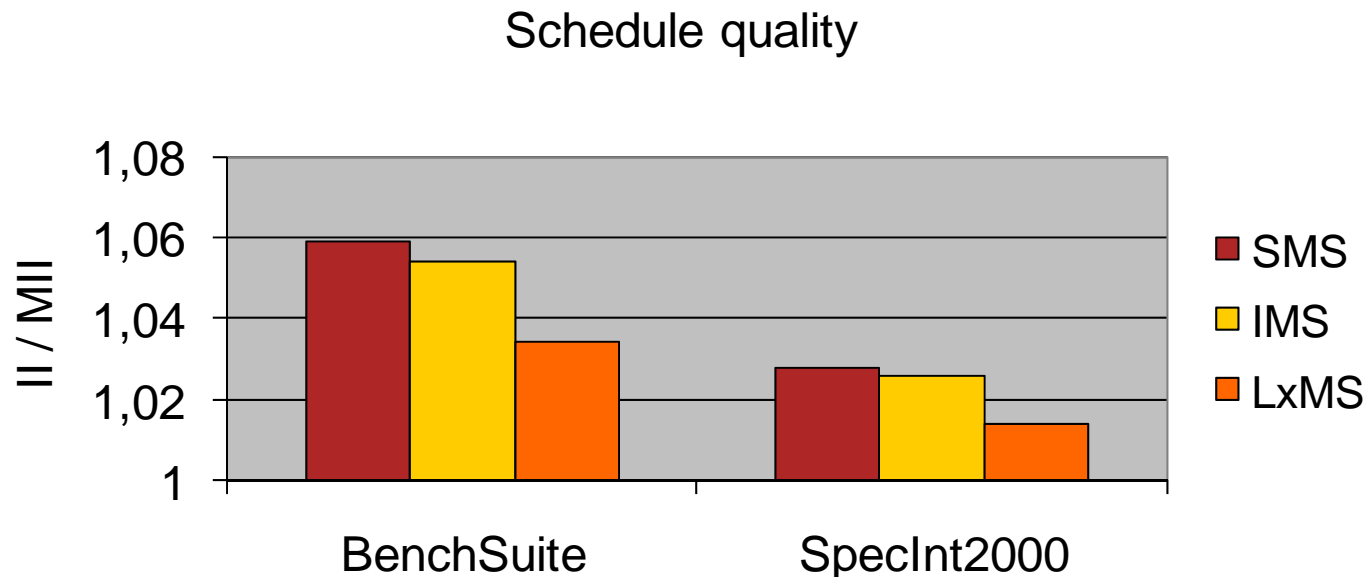
- Code size:
 - ↑ Copy blocks
 - ↓ Epilogues
- Performance:
 - ↓ Copy blocks
 - ↓ Epilog is serialized
 - ↑ Last epilog BB overlaps with next BB

Reduce MVE by copy insertion

- Insert copies in the schedule to split lifetimes. This step is performed after scheduling using the unused slots:
 - Until MVE is reduced do:
 - Select the longest lifetime that can be split by a copy?
 - Insert copy in DDG between the producer and the last user
 - Schedule the copy starting at a midpoint cycle:
 $m, m+1, m-1, m+2, m-2, \dots$
 - Connect all consumers that can read from the copy
 - If MVE is reduced update IR and try to further reduce MVE
 - If any step fails, remove non-necessary copies

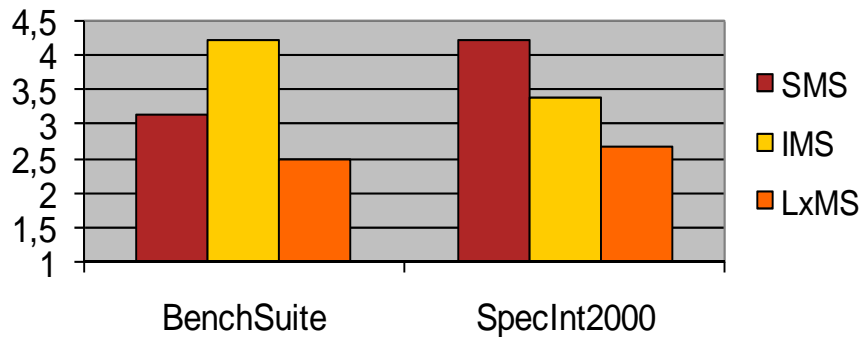
Impact of modulo scheduling heuristics

- SMS – Swing Modulo Scheduling (Llosa et al. PACT96)
 - Used only to show the influence of the modulo scheduler
- IMS – Iterative Modulo Scheduling (Rau Micro27)
 - Used as baseline
- LxMS – LxBE's modulo scheduler (Llosa, Freudenberger & Zalamea) HPLC

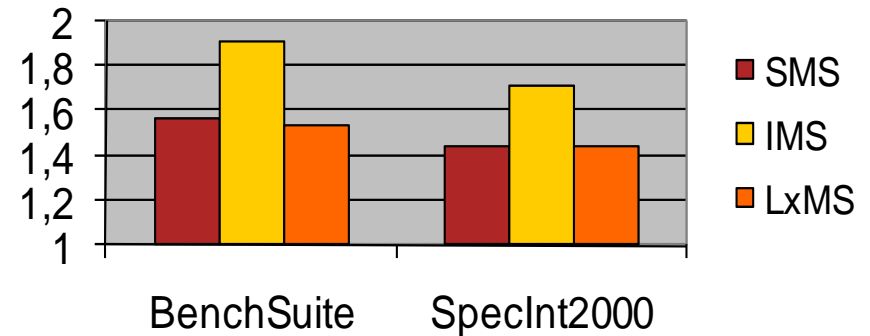


Impact of modulo scheduling heuristics

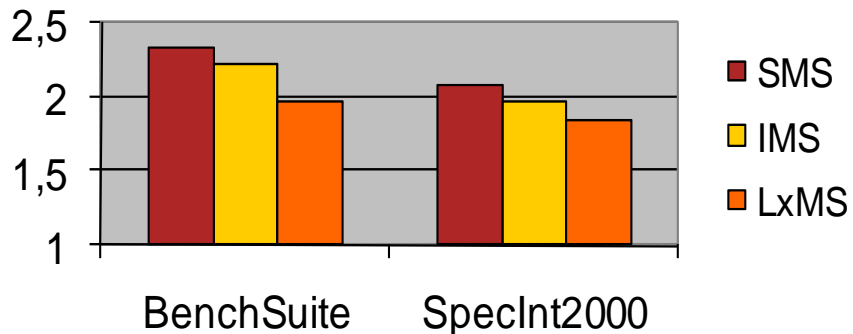
Loop expansion ratio



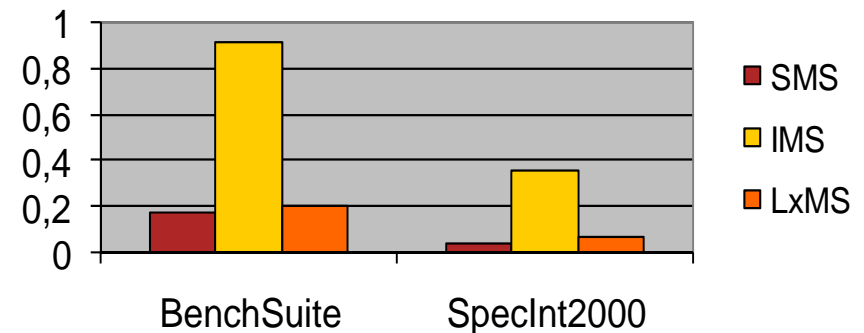
Average Kmin



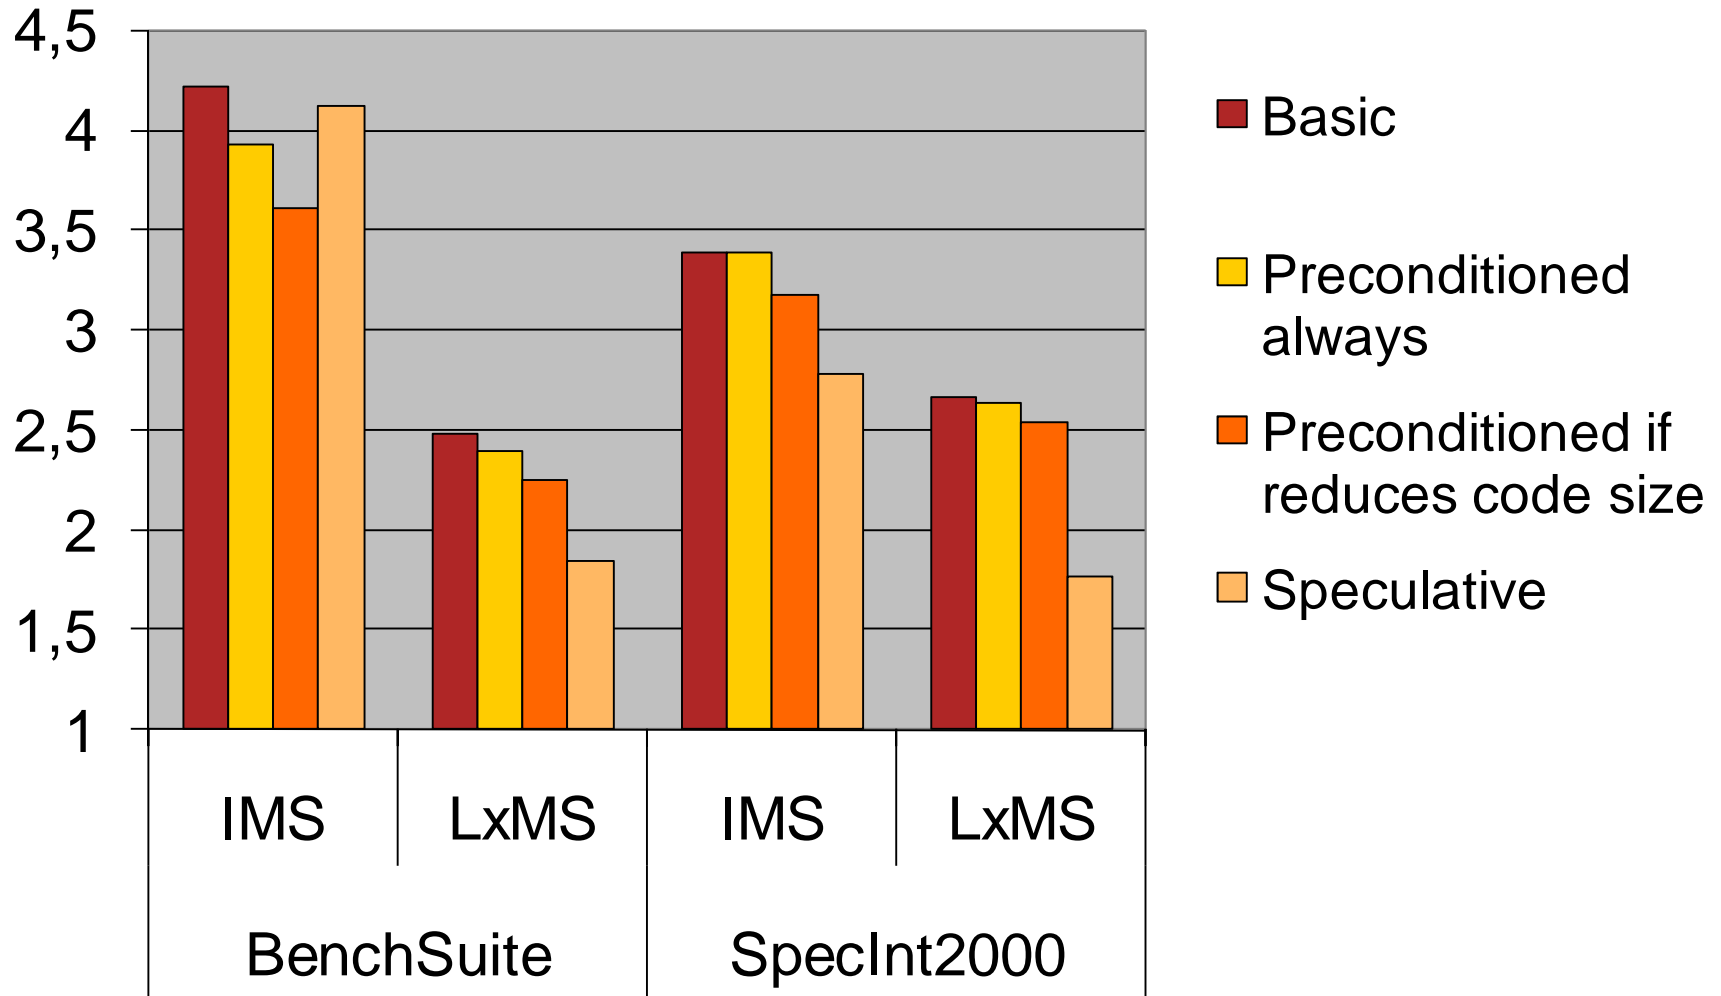
Average Stage Count



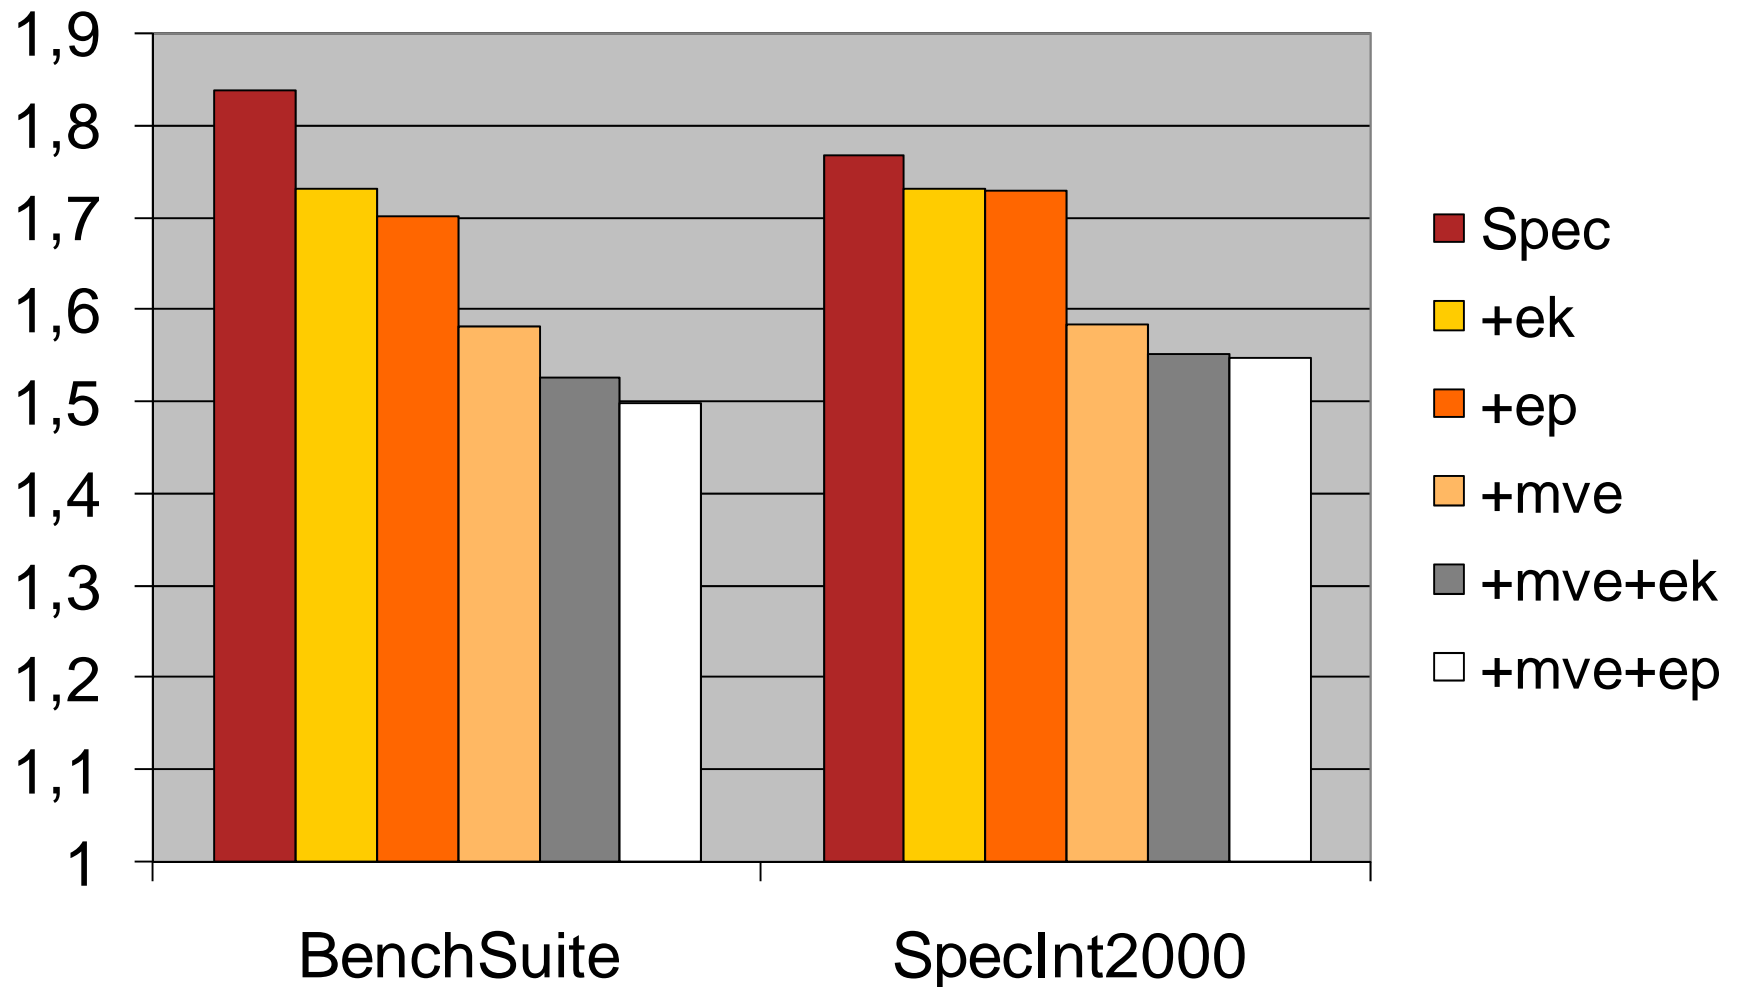
Average Early Exits



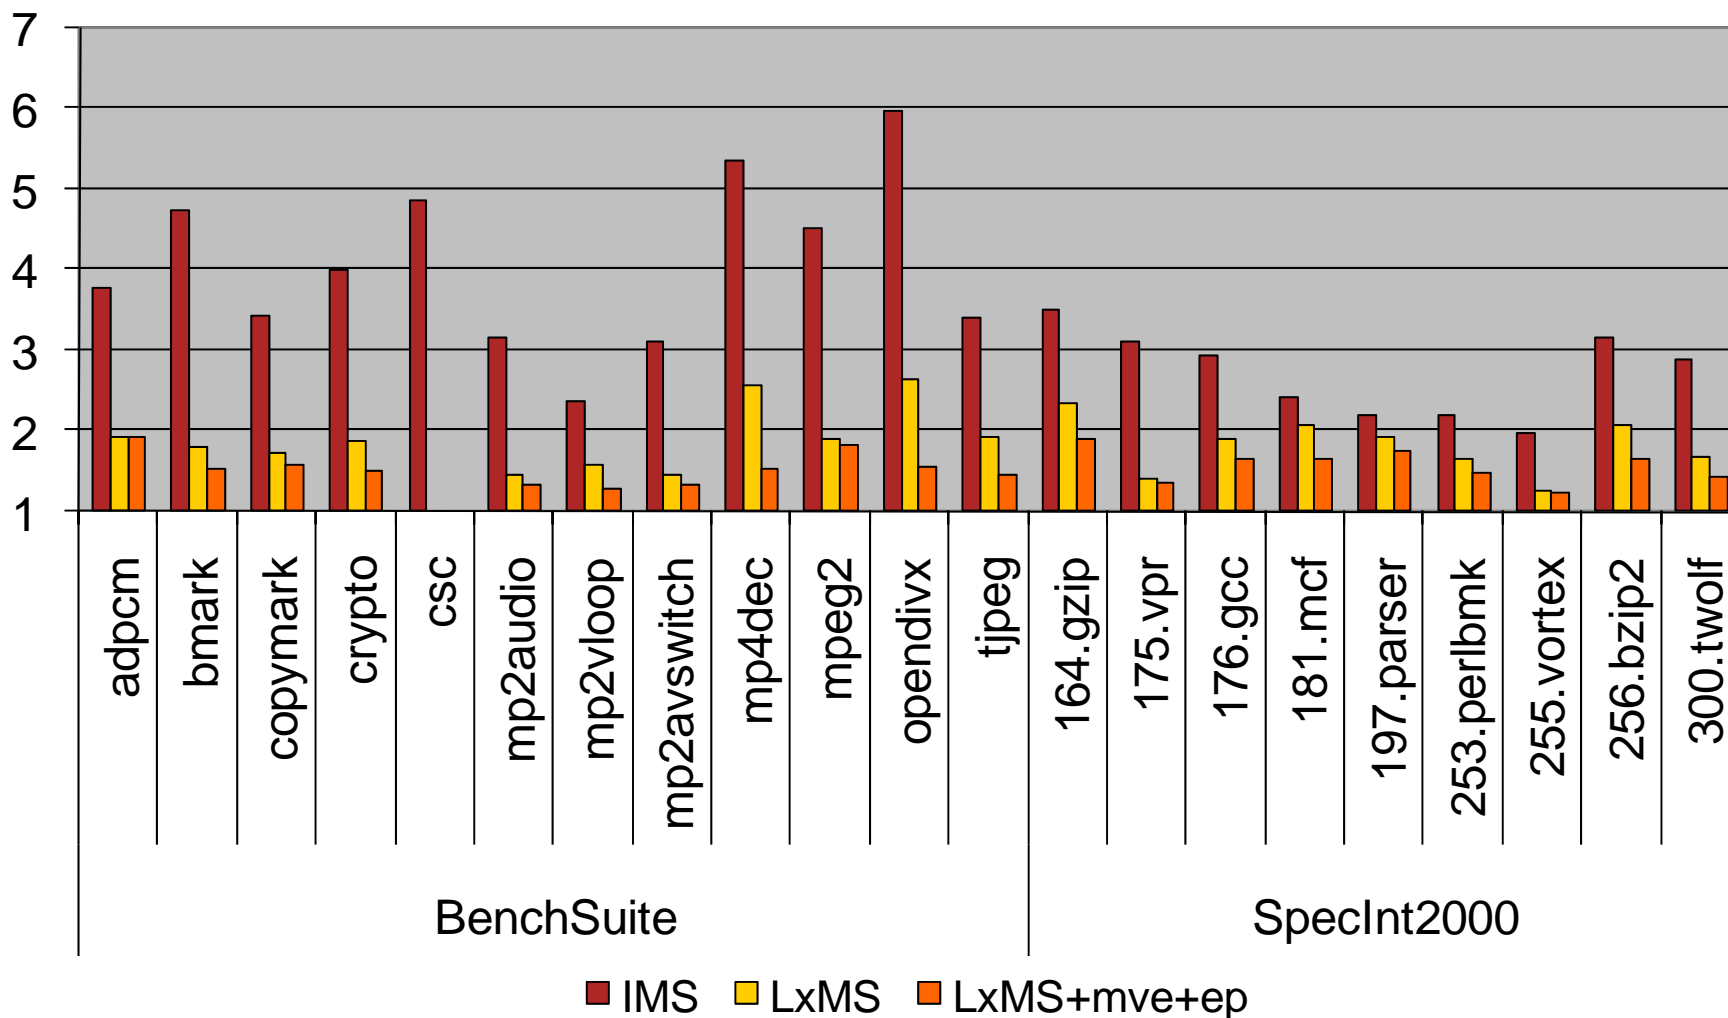
Code size expansion ratio



Code size reduction schemas

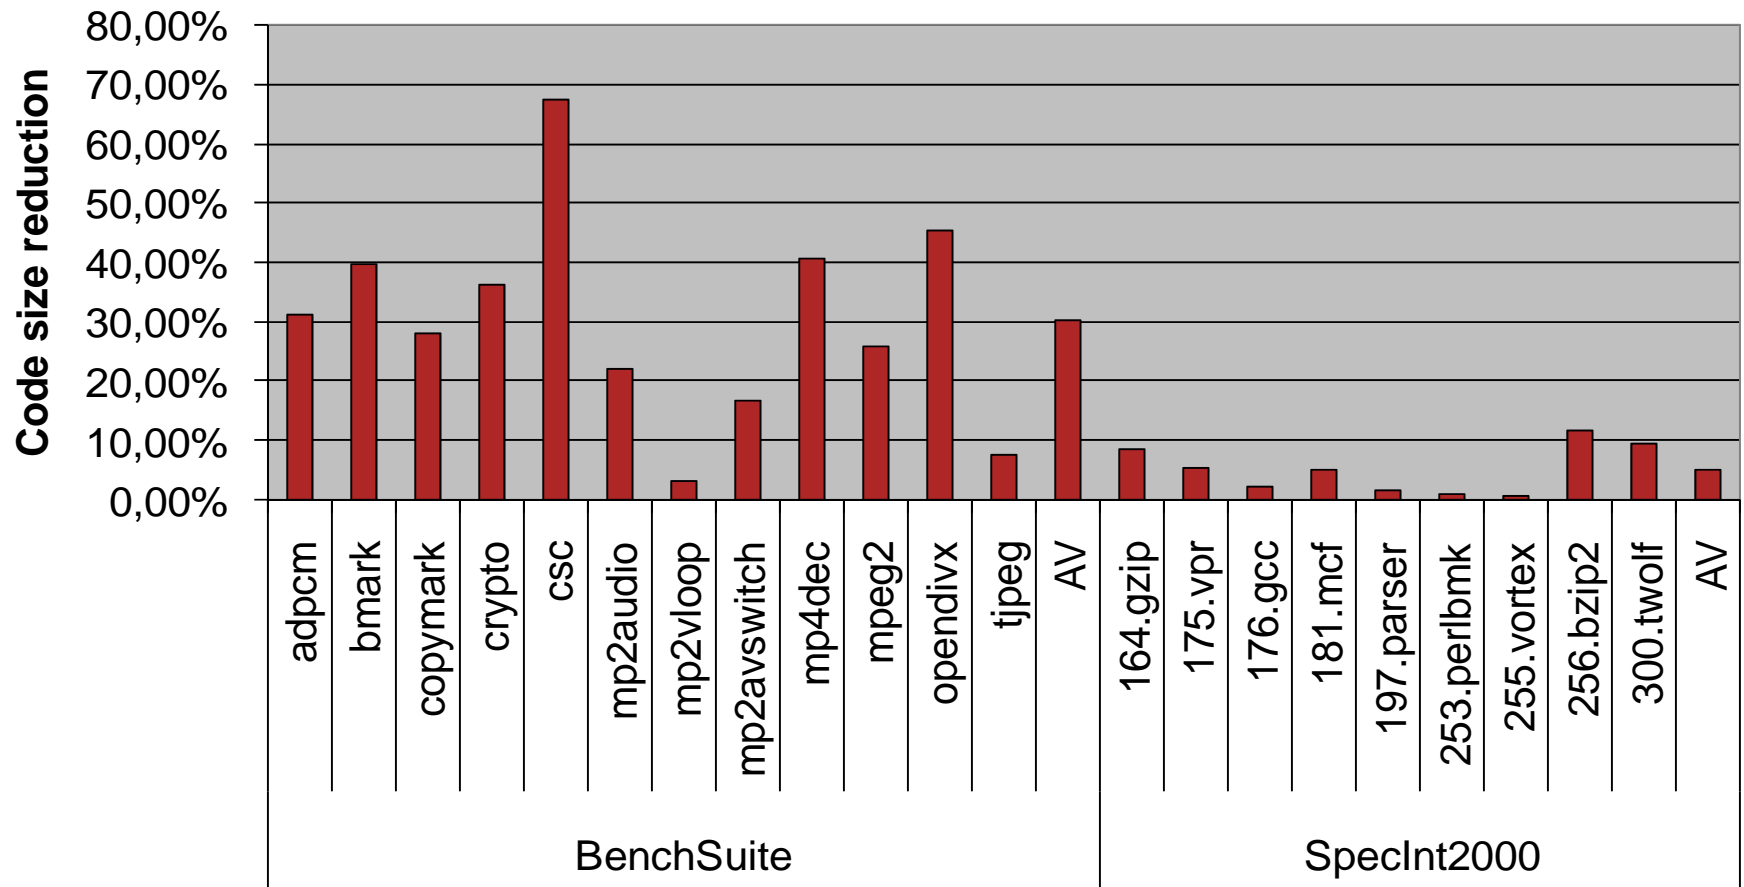


Code expansion ratio

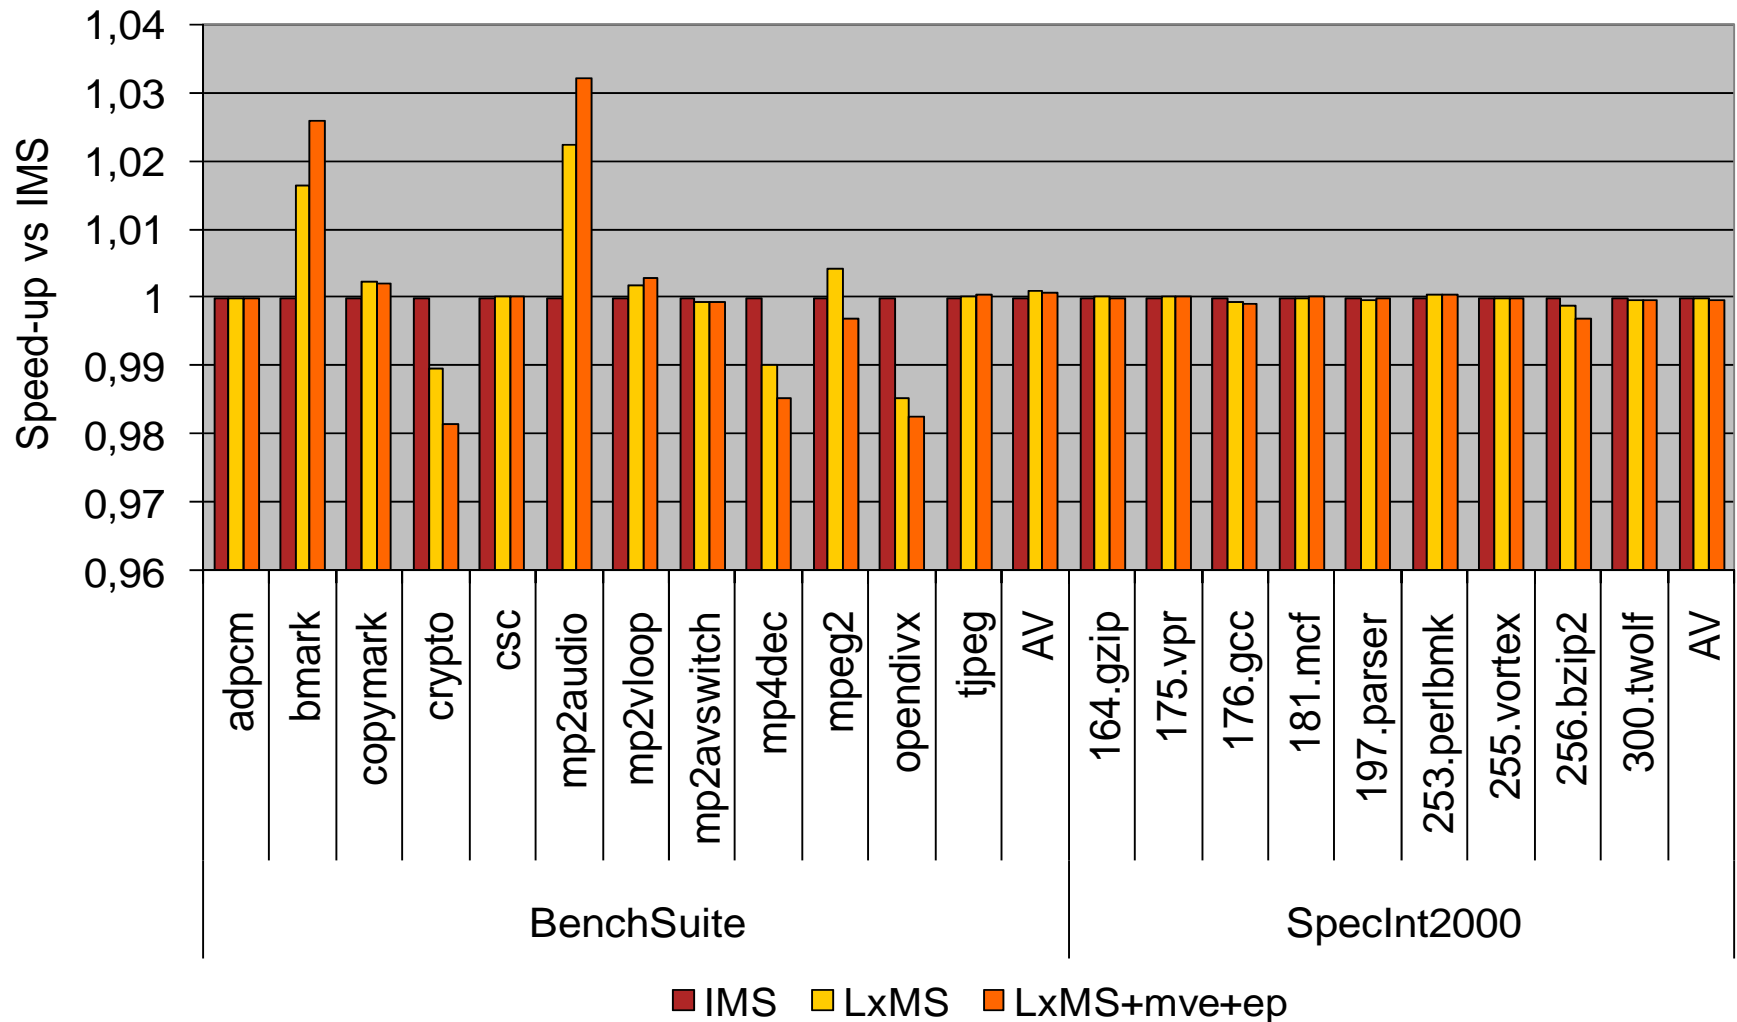


Overall code size reduction

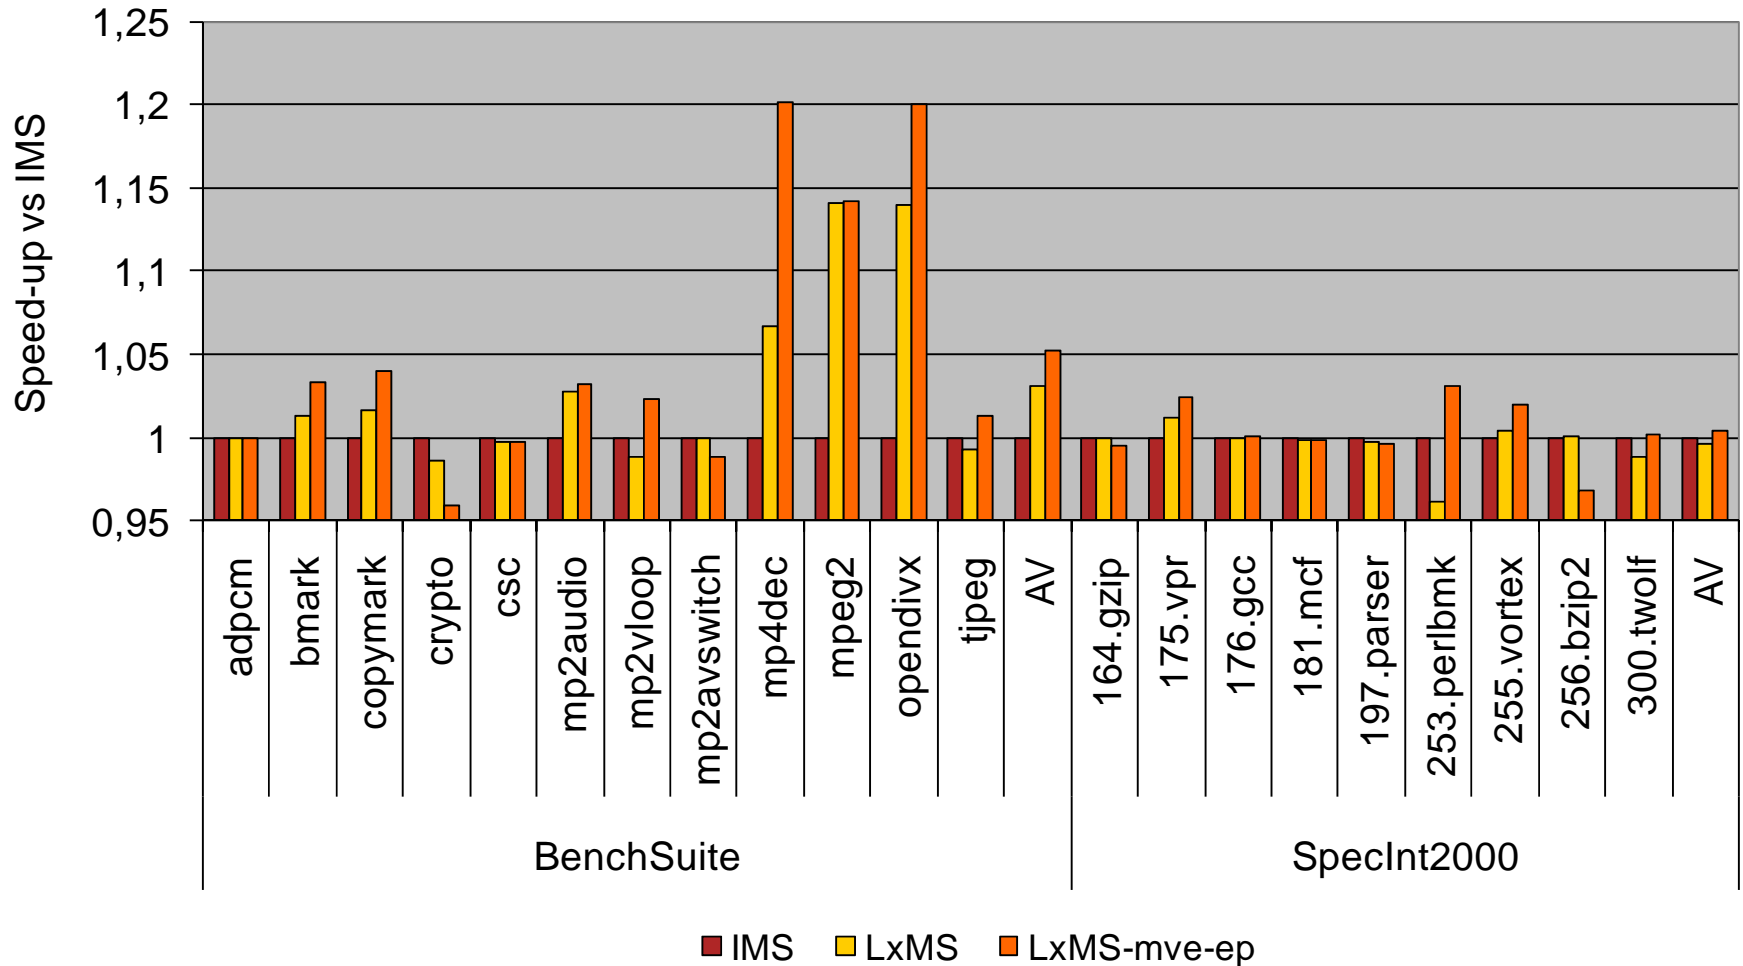
LxMS+mve+ep vs IMS



Impact on performance w/o memory stalls



Impact on performance w. memory stalls



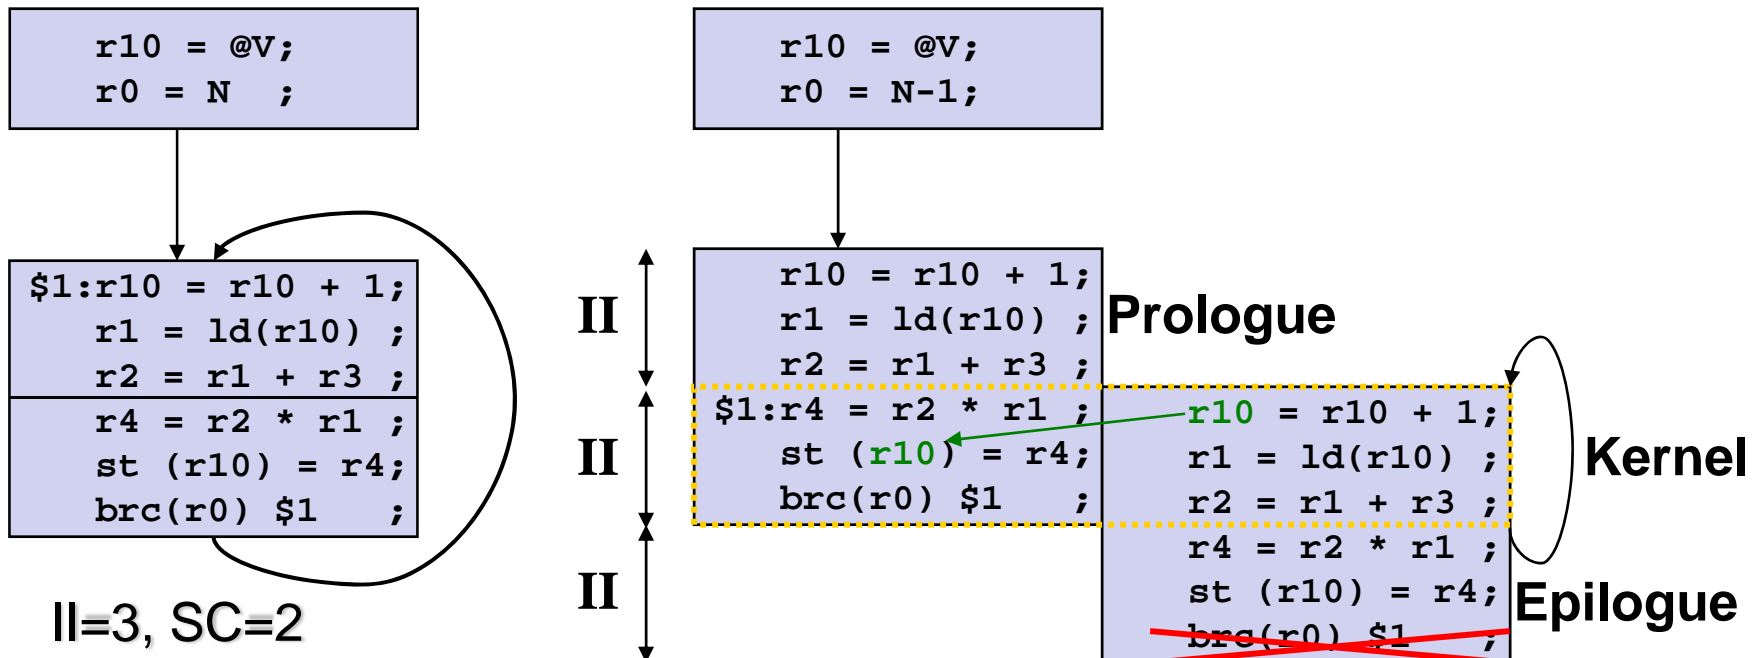
Hardware Support for Modulo Scheduling

Compilers for High Performance
Architectures



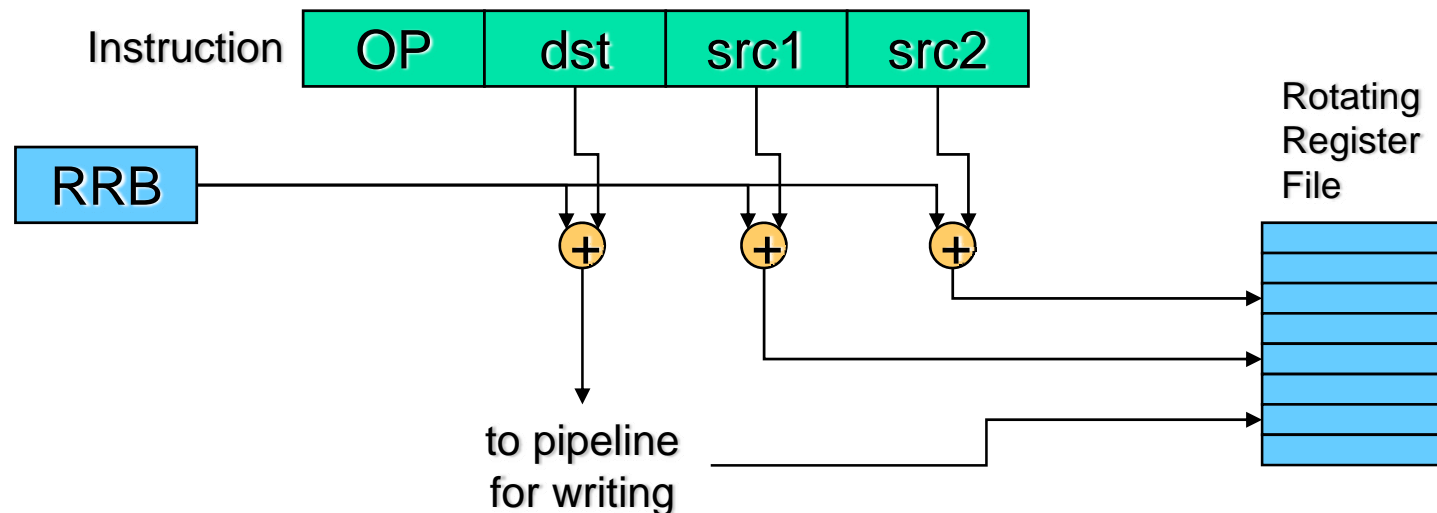
Code generation reminder

- Example is incorrect as is
 - R10 is rewritten by iteration (i+1) before using value generated in iteration (i)

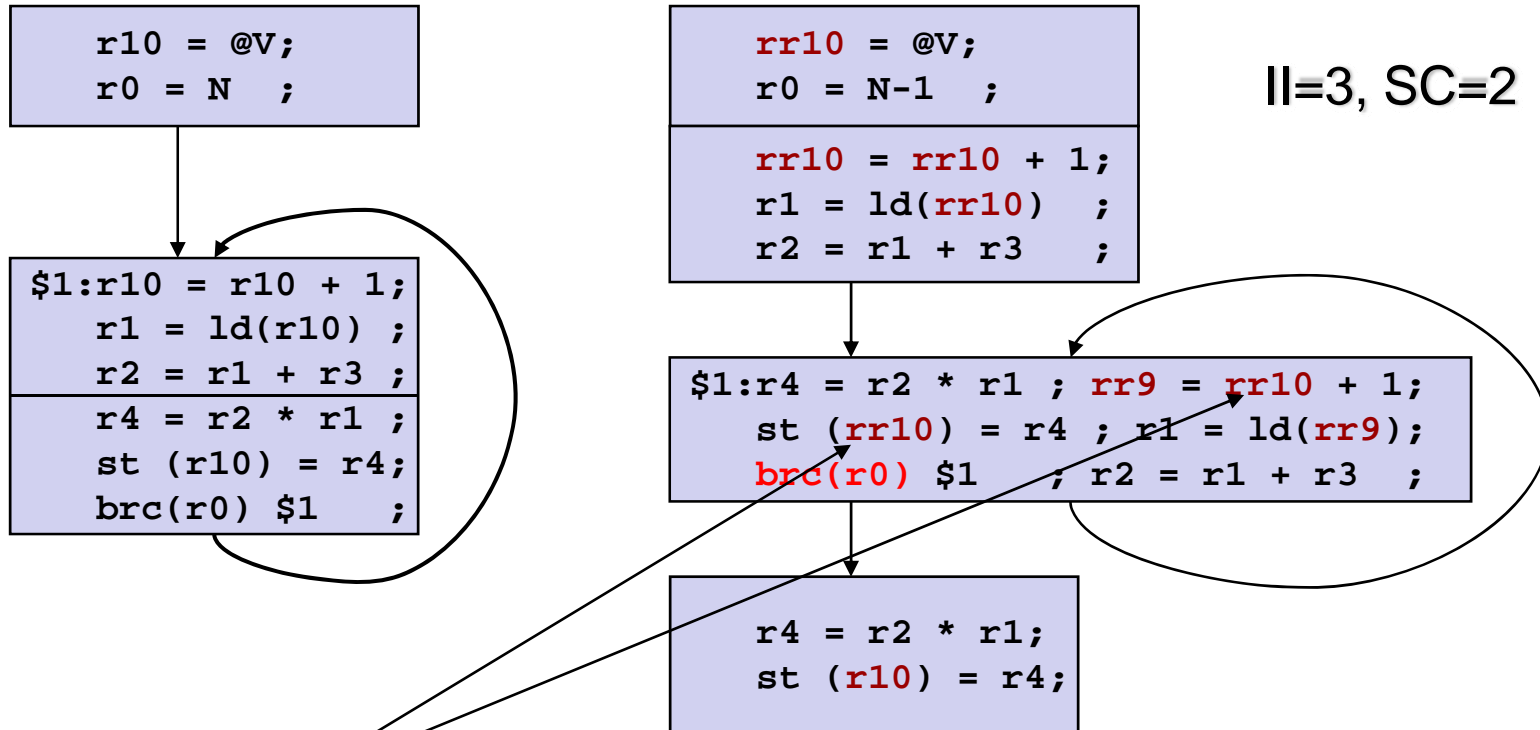


Rotating register files

- Differentiate between Logical & Physical registers
 - Contents of logical register R_i in iteration i
 - Are accessed through logical register R_{i+1} in iteration $i+1$
- Simple renaming using a rotating registers base (RRB)
 - $RP = (RRB + RL) \bmod NR$ (n-bit natural addition without carry)
 - RP = Physical register
 - RL = Logical Register
 - NR = Number of registers
 - RRB is decremented mod NR each kernel iteration



Rotating register files



rr10 refers to rr9 of previous iteration (CORRECT!!!)
loop closing branch instruction (brc) rotates registers

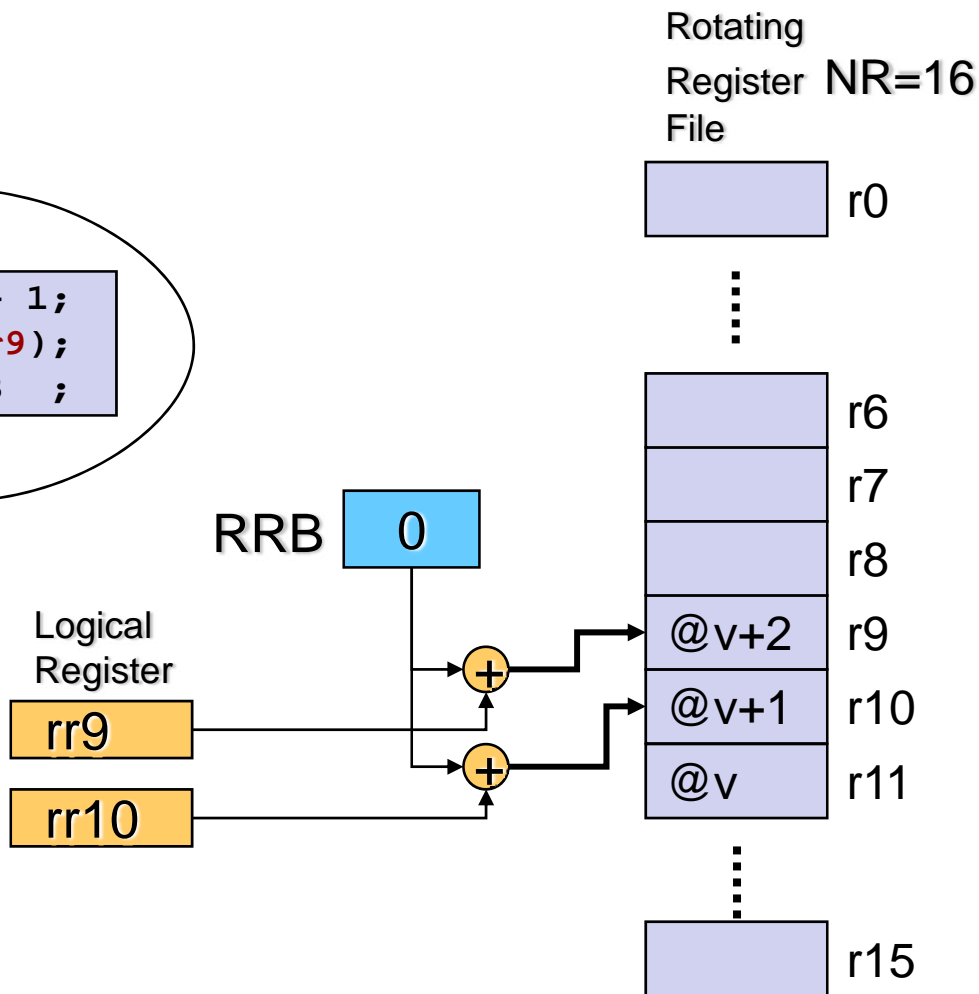
Rotating register files avoid MVE !!!

Rotating register files

```
rr11 = @V;  
r0 = N-1 ;  
  
rr10 = rr11 + 1;  
r1 = ld(rr10) ;  
r2 = r1 + r3 ;
```

```
$1:r4 = r2 * r1 ; rr9 = rr10 + 1;  
st (rr10) = r4 ; r1 = ld(rr9);  
brc(r0) $1 ; r2 = r1 + r3 ;
```

```
r4 = r2 * r1;  
st (r10) = r4;
```

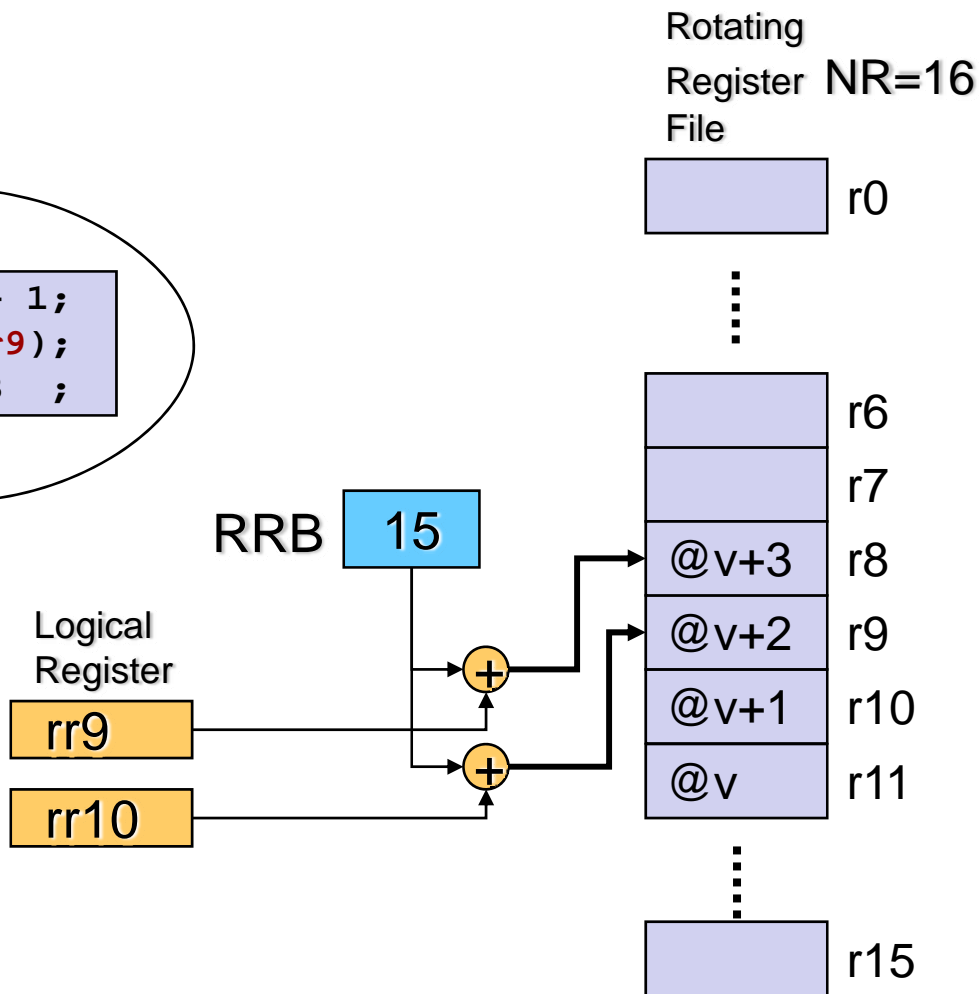


Rotating register files

```
rr11 = @V;  
r0 = N-1 ;  
  
rr10 = rr11 + 1;  
r1 = ld(rr10) ;  
r2 = r1 + r3 ;
```

```
$1:r4 = r2 * r1 ; rr9 = rr10 + 1;  
st (rr10) = r4 ; r1 = ld(rr9);  
brc(r0) $1 ; r2 = r1 + r3 ;
```

```
r4 = r2 * r1;  
st (r10) = r4;
```

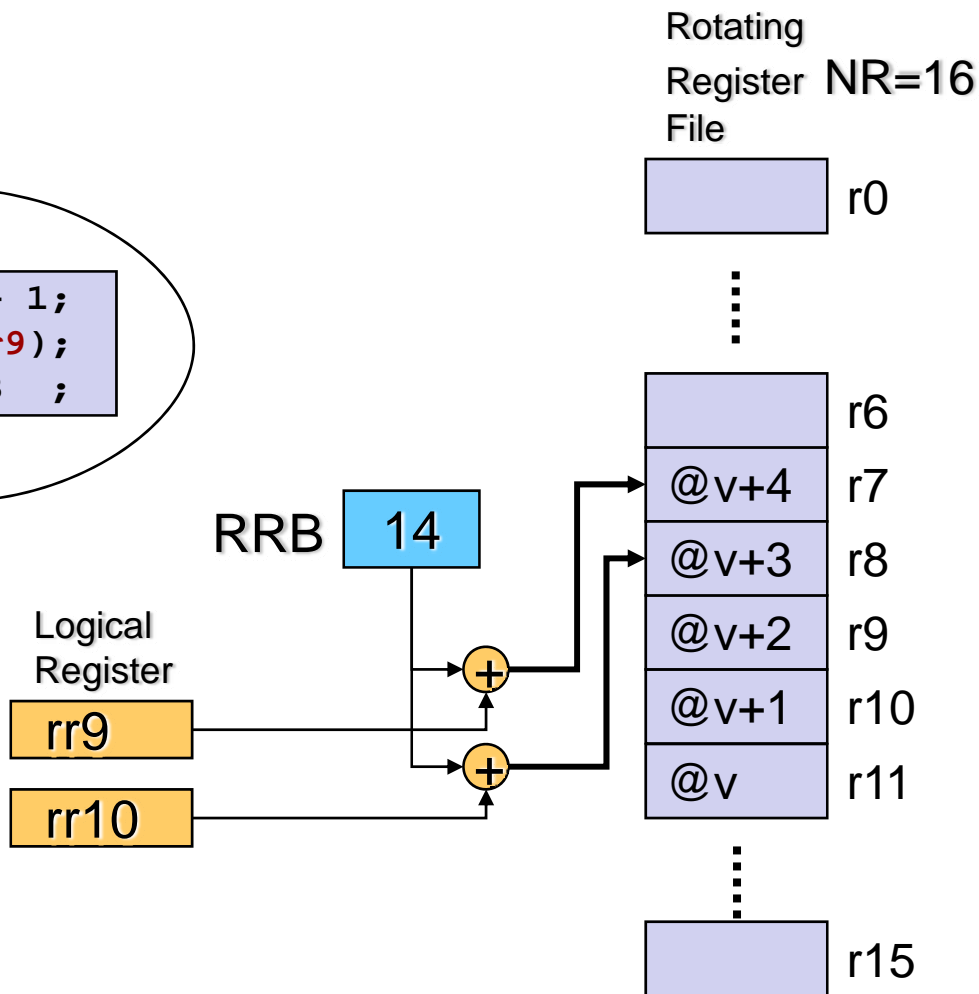


Rotating register files

```
rr11 = @V;  
r0 = N-1 ;  
  
rr10 = rr11 + 1;  
r1 = ld(rr10) ;  
r2 = r1 + r3 ;
```

```
$1:r4 = r2 * r1 ; rr9 = rr10 + 1;  
st (rr10) = r4 ; r1 = ld(rr9);  
brc(r0) $1 ; r2 = r1 + r3 ;
```

```
r4 = r2 * r1;  
st (r10) = r4;
```



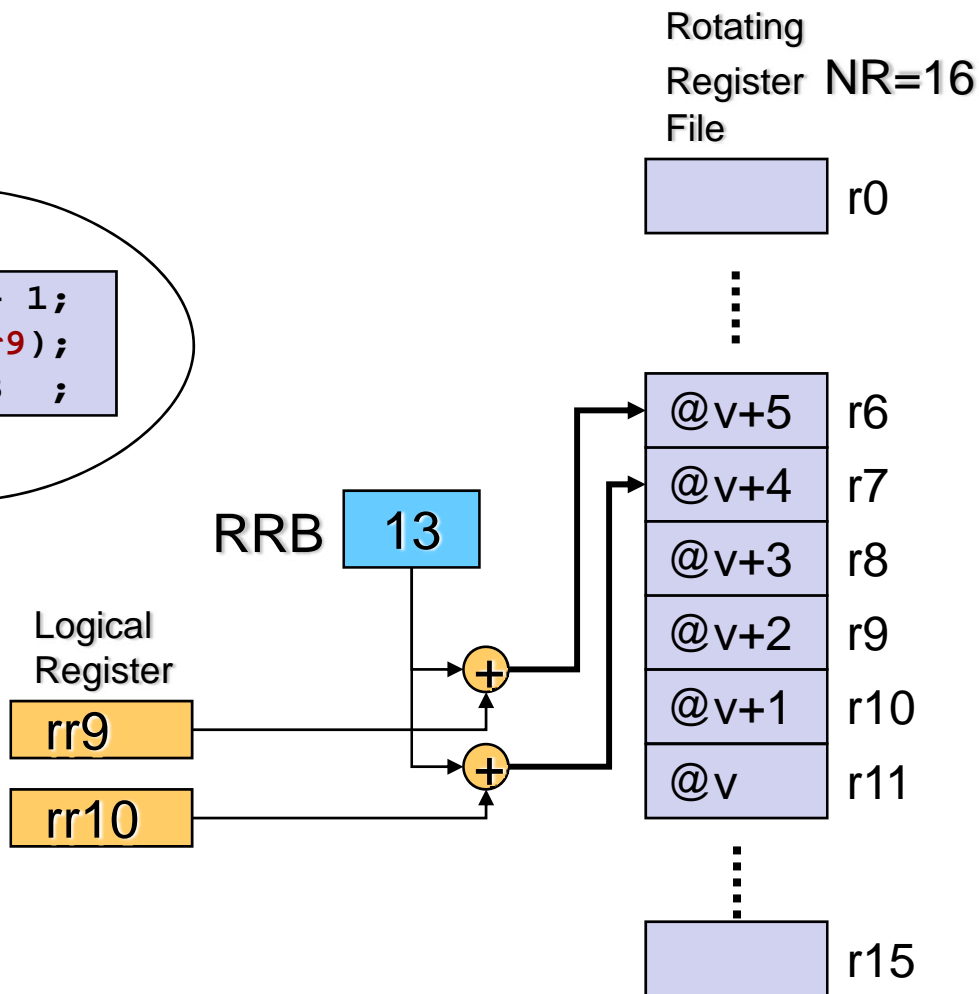
Rotating register files

```
rr11 = @V;  
r0 = N-1 ;
```

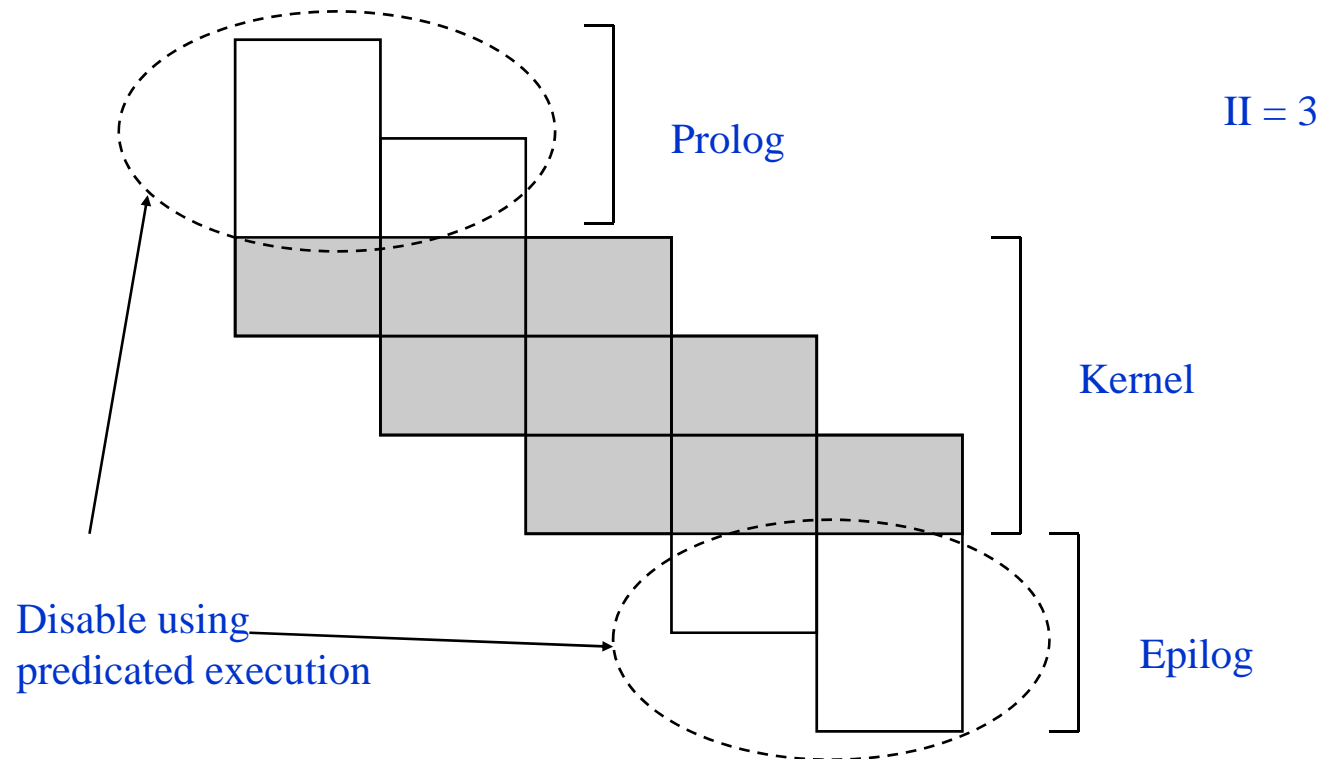
```
rr10 = rr11 + 1;  
r1 = ld(rr10) ;  
r2 = r1 + r3 ;
```

```
$1:r4 = r2 * r1 ; rr9 = rr10 + 1;  
st (rr10) = r4 ; r1 = ld(rr9);  
brc(r0) $1 ; r2 = r1 + r3 ;
```

```
r4 = r2 * r1;  
st (r10) = r4;
```

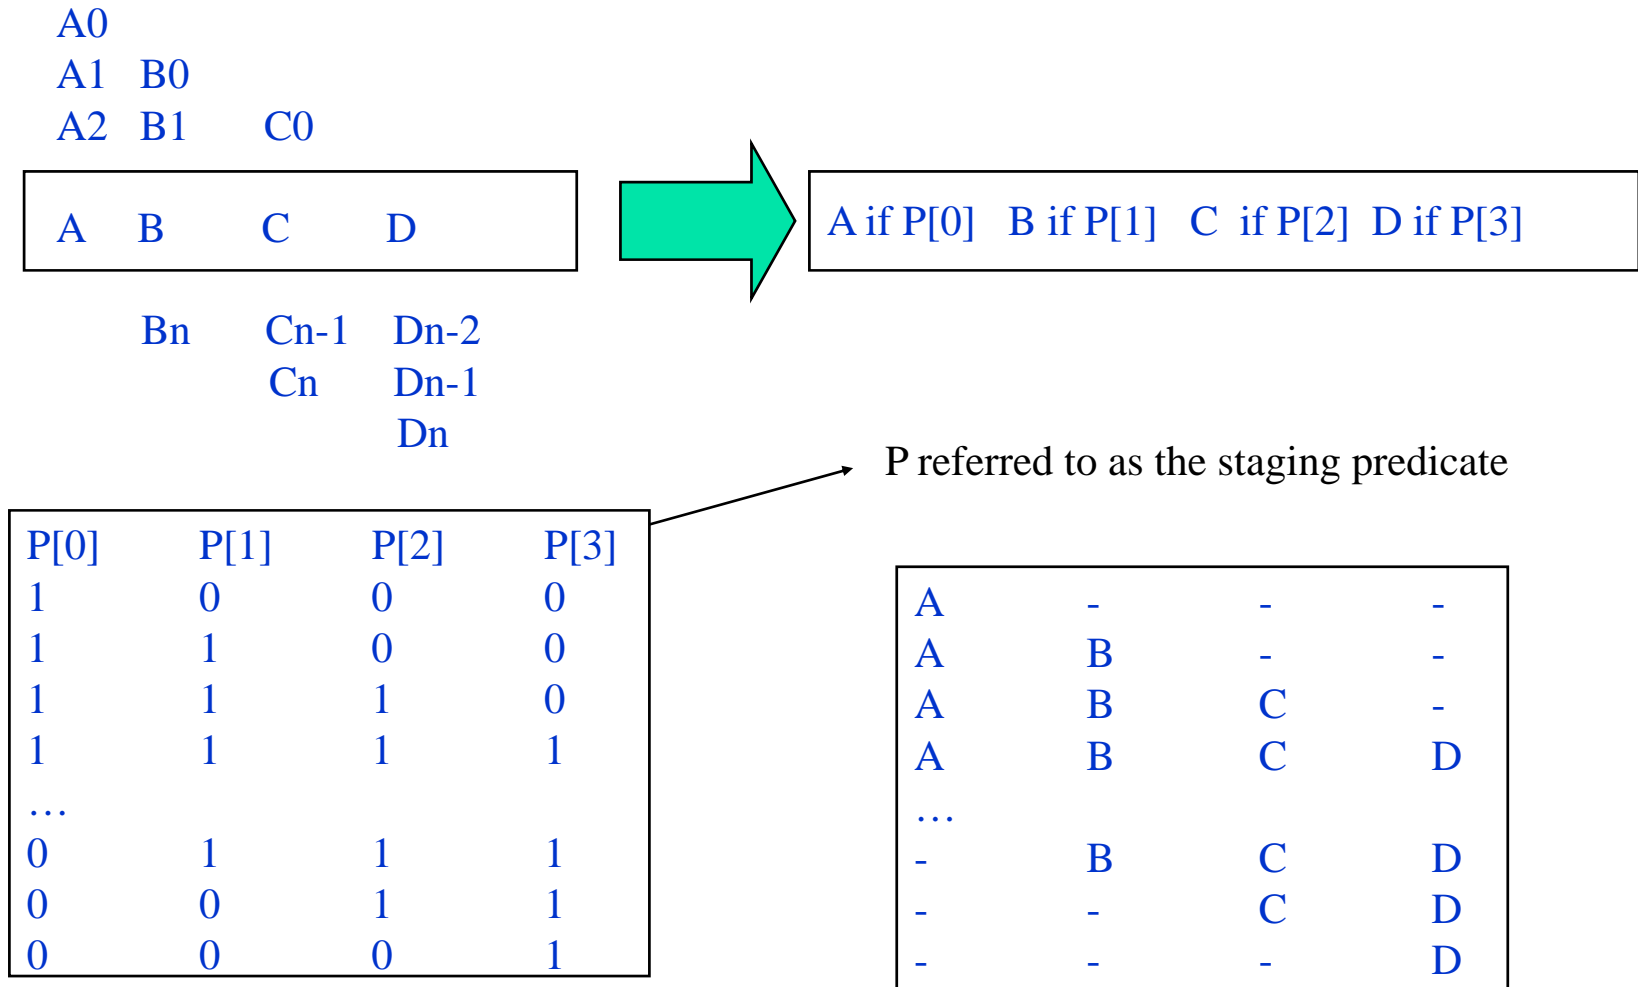


Removing Prolog/Epilog



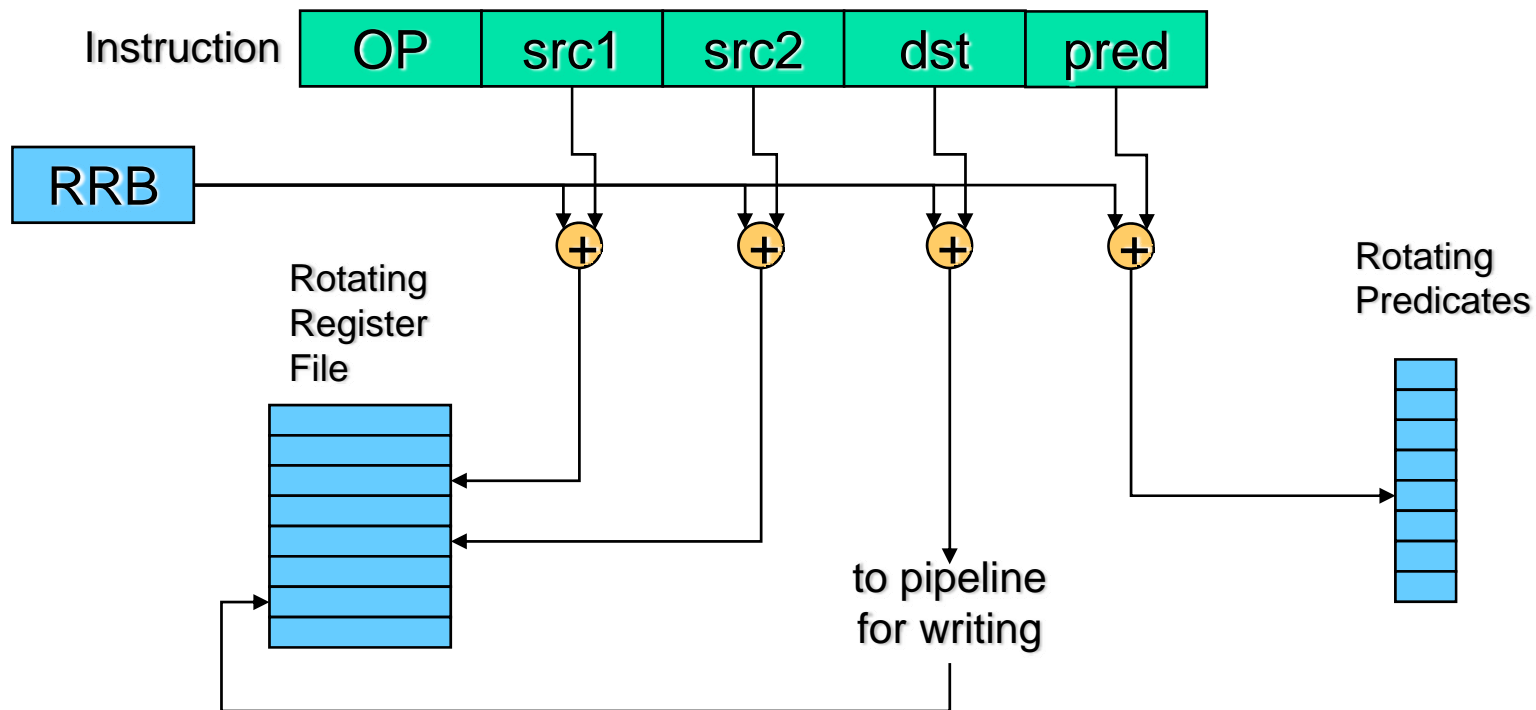
Execute loop kernel on every iteration, but for prolog and epilog selectively disable the appropriate operations to fill/drain the pipeline

Kernel-only Code Using Rotating Predicates



Rotating predicates

- Make predicates to rotate as well as registers
 - Contents of logical predicate P_i in iteration i
 - Are accessed through logical predicate P_{i+1} in iteration $i+1$
- Uses same renaming hardware as rotating registers



Modulo Scheduling Architectural Support

- Loop requiring N iterations
 - Will take $N + (S - 1)$ where S is the number of stages
- 2 special registers created
 - LC: loop counter (holds N)
 - ESC: epilog stage counter (holds S)
- Software pipeline branch operations
 - Initialize $LC = N$, $ESC = S-1$ in loop preheader
 - All rotating predicates are cleared
 - BRtop
 - While $LC > 0$, decrement LC and RRB , $P[0] = 1$, branch to top of loop
 - This occurs for prolog and kernel
 - If $LC = 0$, then while $ESC > 0$, decrement RRB and write a 0 into $P[0]$, and branch to the top of the loop
 - This occurs for the epilog

Execution History With LC/ESC

LC = 4, ESC = 3 /* Remember 0 relative!! */

Clear all rotating predicates

P[0] = 1

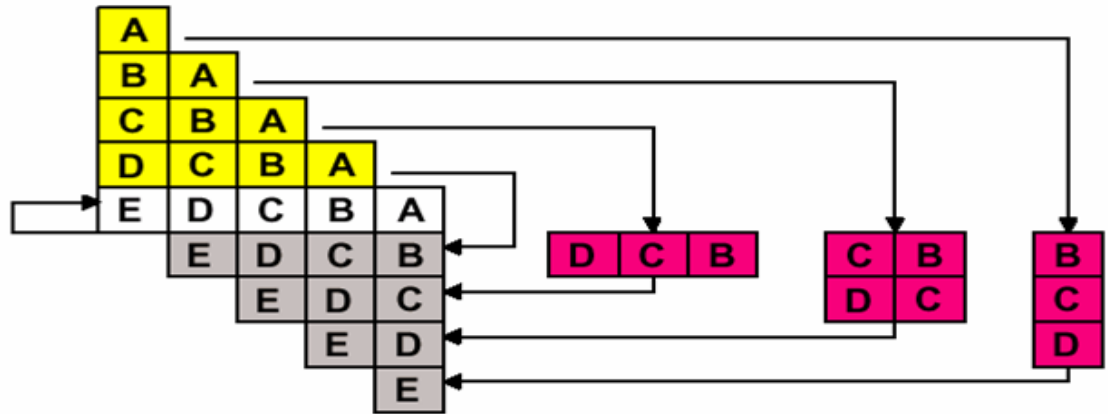
A if P[0]; B if P[1]; C if P[2]; D if P[3]; P[0] = BR_{top};

LC	ESC	P[0]	P[1]	P[2]	P[3]				
4	3	1	0	0	0	A			
3	3	1	1	0	0	A	B		
2	3	1	1	1	0	A	B	C	
1	3	1	1	1	1	A	B	C	D
0	3	1	1	1	1	A	B	C	D
0	2	0	1	1	1	-	B	C	D
0	1	0	0	1	1	-	-	C	D
0	0	0	0	0	1	-	-	-	D

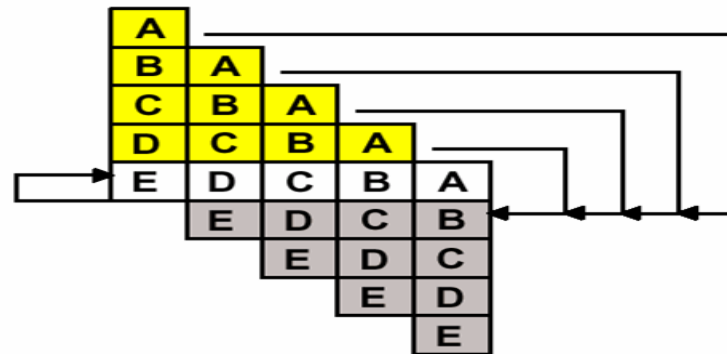
5 iterations, 4 stages, $\Pi = 1$, Note $5 + 4 - 1$ iterations of kernel executed

Code schemas with Hardware Support

With Rotating Registers: we can avoid MVE



With Predicates and Rotating Registers : we can avoid multiple epilogues



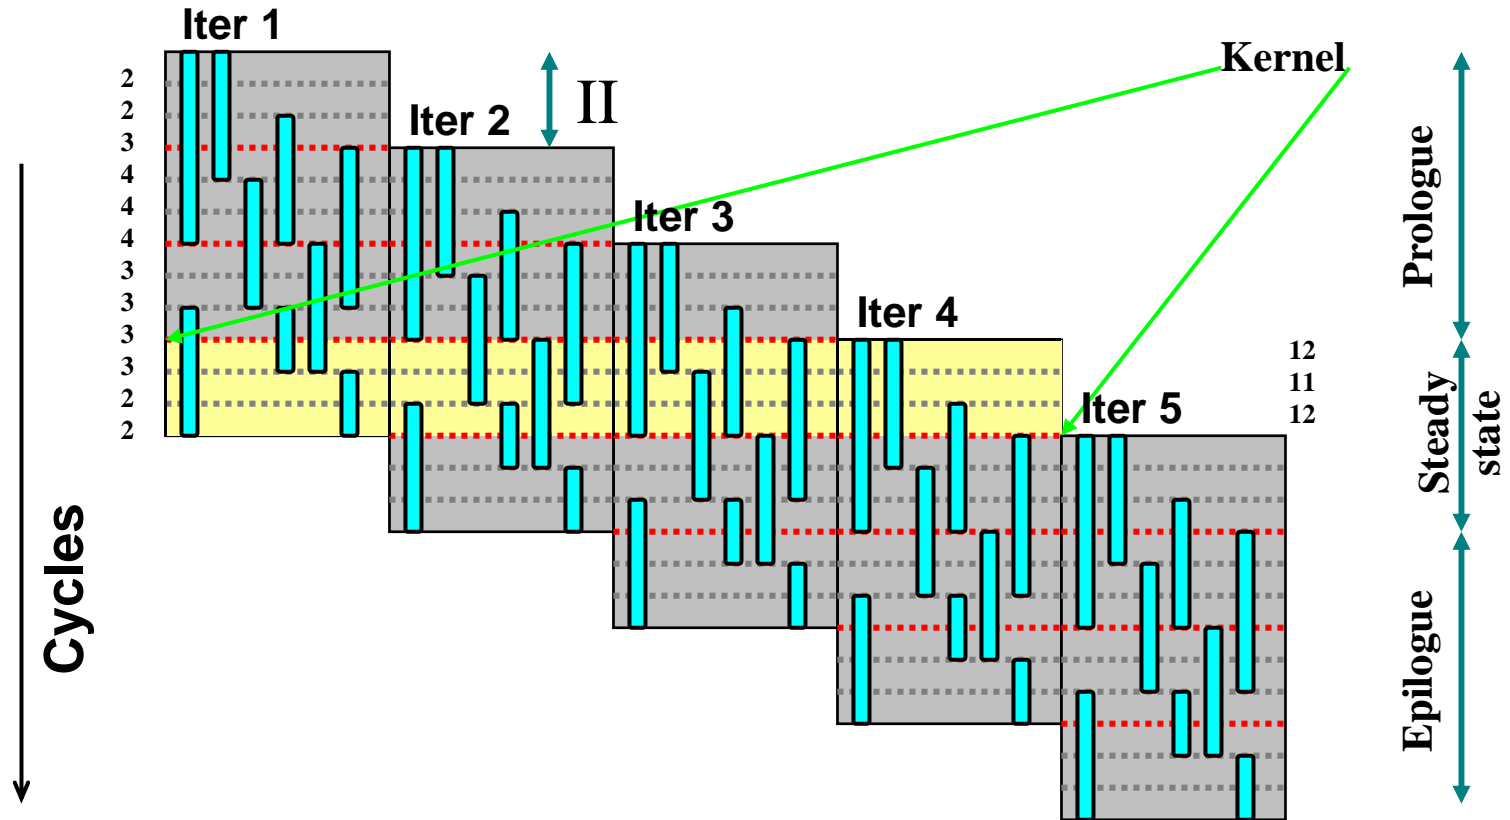
With Rotating Registers and Rotating Predicates: kernel-only code



From UIUC ECE 411 and NTU CA 718-Q
lecture notes © 1999 Wen-mei W. Hwu

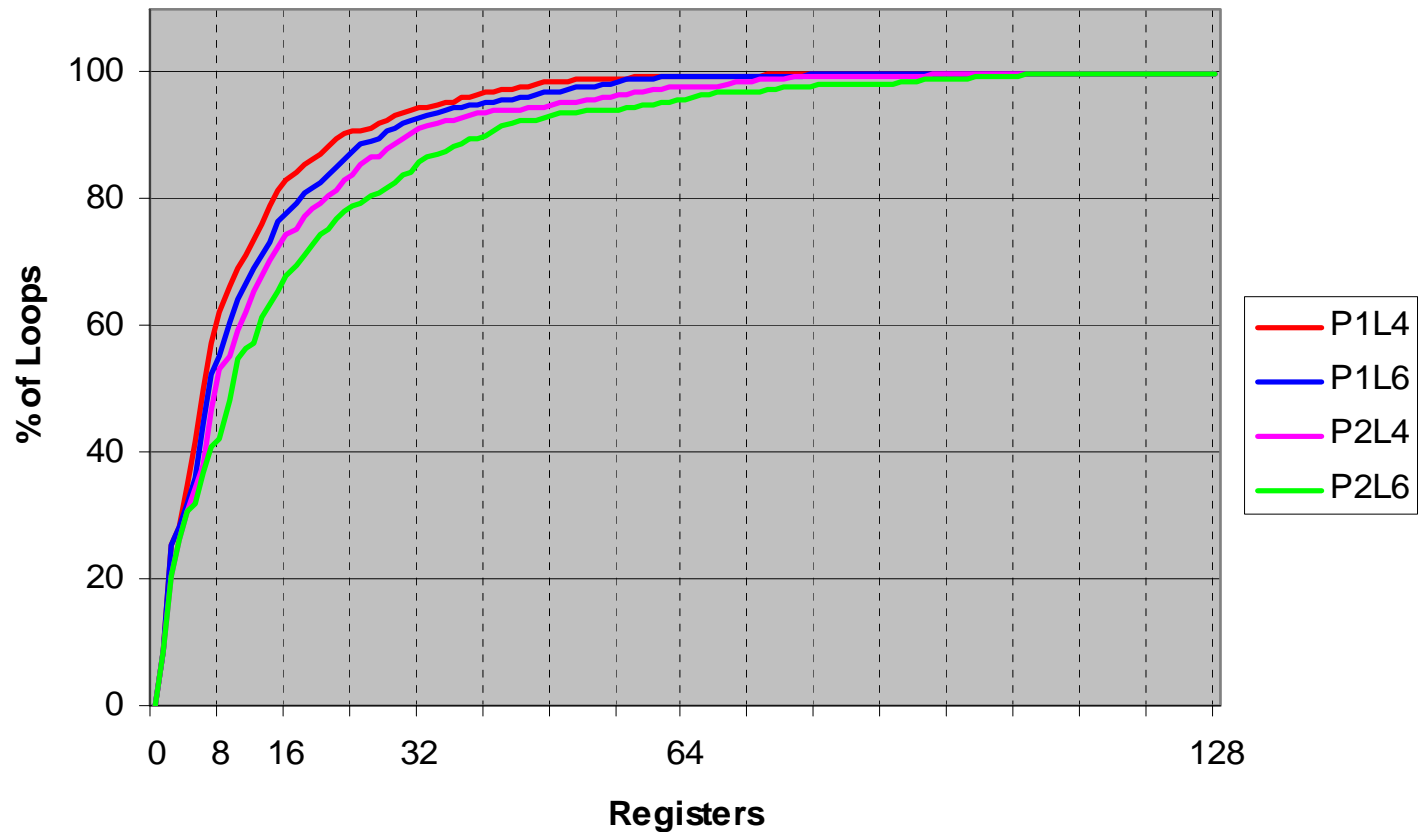
Register Requirements

- Aggressive scheduling techniques: high register requirements



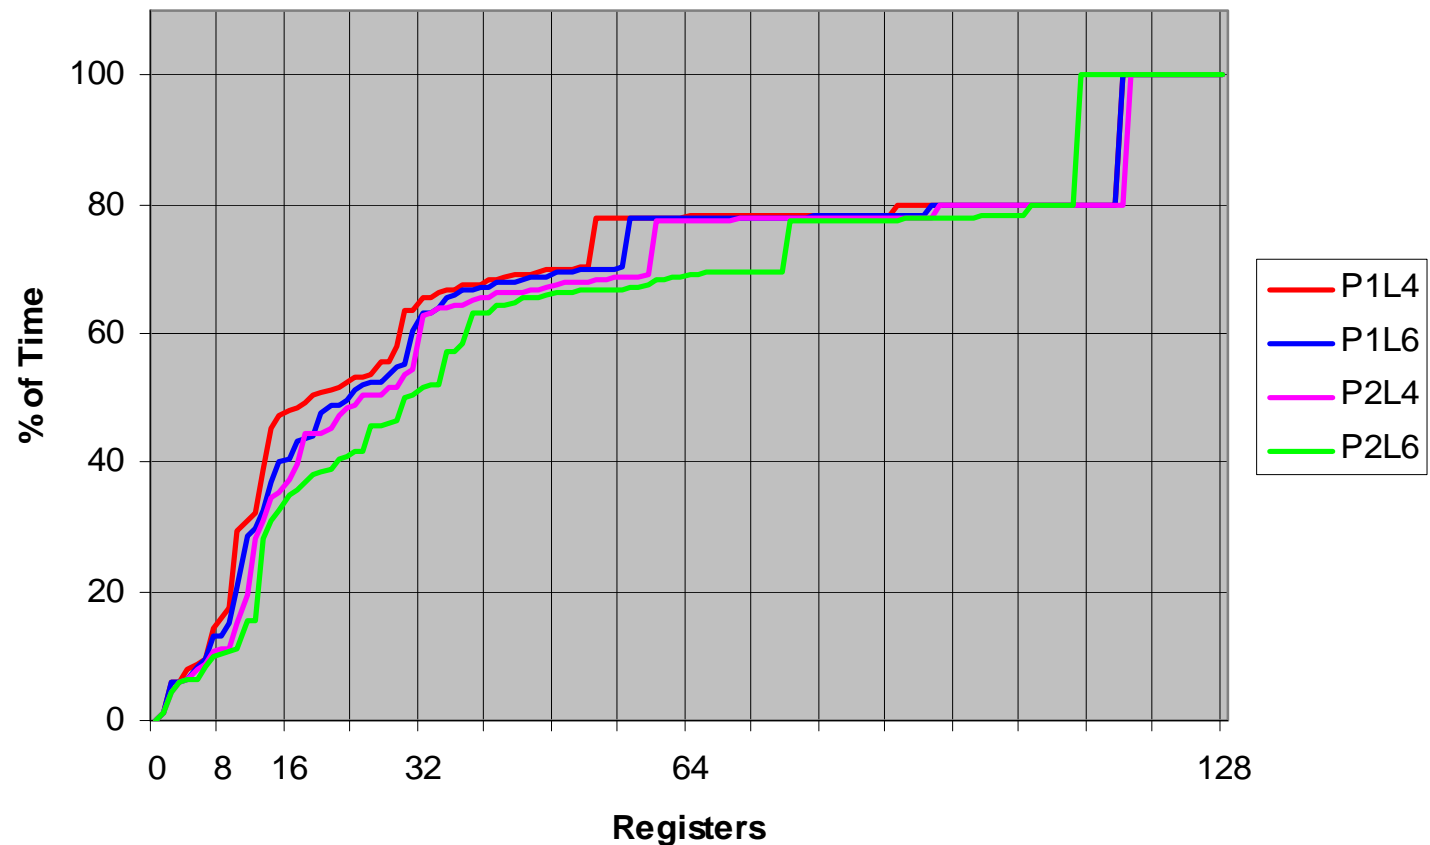
Register requirements

- Static register requirements



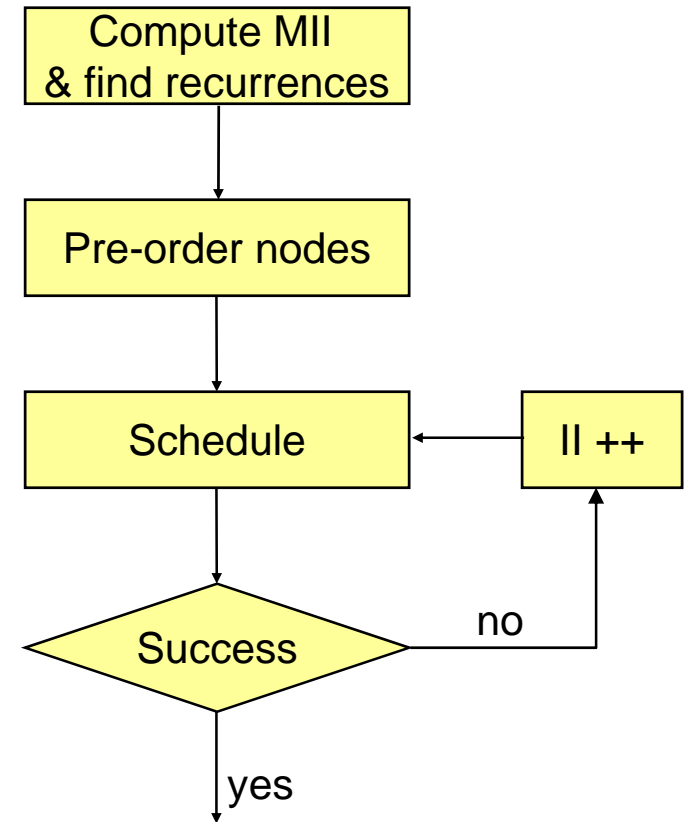
Register requirements

- Dynamic register requirements

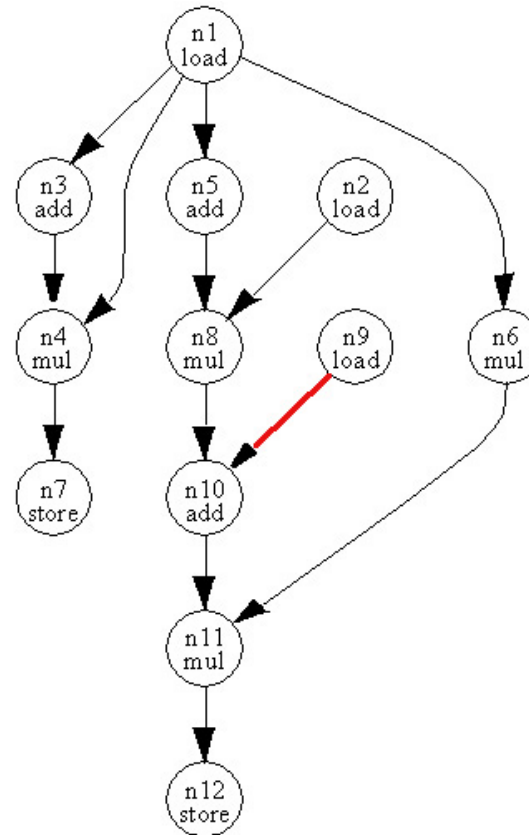


HRMS/SMS

- Proposed techniques:
 - HRMS: Hypernode Reduction Modulo Scheduling [micro95]
 - SMS: Swing Modulo Scheduling [pact96]
- Static priority function to pre-order nodes:
 - Hypernode reduction
 - Bidirectional ordering by adjacency
 - Special priority to recurrence circuits
 - Swing
 - Similar to above but faster and with enhanced critical path priority
- Simple scheduling:
 - bidirectional greedy modulo scheduling
 - Nodes with only predecessors are scheduled top-down
 - Nodes with only successors are scheduled bottom-up



Top-down vs. HRMS



Hardware configuration:

1 add unit
1 mul unit
2 load/store units

Latencies:

add: 2 cycles
mul: 2 cycles
load: 2 cycles
store: 1 cycle

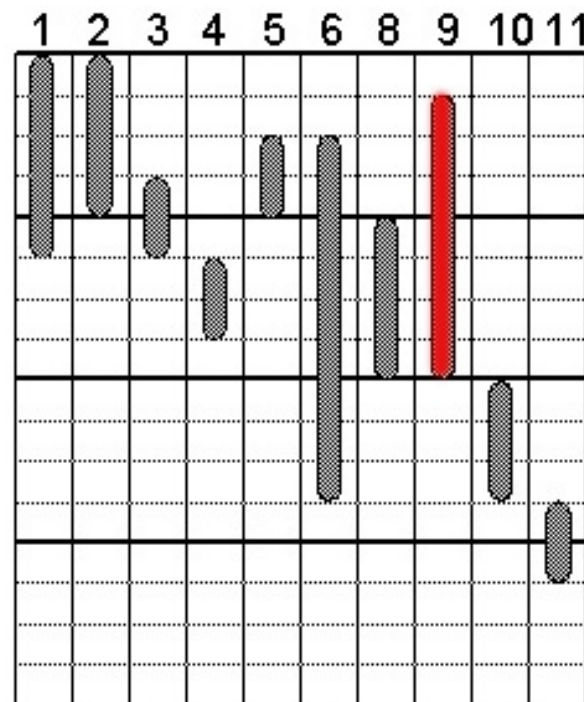
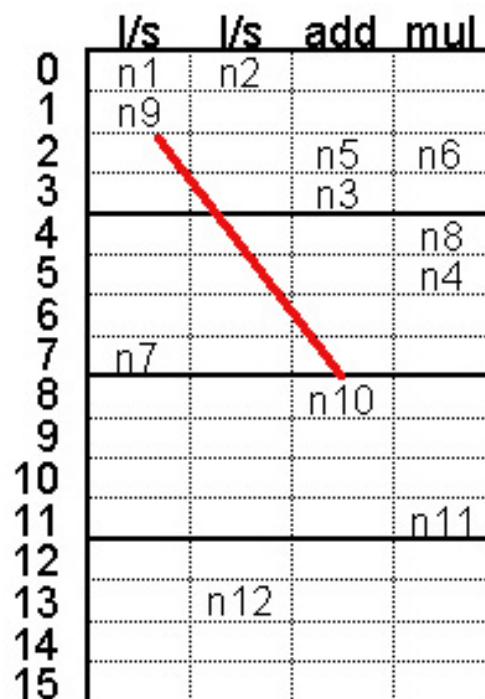
Pre-ordering step:

Top-down={ 1, 2, **9**, 3, 5, 6, 4, 8, 7, 10, 11, 12 }

HRMS={ 1, 3, 5, 6, 4, 7, 8, 10, 11, **9**, 2, 12 }

SMS={ 12, 11, 10, 8, 5, 6, 1, 2, **9**, 3, 4, 7 }

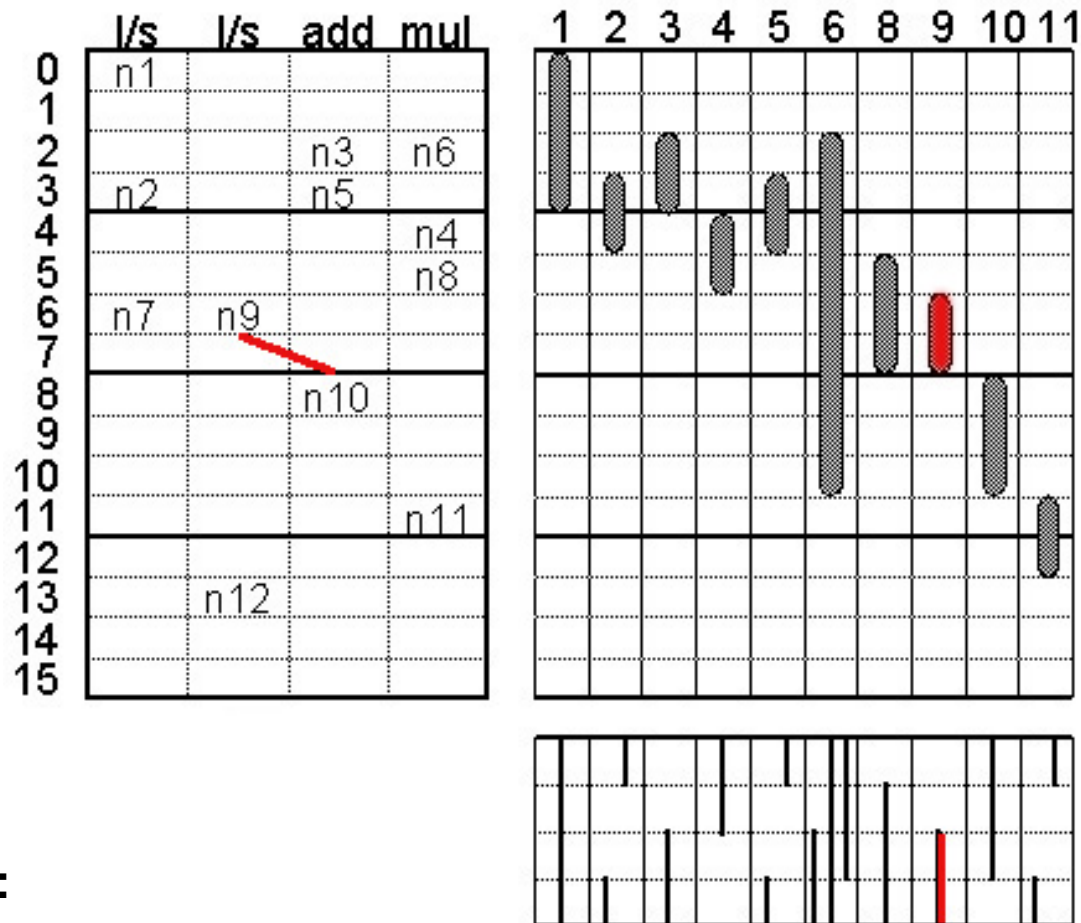
Top-down vs. HRMS



Pre-ordering step:

Top-down={ 1, 2, 5, 8, **9**, 3, 10, 6, 4, 11, 12, 7 }

Top-down vs. HRMS



Pre-ordering step:

HRMS={ 1, 3, 5, 6, 4, 7, 8, 10, 11, **9**, 2, 12 }

HRMS and SMS register requirements

