

# Heterogeneous Computing for Scientific Applications

Marc Gonzalez Tallada

Associate professor

Computer Architecture Department  
Technical University of Catalonia



UNIVERSITAT POLITÈCNICA  
DE CATALUNYA  
BARCELONATECH

# Outline

- Analysis, parallelism characterization of NAS-MZ suite of benchmarks
- NPB-MZ Hybrid Design and Implementation
- Work distribution schemes for hybrid executions
- Evaluation Results

# NPB-MZ Parallel Benchmarks

- Data Structures and Input Size
- Computation Structure
- Sources of Parallelism

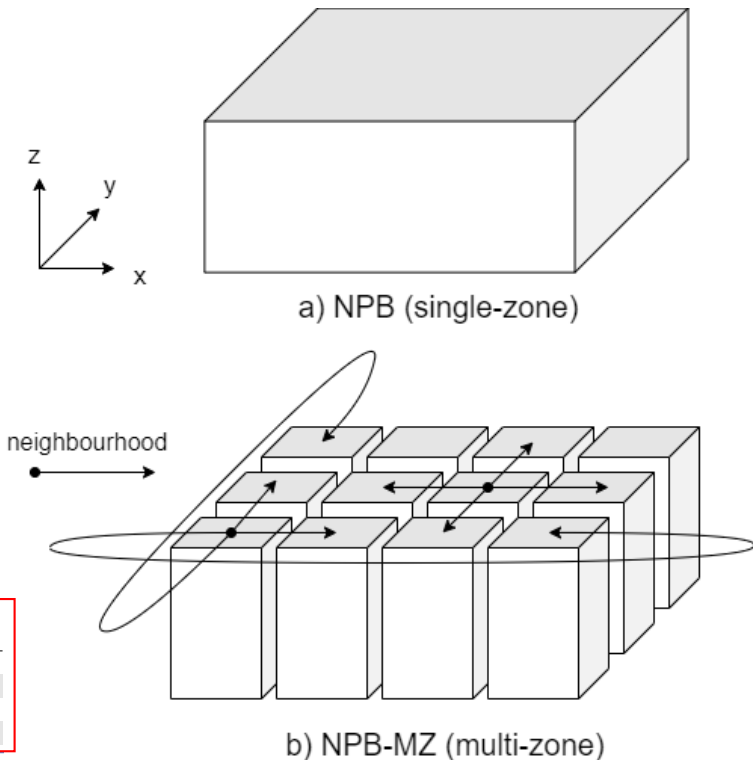
# Data Structures and Input Size

- Data Structures
  - 3D mesh organized in “zones”
    - A zone is a set of multidimensional matrices
- Input Size
 

size increases →

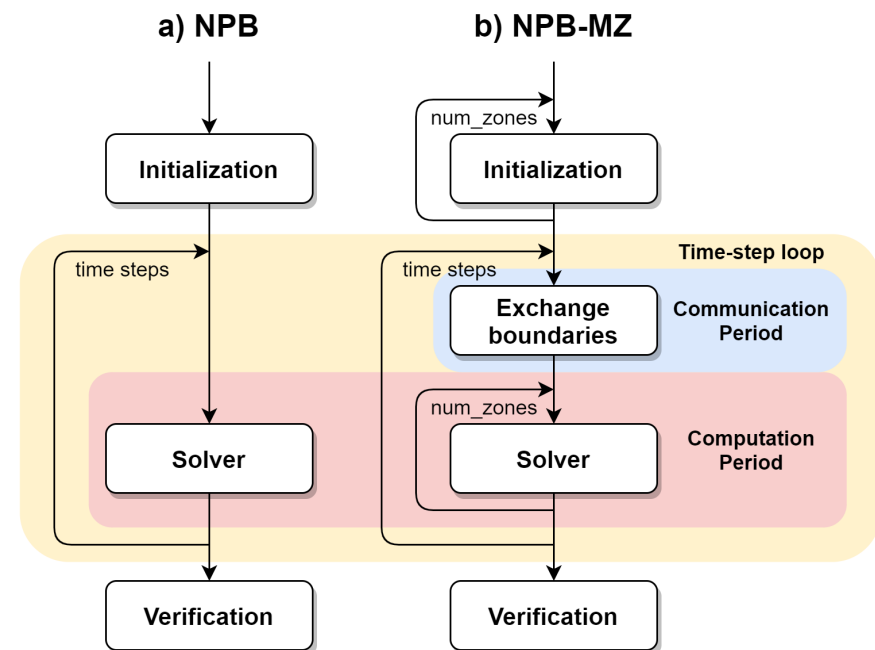
  - Different classes: A, B, C, D, E
  - Number of zones: from 16 to 4096
- Iterative solvers
  - From 250 to 500 iterations

Input Class	3D volume $x \times y \times z$ (points)	Memory (GB)	Num. zones ( $x \times y$ )		Zone size (points per zone)			Time steps		
			LU	SP & BT	LU	SP	BT	LU	SP	BT
B	$304 \times 208 \times 17$	$\approx 0.2$	$4 \times 4$	$8 \times 8$	67 184	16 786	from 2 992 to 59 976	250	400	200
C	$480 \times 320 \times 28$	$\approx 0.8$	$4 \times 4$	$16 \times 16$	268 800	16 800	from 2 912 to 60 648	250	400	200
D	$1\,632 \times 1\,216 \times 34$	$\approx 13$	$4 \times 4$	$32 \times 32$	4 217 088	65 892	from 11 968 to 243 236	300	500	250
E	$4\,224 \times 3\,456 \times 92$	$\approx 250$	$4 \times 4$	$64 \times 64$	83 939 328	327 888	from 59 248 to 1 203 452	300	500	250



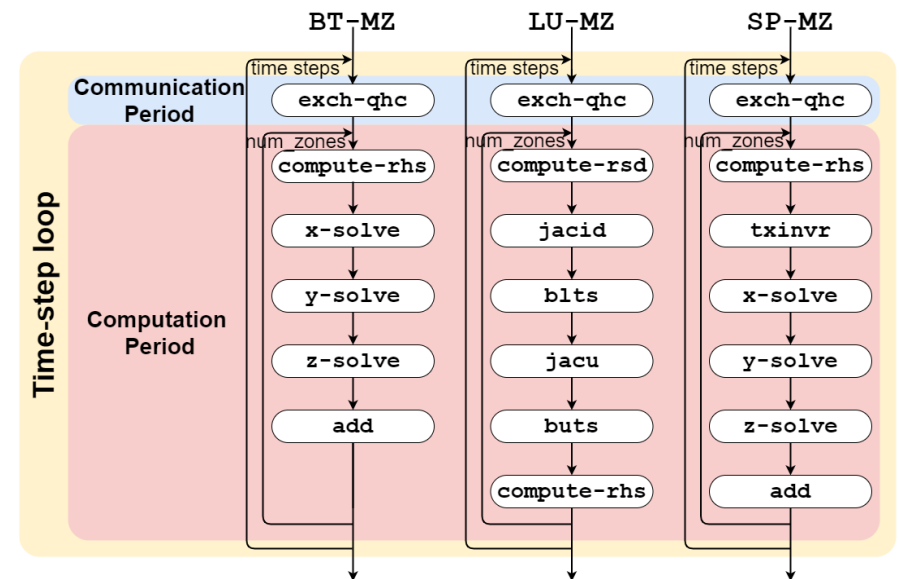
# Computational Periods

- For all benchmarks:
  - Initialization
    - Several stages
  - Computation Period
    - Solver composed of several stages
    - Applied to each zone
  - Communication Period
    - Exchange of border values between zones
    - One single stage
- Verification



# Stages within the Compute Period and Communication Period

- Three benchmarks:
  - BT-MZ
    - “Many zones with small, medium and large sizes”
  - SP-MZ
    - “Many zones with small sizes and all of them equal”
  - LU-MZ
    - “Few zones and with very large size”
- Computation Period
  - Composed of different stages per each benchmark
- Communication Period
  - Single stage and common to each benchmark



# Computation Period(I): Inter-Zone Parallelism

- Traversal of zones
  - Single loop that runs over the set of zones
- Zones are identified
  - Number from 0 up to the NUM-ZONES-1
- Several stages applied per each zone
  - Different in BT-MZ, SP-MZ and LU-MZ

- Example with SP-MZ:

```
for (t=0; t<ITERS; t++) {  
    exch-qbc();  
    for (zone = 0; zone < NUM-ZONES; zone++) {  
        compute-rhs(zone, ...); Inter-zone parallelism  
        txinvr(zone, ...);  
        x-solve(zone, ...);  
        y-solve(zone, ...);  
        z-solve(zone, ...);  
        add(zone, ...);  
    }  
}
```

# Computation Period(II): Intra-Zone Parallelism


- Processing of one stage on one zone
  - Sequence of several loop nests
  - For CPU execution
    - Loop nest is parallelized (e.g.: OpenMP)
  - For GPU execution
    - Loop nest is transformed to CUDA kernel
- Example: BT-MZ, x-solve stage

```
void x-solve(unsigned int zone, ...) {  
    Intra-zone parallelism  
    for (k=0; k<z-dim; k++)  
        for (j=0; j<y-dim; j++)  
            for (i=0; i<x-dim; i++)  
                /* Some matrix-based  
                 computations */  
  
    /* Several computations in the  
       form of loop nests */  
}
```



# Communication Period: Zone Adjacency

- Adjacent zones
  - 4 neighbors
    - north
    - south
    - east
    - west
- Exchange border values between adjacent zones



# Communication Period: Intra-Zone Parallelism

- Processing of zone borders

```
for (zone=0; zone<NUM-ZONES; zone++) {  
    east-zone = adjacency-east[zone]  
    north-zone = adjacency-north[zone];  
  
    copy-face(tmpEast, mesh[east-zone]);  
    copy-face(tmpNorth, mesh[north-zone]);  
  
    compute-border(mesh[zone], tmpEast, tmpNorth);  
  
    copy-face(mesh[east-zone], tmpEast);  
    copy-face(mesh[north-zone], tmpNorth);  
}
```

```
void compute-border(...) {  
  
    for (k=0; k<z-dim; k++)  
        for (j=0; j<y-dim; j++)  
            for (i=0; i<x-dim; i++)  
                /* Some matrix-based  
                 computations */  
}
```


```
void copy-face(...) {  
  
    for (k=0; k<z-dim; k++)  
        for (j=0; j<y-dim; j++)  
            /* Some matrix-based  
             computations */  
}
```

# Summary: Computation and Communication Periods, Sources of Parallelism

- Two levels of parallelism:
  - Coarse grain parallelism: INTER-ZONE parallelism
  - Fine grain parallelism : INTRA-ZONE parallelism

- Parallelism is repeated many times

- Computation Period
  - INTER-ZONE: mapped to CPUs **and** GPUs
  - INTRA-ZONE: mapped to CPUs **or** GPUs
- Communication Period
  - INTRA-ZONE: mapped to CPUs **or** GPUs



# Outline

- Analysis, parallelism characterization of NAS-MZ suite of benchmarks
- NPB-MZ Hybrid Design and Implementation
- Work distribution schemes for hybrid executions
- Evaluation Results

# NPB-MZ Hybrid Design and Implementation

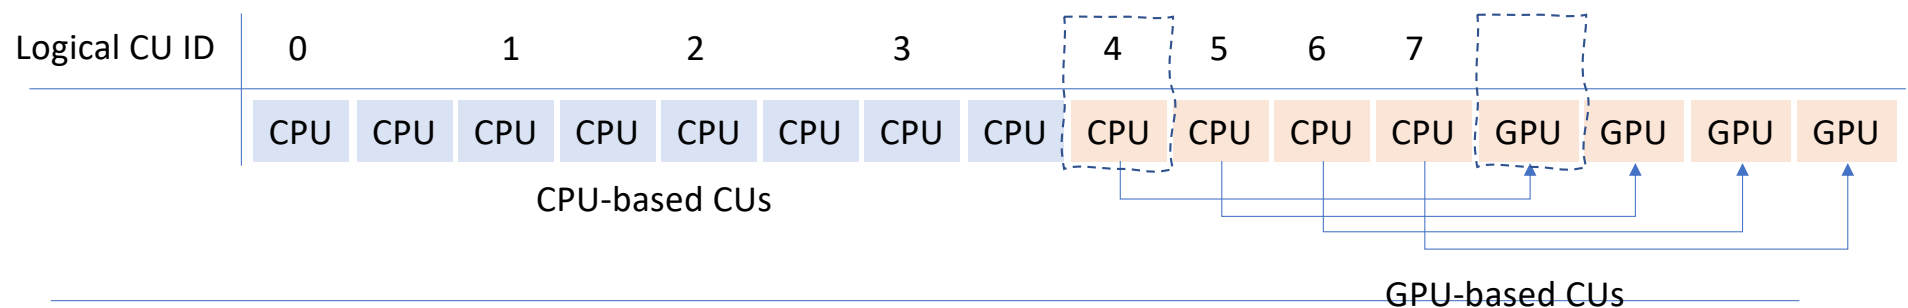
- Hybrid Execution Model: OpenMP+CUDA
  - Computing Unit Abstraction
    - `libCU-rtl`
  - Address Space Abstraction
    - `libAS-rtl`
  - Work Distribution Schemes
    - `libSCHEDULING-rtl`
- NPB-MZ Hybrid Parallelization

# Computing Units (I)

- Computing Units (CUs)
  - Logical view of all available CPUs and GPUs
  - A CU can be of type CPU or GPU
    - If CU type is of CPU
      - CU maps onto one or more physical <CPU>
    - If CU type is of GPU
      - CU maps to one pair of physical <CPU, GPU>
        - The CPU controls the GPU
- Computing Units (CUs)
  - NCU = number of available CUs
  - CUs are numbered from 0 up to the NCU-1
    - NCPU = Number of CPU-based CUs
    - NGPU = Number of GPU-based CUs
    - [0 .. NCPU-1] correspond to CPU-based CUs
    - [NCPU .. NCU-1] correspond to GPU-based CUs
- CUs and Threads
  - We run as many threads as the number of NCU

# Computing Units (II)

- Example:
  - Execution with 8 CUs = 4 x CPU-based CUs + 4 GPU-based CUs
    - CPU-based CU = 2 x CPUs
    - GPU-based CU = CPU + GPU
  - => 12 physical CPUS + 4 physical GPUS
  - 12 threads mapped onto 12 physical CPUs




# Hybrid Execution Model: OpenMP + CUDA

- How do we introduce the abstraction of a CU in OpenMP+CUDA?
  - Implement runtime support to bridge the two execution models
    - `libCU-rtl`
- Parallelism for CUs is created through `omp parallel` construct
  - 1 CPU-based CU: 1 OpenMP thread
  - 1 GPU-based CU: 1 OpenMP thread that controls a device
- CUs are formed through `OMP_PLACES`
- Name spaces for CU, OpenMP threads and CUDA devices
  - Logical CU id  $\Leftrightarrow$  OpenMP Thread id
    - `OpenMP Thread id = CU id`
  - Logical CU id  $\Leftrightarrow$  CUDA Device id
    - `CUDA Device id = CU id - NCPU`



# CUs definition and OpenMP thread places

- Example
  - `OMP_PLACES = "{0, 1}, {2, 3}, {4, 5}, {6, 7}, {8}, {9}, {10}, {11}"`
  - `OMP_PROC_BIND=TRUE`
  - `NCPU = 4` with 2 inner CPUs
  - `NGPU = 4`
  - Logical CU ids match OpenMP logical thread ids



# Runtime Support for CUs

- Small set of runtime primitives

- `RUNTIME::gpus` | Number of GPU-based CUs and CPU-based CUs
- `RUNTIME::cpus` |
- `RUNTIME::isGPU()` | Query if CU is mapped onto CPU or GPU
- `RUNTIME::isCPU()` | Get a CPU/GPU id in OpenMP/CUDA
- `RUNTIME::GPU()` |
- `RUNTIME::CPU()` | Get a number of inner threads for a CPU-based CU
- `RUNTIME::innerCPUs()` |
- `RUNTIME::getTask()` | Get a task mapped to the current CU
- `RUNTIME::getCU()` | Get/set the CU logical ID
- `RUNTIME::setCU()` |
- `RUNTIME::commitTask()` | Start/end of the execution of an assigned task
- `RUNTIME::executeTask()` |
- `RUNTIME::synchronize()` | Synchronize with corresponding device if CU is mapped to a GPU

# Code Transformation for CU support

## CU parallelism

```
#pragma omp parallel \  
    num_threads(RUNTIME::cpus+RUNTIME::gpus)  
{  
    unsigned int task;  
    task = RUNTIME::getTask();  
    while (task!=NO_TASK) {  
  
        unsigned int CU = RUNTIME::getCU();  
  
        RUNTIME::executeTask()  
  
        /* COMPUTATIONS */  
  
        if (RUNTIME::isGPU(CU)) RUNTIME::synchronize();  
  
        RUNTIME::commitTask();  
  
        task = RUNTIME::getTask();  
    } // while  
} // parallel
```

```
OMP_PROC_BIND=TRUE  
OMP_PLACES = "{0, 1},{2, 3},{4, 5},{6, 7},  
              {8},    {9},    {10}, {11}"
```

## CU computations

```
    unsigned int CU = RUNTIME::getCU();  
    if (RUNTIME::isCPU(CU)) {  
        NT = RUNTIME::innerCPUs();  
        #pragma parallel num_threads(NT) proc_bind(master)  
        {  
            /* COMPUTATIONS */  
        } // parallel  
    }  
    else if (RUNTIME::isGPU(CU)) {  
        cudaSetDevice(RUNTIME::GPU(CU));  
        some-GPU-kernel-1<<<grid, block, shared>>>(...);  
    }
```

# Distributed Address Space: OpenMP + CUDA

- How do we introduce the abstraction of a AS in OpenMP+CUDA?
  - Implement runtime support to bridge the two execution models
    - `libAS-rtl`
- Several Address Spaces (AS)
  - Host
  - Devices
- Name space for AS
  - $[0 \dots \text{NCPU}-1] \Rightarrow \text{HOST}$
  - $[\text{NCPU} \dots \text{NGPU}-1] \Rightarrow \text{DEVICES}$
- Monitor data placement
  - Maintain a map of data structures and AS

# Runtime Support for Data Placement

- Small set of runtime primitives

- `PLACEMENT::getAS()` | Determine AS for a memory region (e.g.: a zone)
- `PLACEMENT::alloc()` | Memory allocation (e.g.: a zone) in an AS
- `PLACEMENT::migrate()` | Migrate a memory (e.g.: a zone) to an AS
- `PLACEMENT::copy()` | Generate a copy of a zone to an AS
- `PLACEMENT::isGPU()` | Query if memory (e.g.: a zone) is mapped onto CPU or GPU
- `PLACEMENT::isCPU()`

# Work Distribution: OpenMP + CUDA

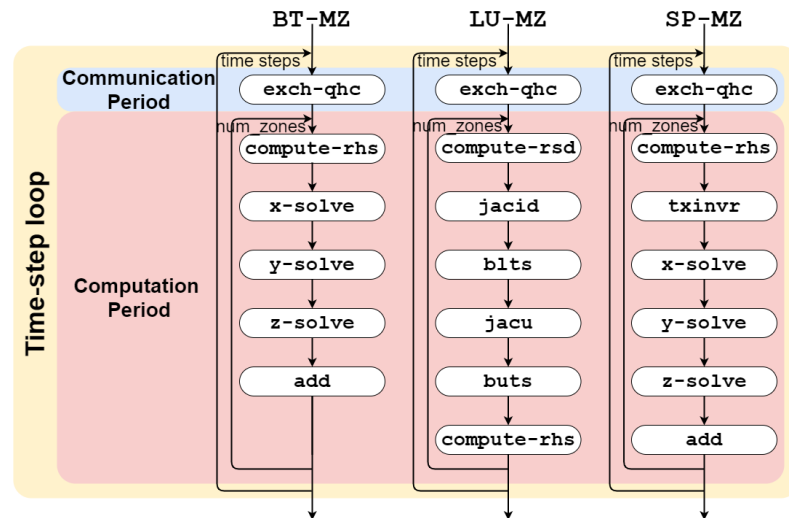
- OpenMP work distribution schemes are not suitable
  - TASKS: you do not have control over which thread will execute a task
  - LOOPS: available loop schedulers are not suitable, not designed for heterogeneity
    - STATIC, DYNAMIC, GUIDED
- Implement a runtime support to bridge the two execution models at the application level
  - libSCHEDULING-rtl

# Runtime Support for Scheduling

- Assume tasks are numbered from  $[0 \dots N_{\text{tasks}} - 1]$ 
  - Inputs for schedulings
    - $N_{\text{tasks}}$
    - NCPU
    - NGPU
- SCHEDULING-rtl
  - SCHEDULING::Static()
  - SCHEDULING::Dynamic()
  - SCHEDULING::pcf-Static()
  - SCHEDULING::pcf-Guided()
  - SCHEDULING::ClusteredGuided()

# Computation Period: Code Transformation

- INTER-ZONE parallelism



- Example: SP-MZ

```
SCHEDULE::Static(NUM_ZONES, RUNTIME::cpus, RUNTIME::gpus);
#pragma omp parallel \
    num_threads(RUNTIME::cpus+RUNTIME::gpus)
{
    unsigned int task;
    task = RUNTIME::getTask();
    while (task!=NO_TASK) {
        unsigned int zone = task;

        unsigned int CU = RUNTIME::getCU();
        if (PLACEMENT::getAS(mesh[zone])!=CU) PLACEMENT::migrate(mesh[zone], CU);

        RUNTIME::executeTask();

        compute-rhs(zone, ...);
        txinvr(zone, ...);
        x-zolve(zone, ...);
        y-solve(zone, ...);
        z-solve(zone, ...);
        add(zone, ...);

        if (RUNTIME::isGPU(CU)) RUNTIME::synchronize();

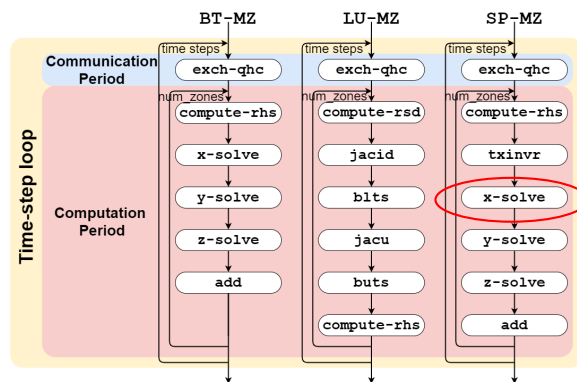
        RUNTIME::commitTask();
        task = RUNTIME::getTask();
    } // while
} // parallel
```

OMP\_PROC\_BIND=TRUE



# Computation Period: Code Transformation

- INTRA-ZONE parallelism



```
void x-solve(unsigned int zone, ...) {
    unsigned int CU = RUNTIME::getCU();
    if (RUNTIME::isCPU(CU)) {
        NT = RUNTIME::innerCPUs();
        #pragma parallel for schedule(static)\
            num_threads(NT) proc_bind(master)
        for (k=0; k<z-dim; k++)
            for (j=0; j<y-dim; j++)
                for (i=0; i<x-dim; i++)
                    /* Some matrix-based
                     * computations */
    }
    else if (RUNTIME::isGPU(CU)){
        x-xolve-kernel-1<<<grid, block, shared>>>(zone,...);
    }

    /* Several other computations computations in the
     * previous form */
}
```

# Communication Period: Border Processing

```
for (zone=0; zone<NUM-ZONES; zone++) {
    east-zone = adjacency-east[zone]
    north-zone = adjacency-north[zone];
    east-AS = PLACEMENT::getAS(mesh[east-zone]);
    north-AS = PLACEMENT::getAS(mesh[north-zone]);
    zone-AS = PLACEMENT::getAS(mesh[zone]);

    PLACEMENT::pack-face(tmpEast, east-AS, mesh[east-zone]);
    PLACEMENT::pack-face(tmpNorth, north-AS, mesh[north-zone]);
    if (PLACEMENT::isGPU(zone-AS)) {
        CUDA::setDevice(RUNTIME::GPU(zone-AS)):
        gpu-compute-border<<< grid, block, shared>>>(mesh[zone],
                                                    tmpEast,
                                                    tmpNorth);
    }
    else if (PLACEMENT::isCPU(zone-AS)) {
        cpu-compute-border(mesh[zone], tmpEast, tmpNorth);
    }
    PLACEMENT::unpack-face(tmpEast, east-AS, mesh[east-zone]);
    PLACEMENT::unpack-face(tmpNorth, north-AS, mesh[north-zone]);
}
```

# Communication Period: Data Packing

```
void PLACEMENT::pack-face(buffer, zoneAS, mesh-zone) {  
    myAS = PLACEMENT::getAS();  
    if (myAS!=zoneAS) {  
        // remote packing of border data and transfer to myAS  
        if (PLACEMENT::isGPU(zoneAS) {  
            CUDA::setDevice(RUNTIME::GPU(zone-AS));  
            gpu-pack-face<<<grid, block, shared>>>(...);  
            CUDA::cudaMemCpy(buffer, ..., cudaDeviceToHost);  
        } else if (PLACEMENT::isCPU(zoneAS){  
            cpu-pack-face(...);  
        }  
    } else {  
        // pack border data  
    }  
}  
void PLACEMENT::unpack-face(buffer, zoneAS, mesh-zone);
```


```
void cpu-pack-face(buffer, mesh-zone);  
void cpu-unpack-face(buffer, mesh-zone);
```

```
__global__ gpu-pack-face(buffer, mesh-zone);  
__global__ gpu-unpack-face(buffer, mesh-zone);
```

- In a hybrid execution the **CUDA::cudaMemCpy** is the source of overhead

# Communication Period: Data placement and Communication

- Trade off between the number of data transfers and data placement
  - $NCU = 4$
  - Zone B
    - Assigned to a CU where 1 adjacent zone is assigned to other CUs
  - Zone A
    - Assigned to a CU where all 4 adjacent zones are assigned to the same CU
  - Zone C
    - Assigned to a CU where 3 adjacent zones are assigned to other CUs



# Outline

- Analysis, parallelism characterization of NAS-MZ suite of benchmarks
- NPB-MZ Hybrid Design and Implementation
- Work distribution schemes for hybrid executions
- Evaluation Results

# Work Distribution Schemes For OpenMP+CUDA

- Baseline Schedulers
  - STATIC
  - DYNAMIC
- Performance Factor Conversion (PCF) Schedulers
  - PCF-STATIC
  - PCF-GUIDED
- Self Adaptive Schedulers
  - CLUSTERED GUIDED

# STATIC scheduler

- Tasks are identified from 0 up to the number of tasks -1
- `NTASKS = NUM_ZONES;`
- `CU-WORK = NTASKS/NCU;`
- If remaining tasks, those are assigned to lower CUs within the numbering
  - CPUs first, then the GPUs
- `<start,end>`: identifies the set of tasks assigned to a CU
  - `start` = initial task
  - `end` = final task
- Maximizes adjacency, so minimizes communications

# DYNAMIC scheduler

- Fast CUs (e.g.: GPUs) will tend to get more work
- Slow CUs (e.g.: CPUs) will tend to get less work
- `CHUNK = 1;`
- Memorizes TASK-CU mapping
  - Avoids memory migration overheads
- Breaks adjacency, so will tend to increase the communications



# PCF-STATIC scheduler

- Variant of STATIC scheduler
- Based on a Performance Conversion Factor (PCF)
  - PCF = relation between the computational power of CPU-based CUs and GPU-based CUs
  - Example
    - PCF = 2, means that GPU-based CUs are twice faster than CPU-based CUs
    - PCF = 1, means both CPU and GPU based CUs have same computation power

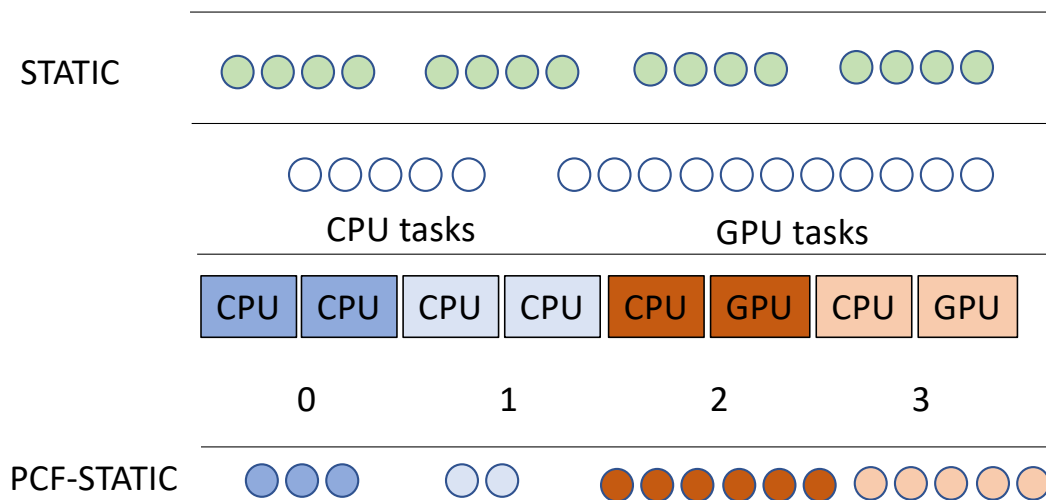
# PCF-STATIC scheduler

- Divide in two parts the set of tasks
  - $NTASKS = PCF \times CPU-TASKS + CPU-TASKS = (PCF+1) \times CPU-TASKS$
  - $GPU-TASKS = NTASKS - CPU-TASKS$
- Apply STATIC over  $CPU-TASKS$  and among all CPU-based CUs
- Apply STATIC over  $GPU-TASKS$  and among all GPU-based CUs
- Maximizes adjacency, so minimizes communications

# PCF-STATIC scheduler

- Example:

- PCF = 2
- NCU = 4
  - CPU-based CU = 2
  - GPU-based CU = 2
- NTASKS = 16
- CPU-TASKS = 5
- GPU-TASKS = 11



PCF-STATIC

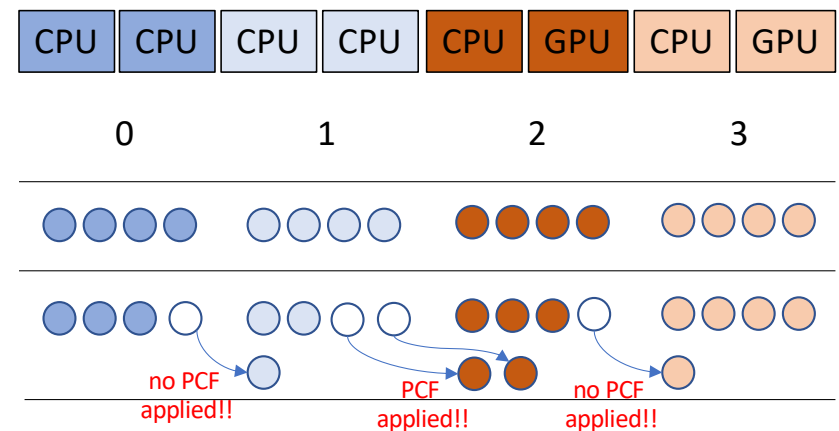
<start,end>: interval of tasks assigned to a CU

STATIC: <0,3> <4,7> <8,11> <12,15>

PCF-STATIC: <0,2> <3,4> <5,10> <11,15>

# PCF-GUIDED scheduler

- Initially, apply a STATIC scheduling
- Execute and monitor task execution time
- Adapt the work distribution so that it gets balanced
  - ALL-WORK = addition of all task execution times factorized with PCF
  - $WORK-CU = ALL-WORK / NCU$
  - Keep the assigned work per CU close to WORK-CU
- Maximizes adjacency, so minimizes communications



<start,end>: interval of tasks assigned to a CU

STATIC:	<0,3>	<4,7>	<8,11>	<12,15>
PCF-GUIDED:	<0,2>	<3,5>	<6,10>	<11,15>

# PCF-GUIDED scheduler

- Samples of execution time are taken from `RUNTIME::executeTask()` and `RUNTIME::commitTask()`
- Code scheme for work balance

```
for cu = 0, NCU-1
    WORK[cu]= "add tasks execution times"

for cu = 0, NCU-2
    while WORK[cu]>WORK-CU /* With some threshold */
        <start[cu],end[cu]-->
        WORK[cu] -= "execution time of end task"
        WORK[cu+1] += "execution time of end task refactorized with PCF"

    while WORK[cu]<WORK-CU /* With some threshold */
        <start[cu],end[cu]++>
        WORK[cu] += "execution time of end task refactorized with PCF"
        WORK[cu+1] -= "execution time of end task"

    start[cu+1] = end[cu]
```

```
#pragma omp parallel \
    num_threads(RUNTIME::cpus+RUNTIME::gpus)\
    proc_bind(true)
{
    unsigned int task;
    task = RUNTIME::getTask();
    while (task!=NO_TASK) {
        unsigned int zone = task;

        unsigned int CU = RUNTIME::getCU();
        if (PLACEMENT::zonePlacement[zone]!=CU &&
            RUNTIME::isGPU(CU))
            PLACEMENT::migrateZone(CU);

        RUNTIME::executeTask();

        /* Computational Stages */

        if (RUNTIME::isGPU(CU)) RUNTIME::synchronize();


        RUNTIME::commitTask();

        task = RUNTIME::getTask();
    } // while
} // parallel
```

g for


# CLUSTERED GUIDED scheduler

- Set of tasks is divided in two clusters
  - One for CPU-based CUs, the other for GPU-based CUs
  - PIVOT: separating point for the two clusters
- At each instance of the scheduler, the PIVOT moves left/right to balance the work assigned to each cluster
  - Sort of a dichotomic search
- The scheduler balances the work distribution within each cluster



# CLUSTERED GUIDED scheduler

- The scheduler evolves at runtime, switching between different states:
  - **INIT**: PIVOT is initialized (e.g.:  $\text{PIVOT} = \text{NCPU}$ ). Applies a **STATIC** scheduling in both clusters.
  - **MOVE**: Moves PIVOT, according to where the maximum execution time has occurred.
  - **PROBE**: Samples execution time for each CU. Determines the maximum value observed in each cluster.
  - **BALANCE**: Applies a PCF-GUIDED scheduler with  $\text{PCF}=1$  to each cluster.
  - **STEADY**: Records final configuration for work distribution.



# CLUSTERED GUIDED scheduler

## Data Structures:

Pivot: Index that indicates the border between tasks assigned to CPUs and tasks assigned to GPUs

DEC: Value to decrement the Pivot


INC: Value to increment the Pivot

TaskTime []: Vector of  $N_{zones}$  elements, each one describing the execution time of a task

CUtime []: Vector of  $N_{CUs}$  elements, each one describing the execution time of CU

FirstTask []: Vector of  $N_{CUs}$  elements, each one describing the  $T_{first}$  task assigned to a CU

LastTask []: Vector of  $N_{CUs}$  elements, each one describing the  $T_{last}$  task assigned to a CU



## MOVE:

$T_{cpu}$  = max time for CPUs

$T_{gpu}$  = max time for GPUs

if  $T_{cpu} > T_{gpu}$

INC = INC / 2

goingLEFT = false

goingRIGHT = true

pivot += INC

if  $T_{gpu} > T_{cpu}$

DEC = DEC / 2

goingLEFT = true

goingRIGHT = false

pivot -= DEC

## BALANCE (CPUs/GPUs):

for  $cu = 1, N_{CUs}$

TotalWork = CUtime[cu]

WorkPerCU = TotalWork /  $N_{CUs}$

for  $cu = 1, N_{CUs}$

First = FirstTask[cu]

Last = LastTask[cu]

Work =  $\sum_{First}^{Last} T_i$

DIFF = | Work - WorkPerCU |

if ( Work < WorkPerCU && DIFF > TH )

while ( Work < WorkPerCU ) Last++; Work +=  $T_{last}$

else if ( Work > WorkPerCU && DIFF > TH )

while ( Work > WorkPerCU ) Last--; Work -=  $T_{last}$


LastTask[cu] = Last

FirstTask[cu] = Last + 1

## INIT:

Pivot =  $N_{CPUs}$


INC = DEC =  $N_{tasks} / 2$




Number of steps to STEADY state:  $\log_2(N_{tasks}) \times 4$




# Examples of Clustered Guided



# Examples of Clustered Guided



# Examples of Clustered Guided









# Outline

- Analysis, parallelism characterization of NAS-MZ suite of benchmarks
- NPB-MZ Hybrid Design and Implementation
- Work distribution schemes for hybrid executions
- Evaluation Results


# Environment

- AMD CTE @ BSC
  - AMD EPYC 7742 @ 2.250GHz (64 cores)
  - 2 x GPU AMD Radeon Instinct MI50 with 32GB
- Software stack
  - GCC 8.3.1–4
  - Radeon Open Compute (ROCm) Runtime software stack 3.5.0
- Hybrid NPB-MZ implementation
  - Parallel CUDA multi-gpu version of the NPB-MZ in C++
  - Parallel (OpenMP) implementation in Fortran of NPB-MZ
  - Runtime Libraries: `libCU-rtl`, `libPLACEMENT-rtl`, `libSCHEDULING-rtl`







# LU-MZ: Compute and Communication



# LU-MZ: Overall Performance




# BT-MZ: Compute and Communication

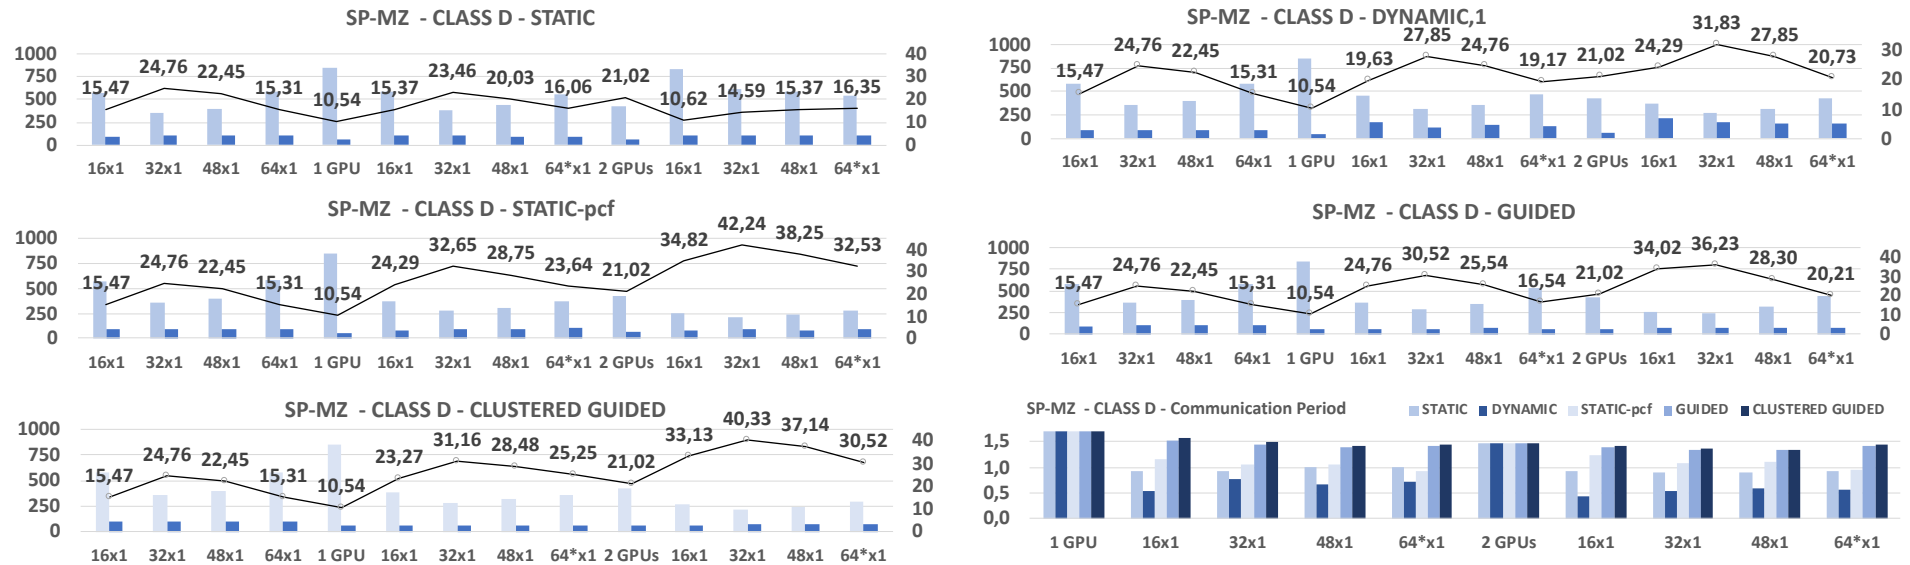





# BT-MZ: Overall Performance



# SP-MZ: Compute and Communication



# SP-MZ: Overall Performance



# Conclusions

- Case of study for hybrid computing (NPB-MZ), using OpenMP+CUDA
  - Methodology based on runtime implementation for bridging the two execution models
    - CU abstraction, AS abstraction at the programming model level
    - Work distribution schemes for hybrid executions
      - Based on rough performance comparison (PFC)
      - Monitoring task execution times (Clustered Guided)
- In general, the more GPUs in place, less space for using the CPUs
  - From a pure “FLOPS” perspective, we should try to use the CPUs for other purposes
    - Design runtime systems that solve programming limitations: memory allocators for dynamically changing data layouts, work distribution schemes, even code optimization
    - ...

# Heterogeneous Computing for Scientific Applications

Marc Gonzalez Tallada

Associate professor

Computer Architecture Department  
Technical University of Catalonia



UNIVERSITAT POLITÈCNICA  
DE CATALUNYA  
BARCELONATECH